

Class Evaluation

- The class will start 5 minutes late
- Please give us the feedback on 6.837
- Thank you!!!

<http://web.mit.edu/subjectevaluation>

Graphics Hardware



MIT EECS 6.837
Wojciech Matusik

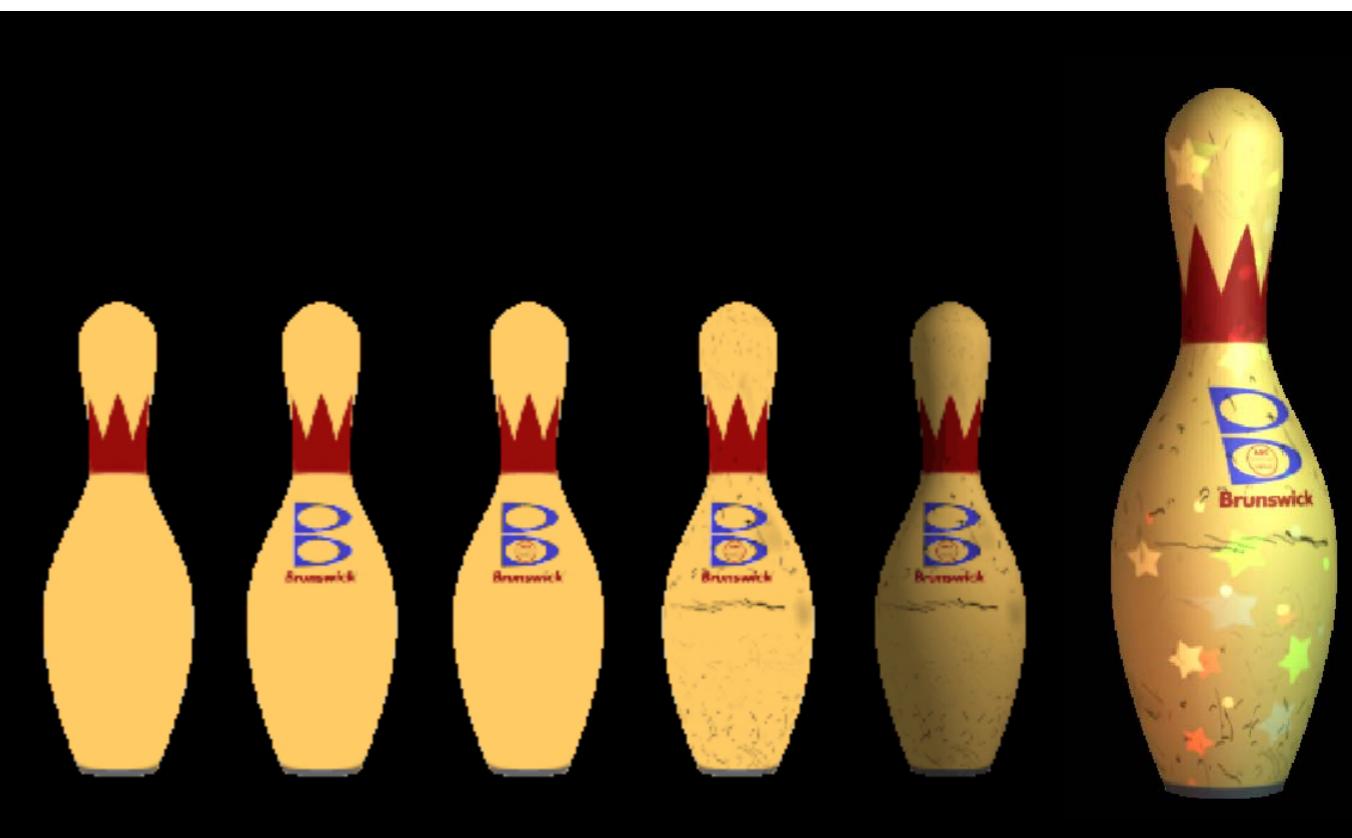
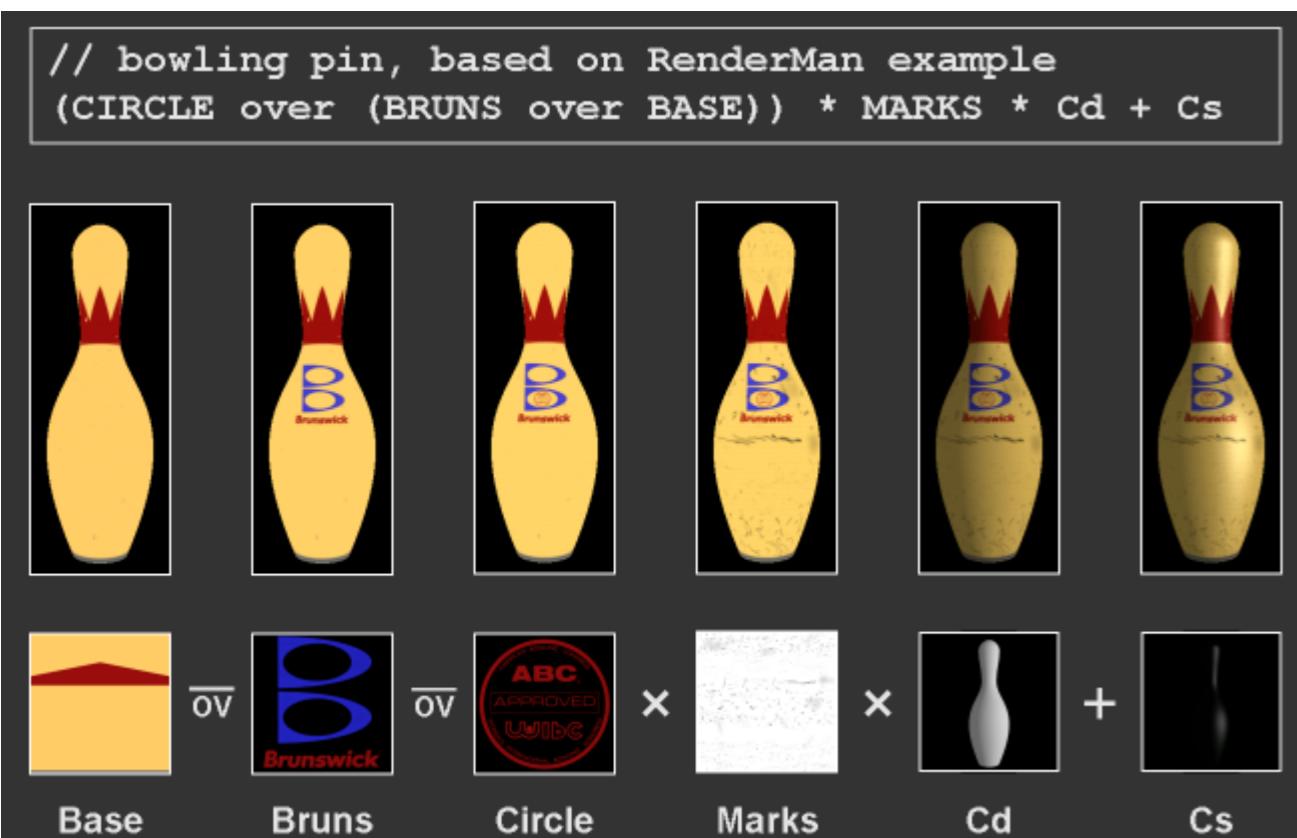
with slides from
Fredo Durand, Jonathan Ragan-Kelley, Hanrahan & Akeley, Gary
McTaggart, NVIDIA, ATI

Final exam

- 4 pages (2 double-sided sheets) of notes
- Everything we saw this semester
 - With emphasis on 2nd part
- Similar to quiz 1
- Will target a 2-hour subject

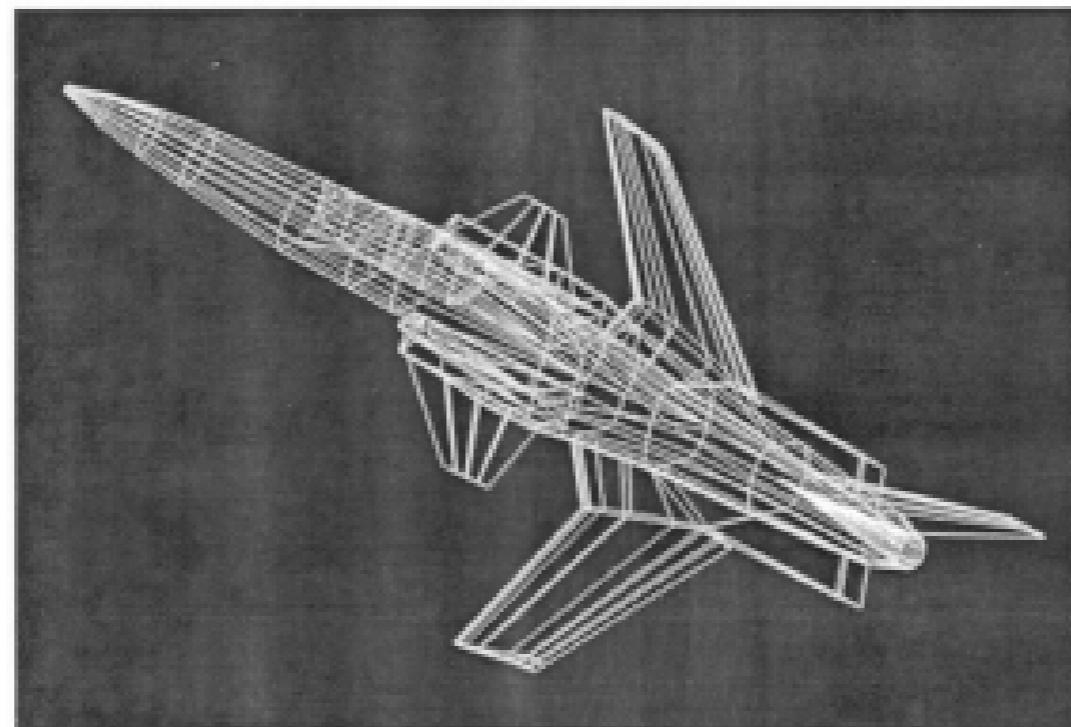
Modern Graphics Hardware (“GPU”)

- Hardware implementation of the *rendering pipeline*
- Programmability (“shaders”) within specific stages of the pipeline
 - Recent, last 10 years
 - At the vertex and pixel level



First Generation - Wireframe

- **Vertex:** transform, clip, and project
- Rasterization: color interpolation (points, lines)
- Fragment: overwrite
- Dates: prior to 1987



Slide from Stanford CS448A, Akeley & Hanrahan, 2007

Second Generation - Shaded Solids

- Vertex: lighting calculation
- Rasterization: depth interpolation (triangles)
- Fragment: depth buffer, color blending
- Dates: 1987-1992



Slide from Stanford CS448A, Akeley & Hanrahan, 2007

Third Generation - Texture Mapping

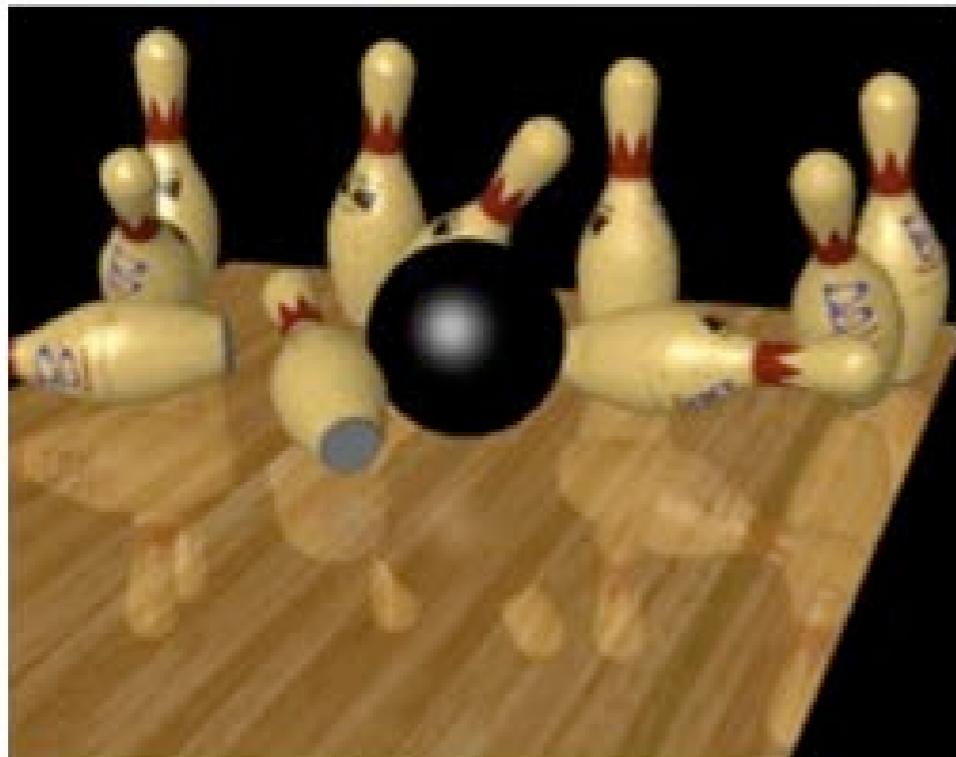
- Vertex: texture coordinate transformation
- **Rasterization:** texture coordinate interpolation
- **Fragment:** texture evaluation, antialiasing
- Dates: 1992-2000



Slide from Stanford CS448A, Akeley & Hanrahan, 2007

Fourth Generation - Programmable Shading

- **Vertex:** programmable transformation
- Rasterization: user attribute interpolation
- **Fragment:** programmable color computation
- Dates: 2000-2009



Fifth Generation - Programmable Pipeline

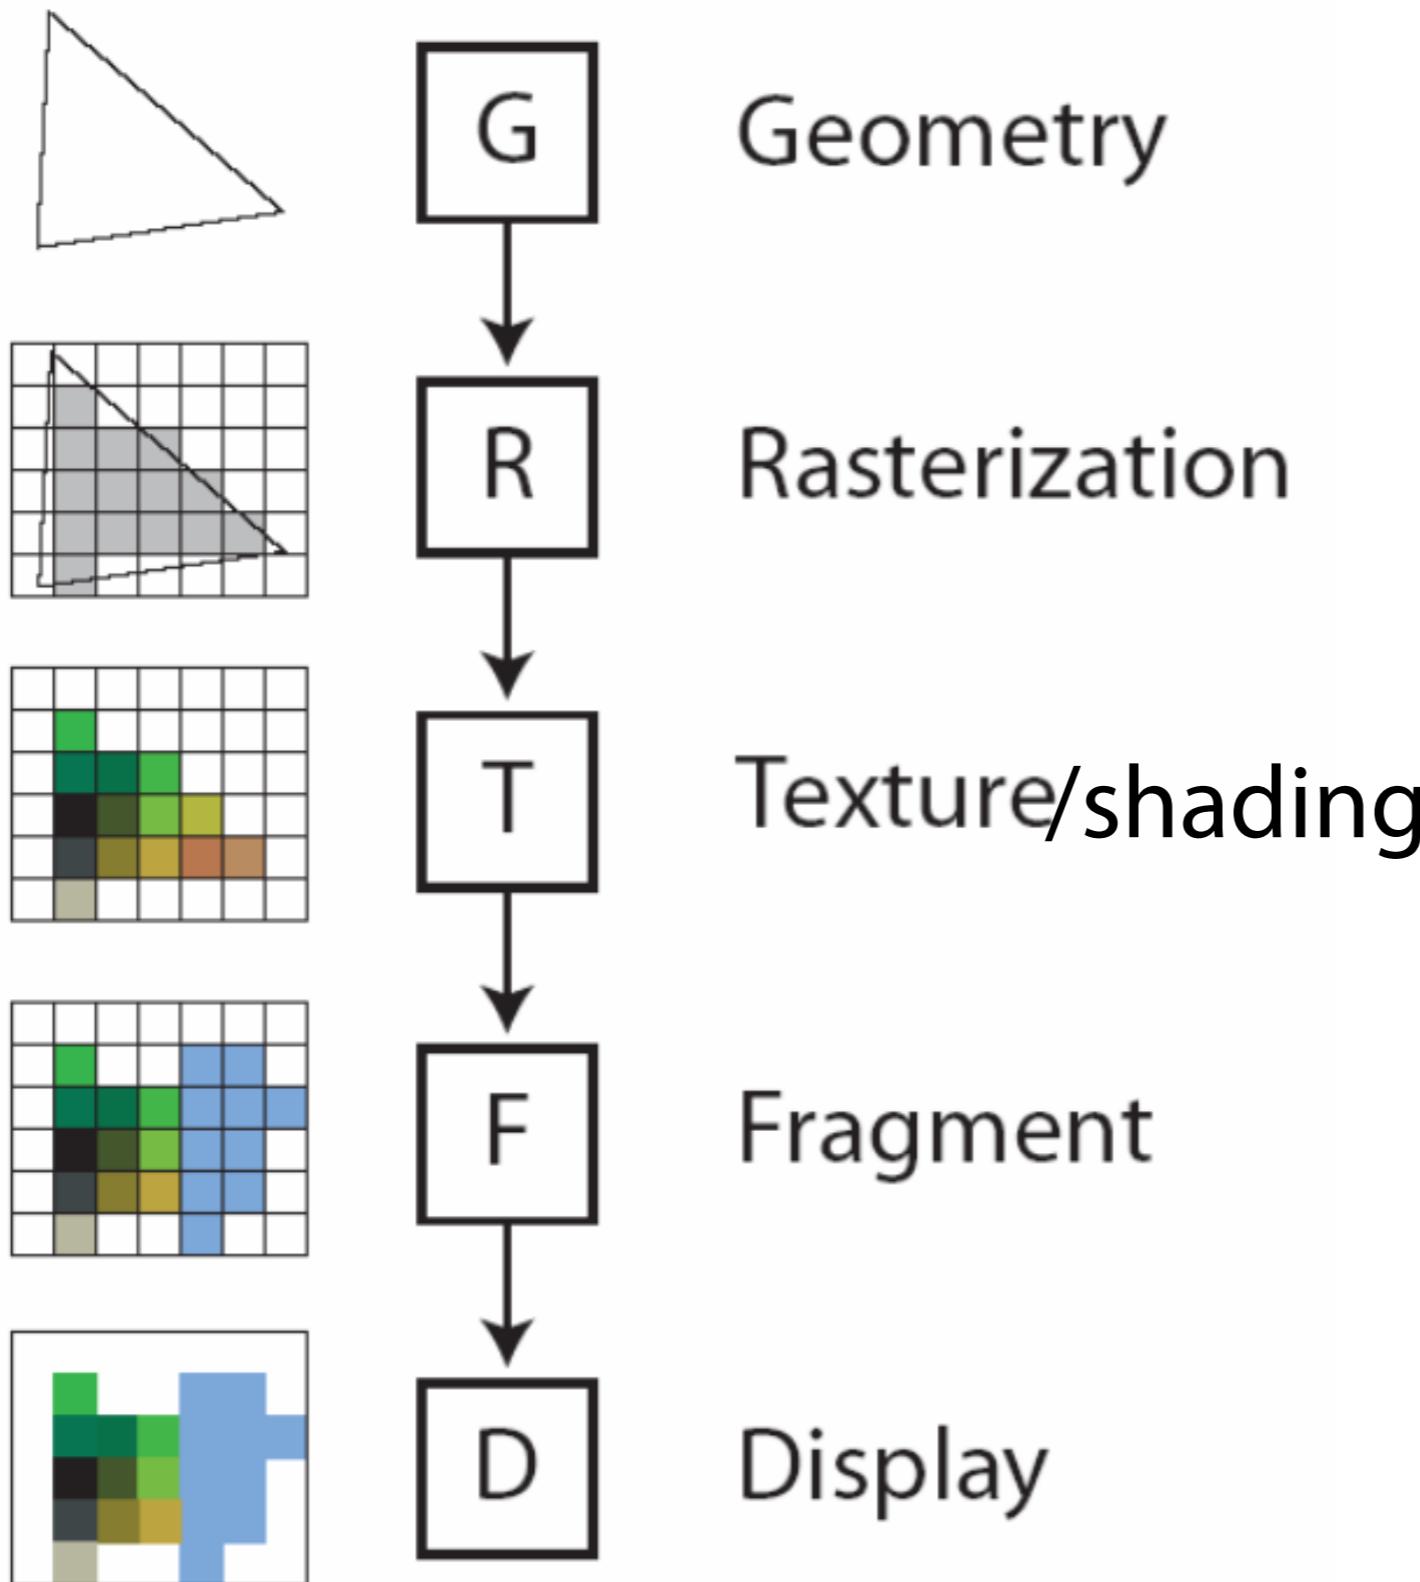
- Vertex,
Rasterization,
Fragment, ???: arbitrary parallel programs
- Dates: 2009+



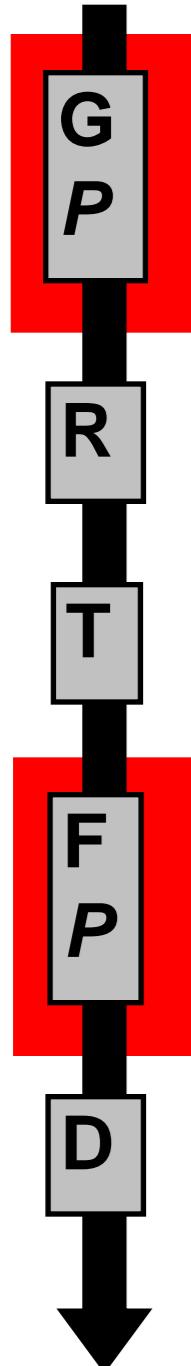
Questions?



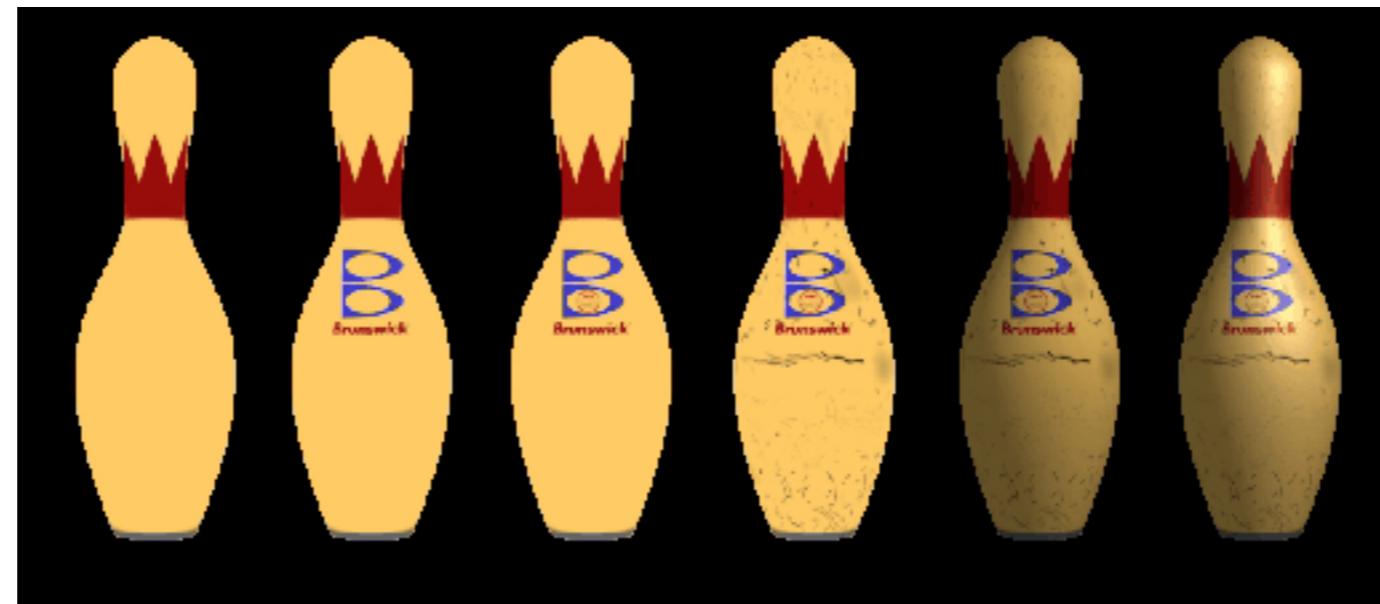
The Graphics Pipeline



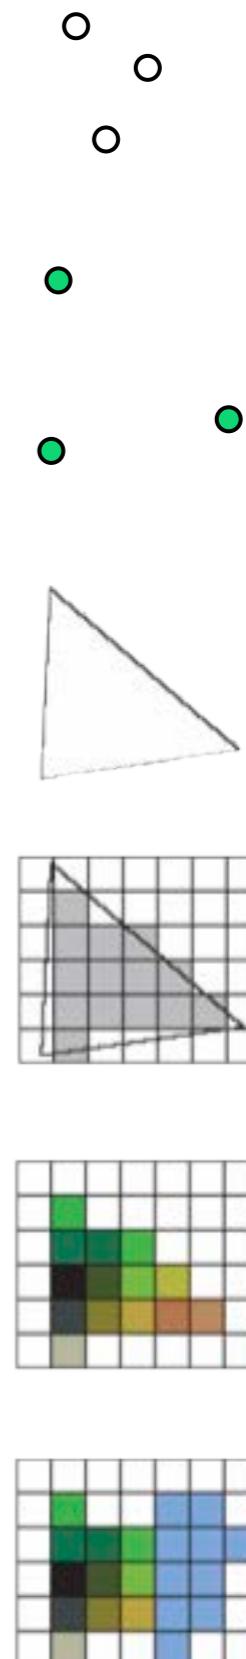
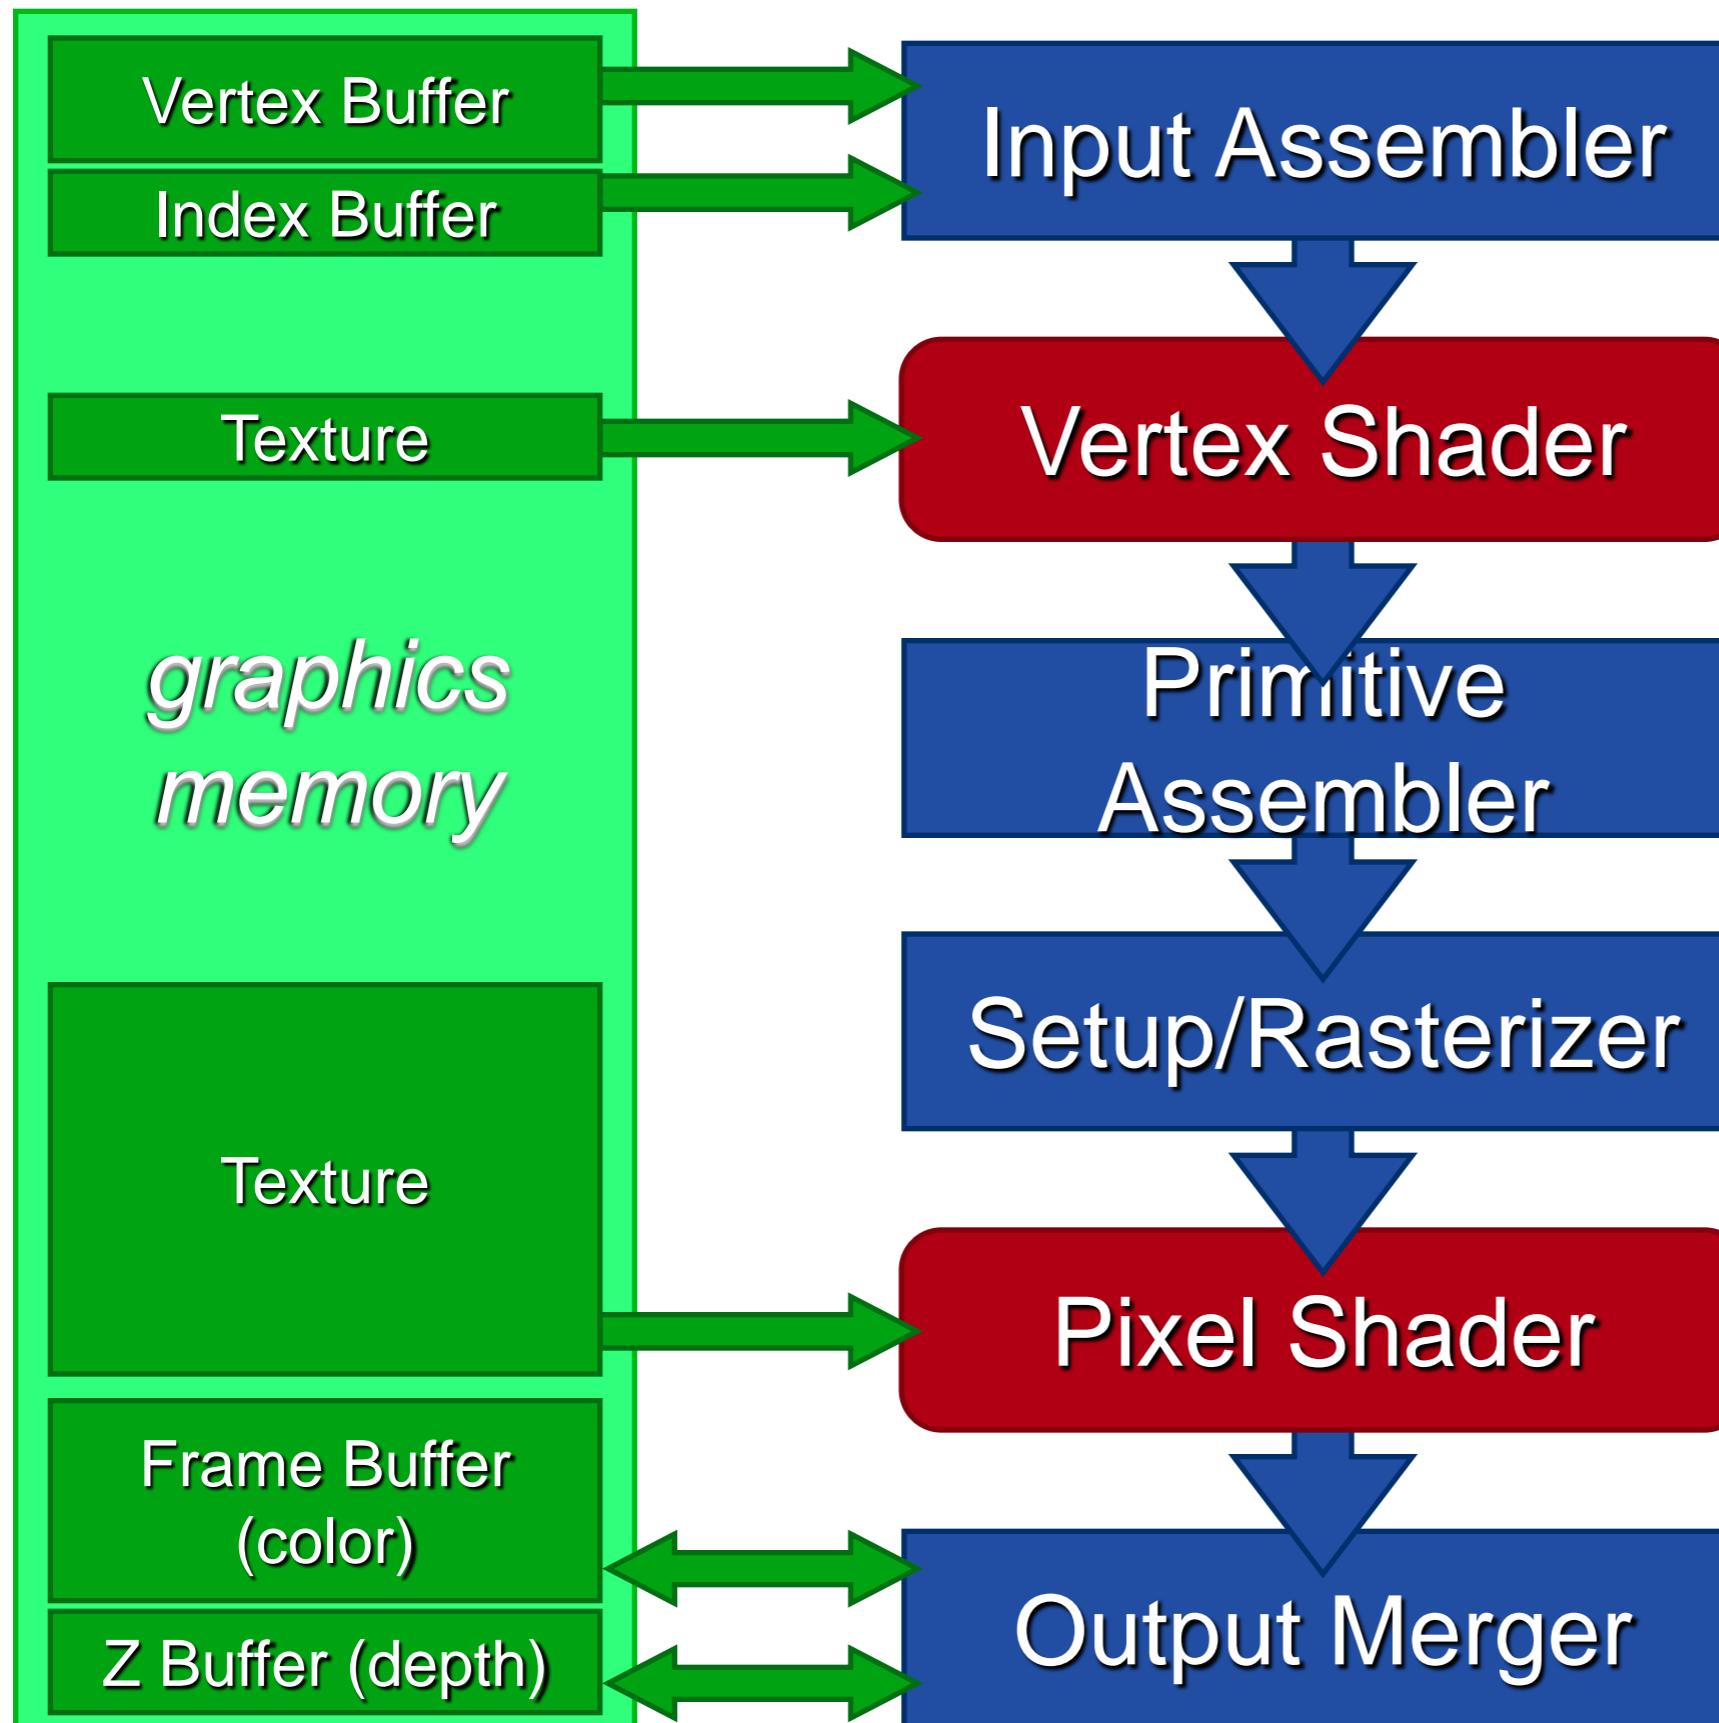
Programmable Graphics Hardware



- Geometry and Pixel (fragment) stage become programmable
 - Elaborate appearance
 - More and more general-purpose computation (GPGPU)



The Direct3D 9 Graphics Pipeline



Questions?

Vertex Shaders

$f(\text{position, attributes}) \longrightarrow \text{new position, attributes}$
purely functional (no side-effects)

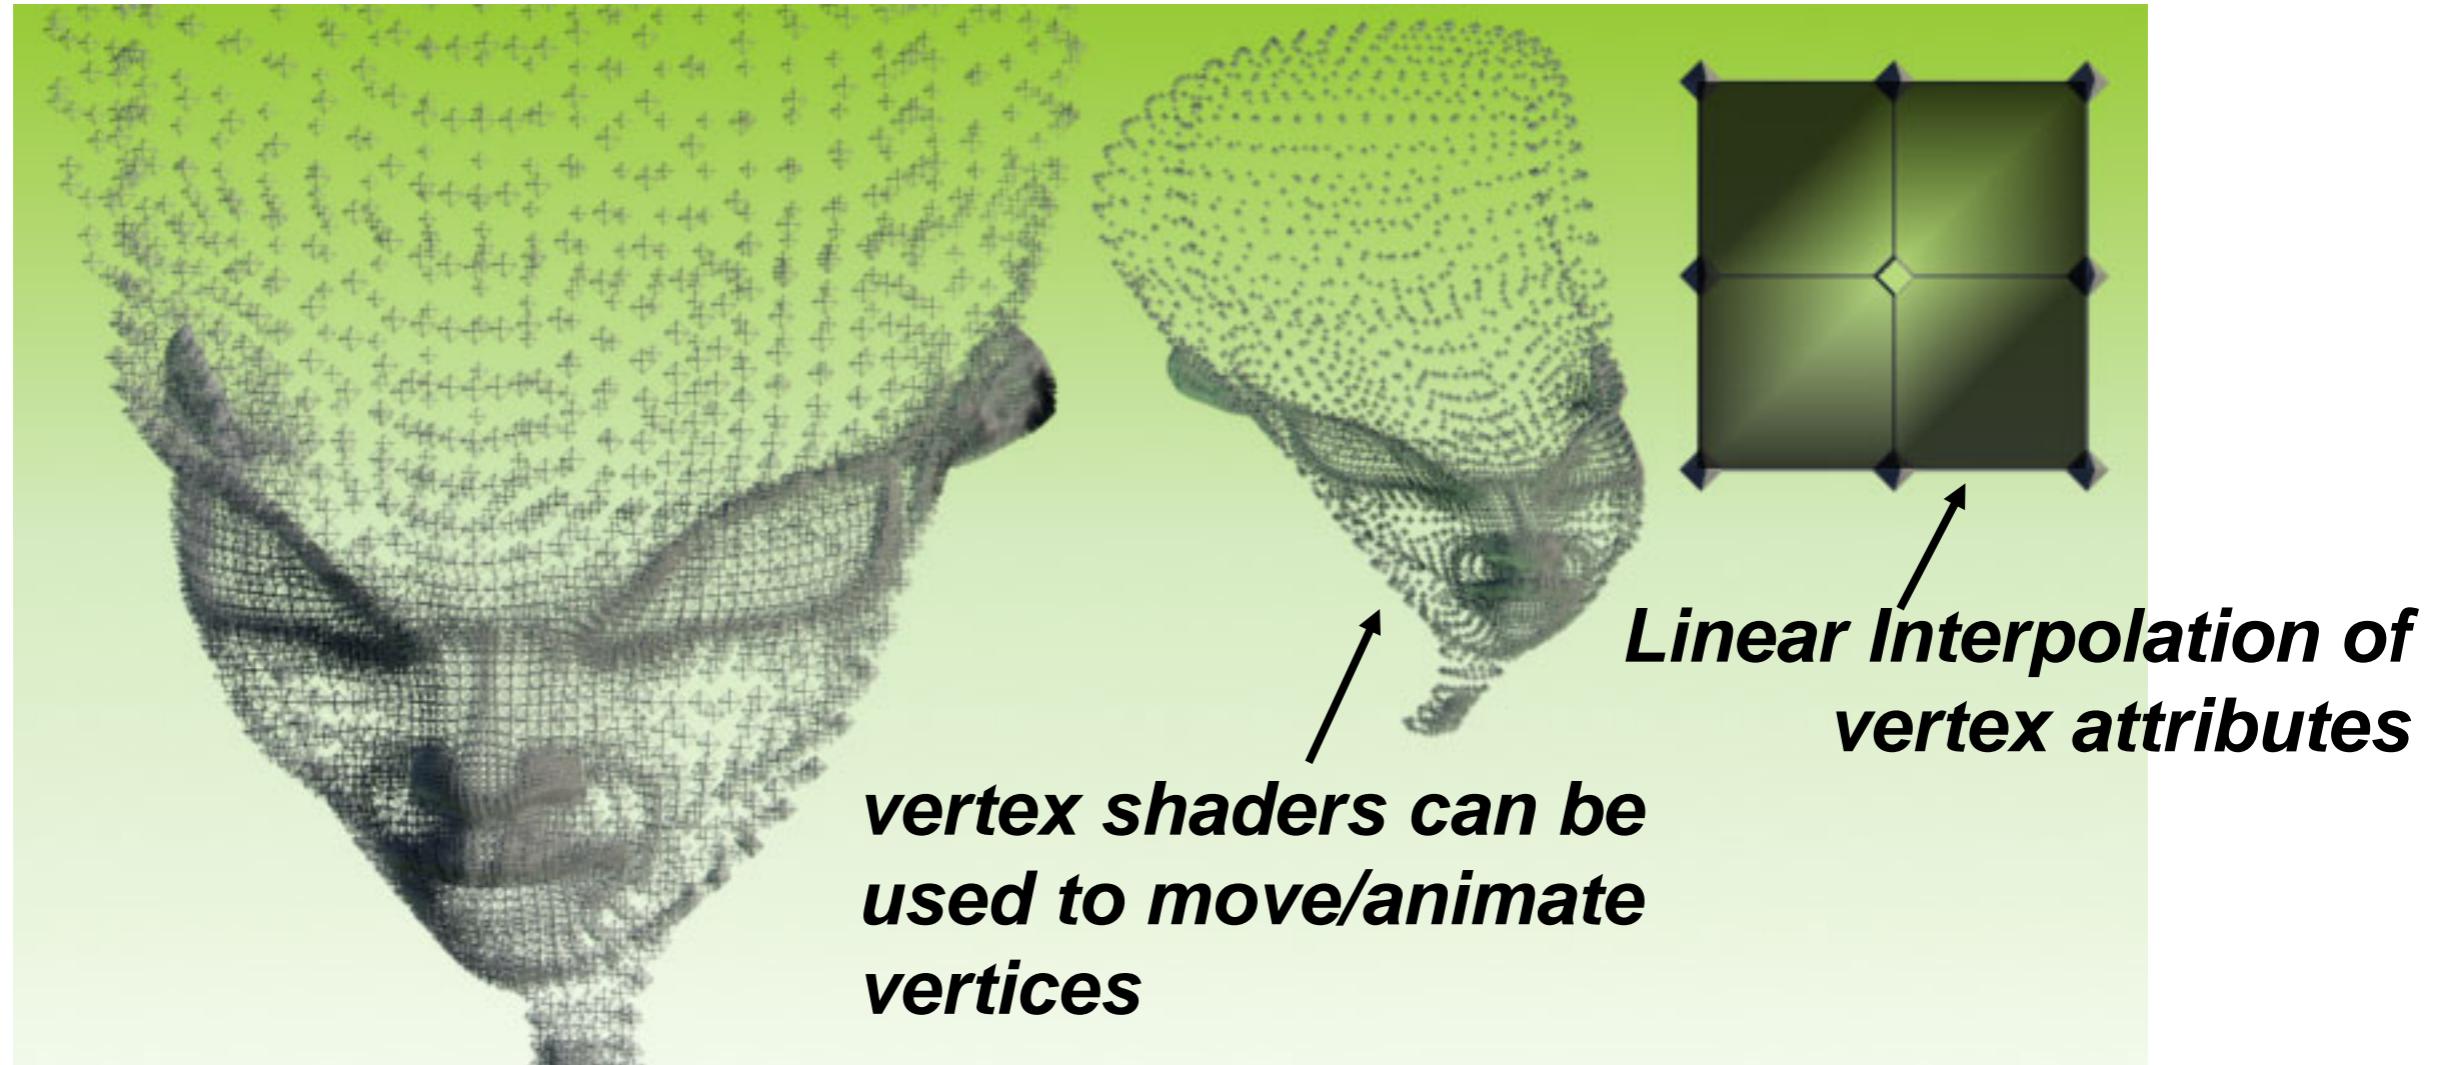
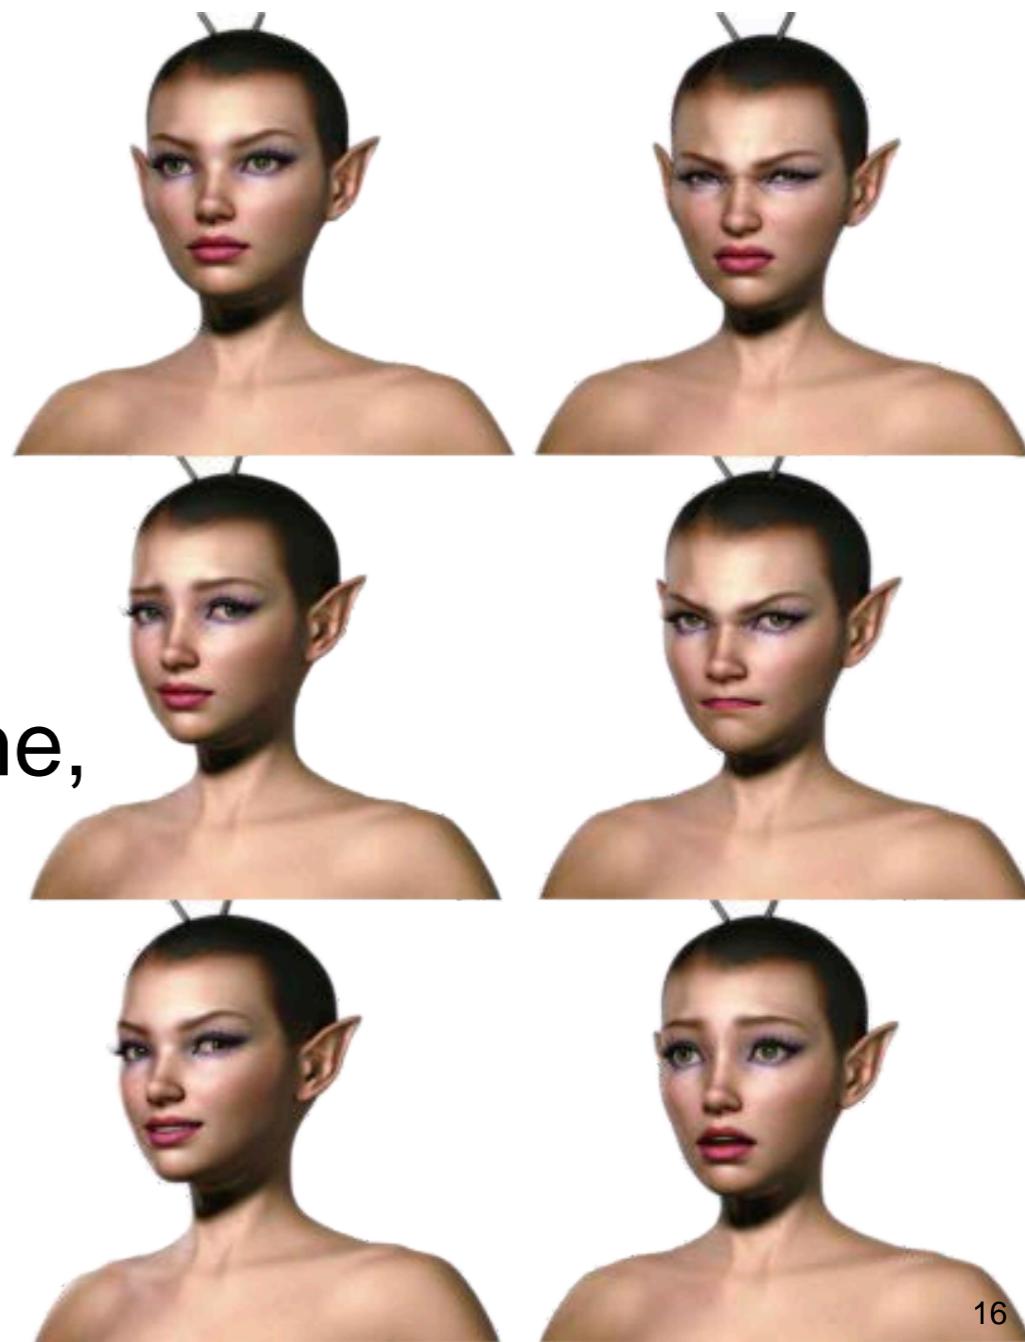


Figure from NVIDIA

Vertex Shader: Blendshapes

- 50 face geometries
 - angry, happy, sad, move eyebrow, ...
- Each target stored as difference vector
 - For each vertex:
average position + 50 differences
- Result is a weighted sum of all targets
 - Only transmit new weights each frame,
the targets remain in GPU memory
 - Big multiply-add
 - Per active blend target
 - Per attribute



Vertex Shader: Skinning (SSD)

1. Transform each vertex p_i with each bone as if it was rigidly tied to it
2. Blend the results using bone weights

```
float4 skin(float4 restPos,  
          uniform float4x4 xform[N_BONES],  
          uniform float weight[N_BONES])  
{  
    float4 outPos = float4(0,0,0,0)  
  
    for (int b = 0; b < N_BONES; b++) {  
        outPos += weight[b]*mul(xform[b], restPos)  
    }  
  
    return outPos  
}
```



Other jobs for vertex shaders

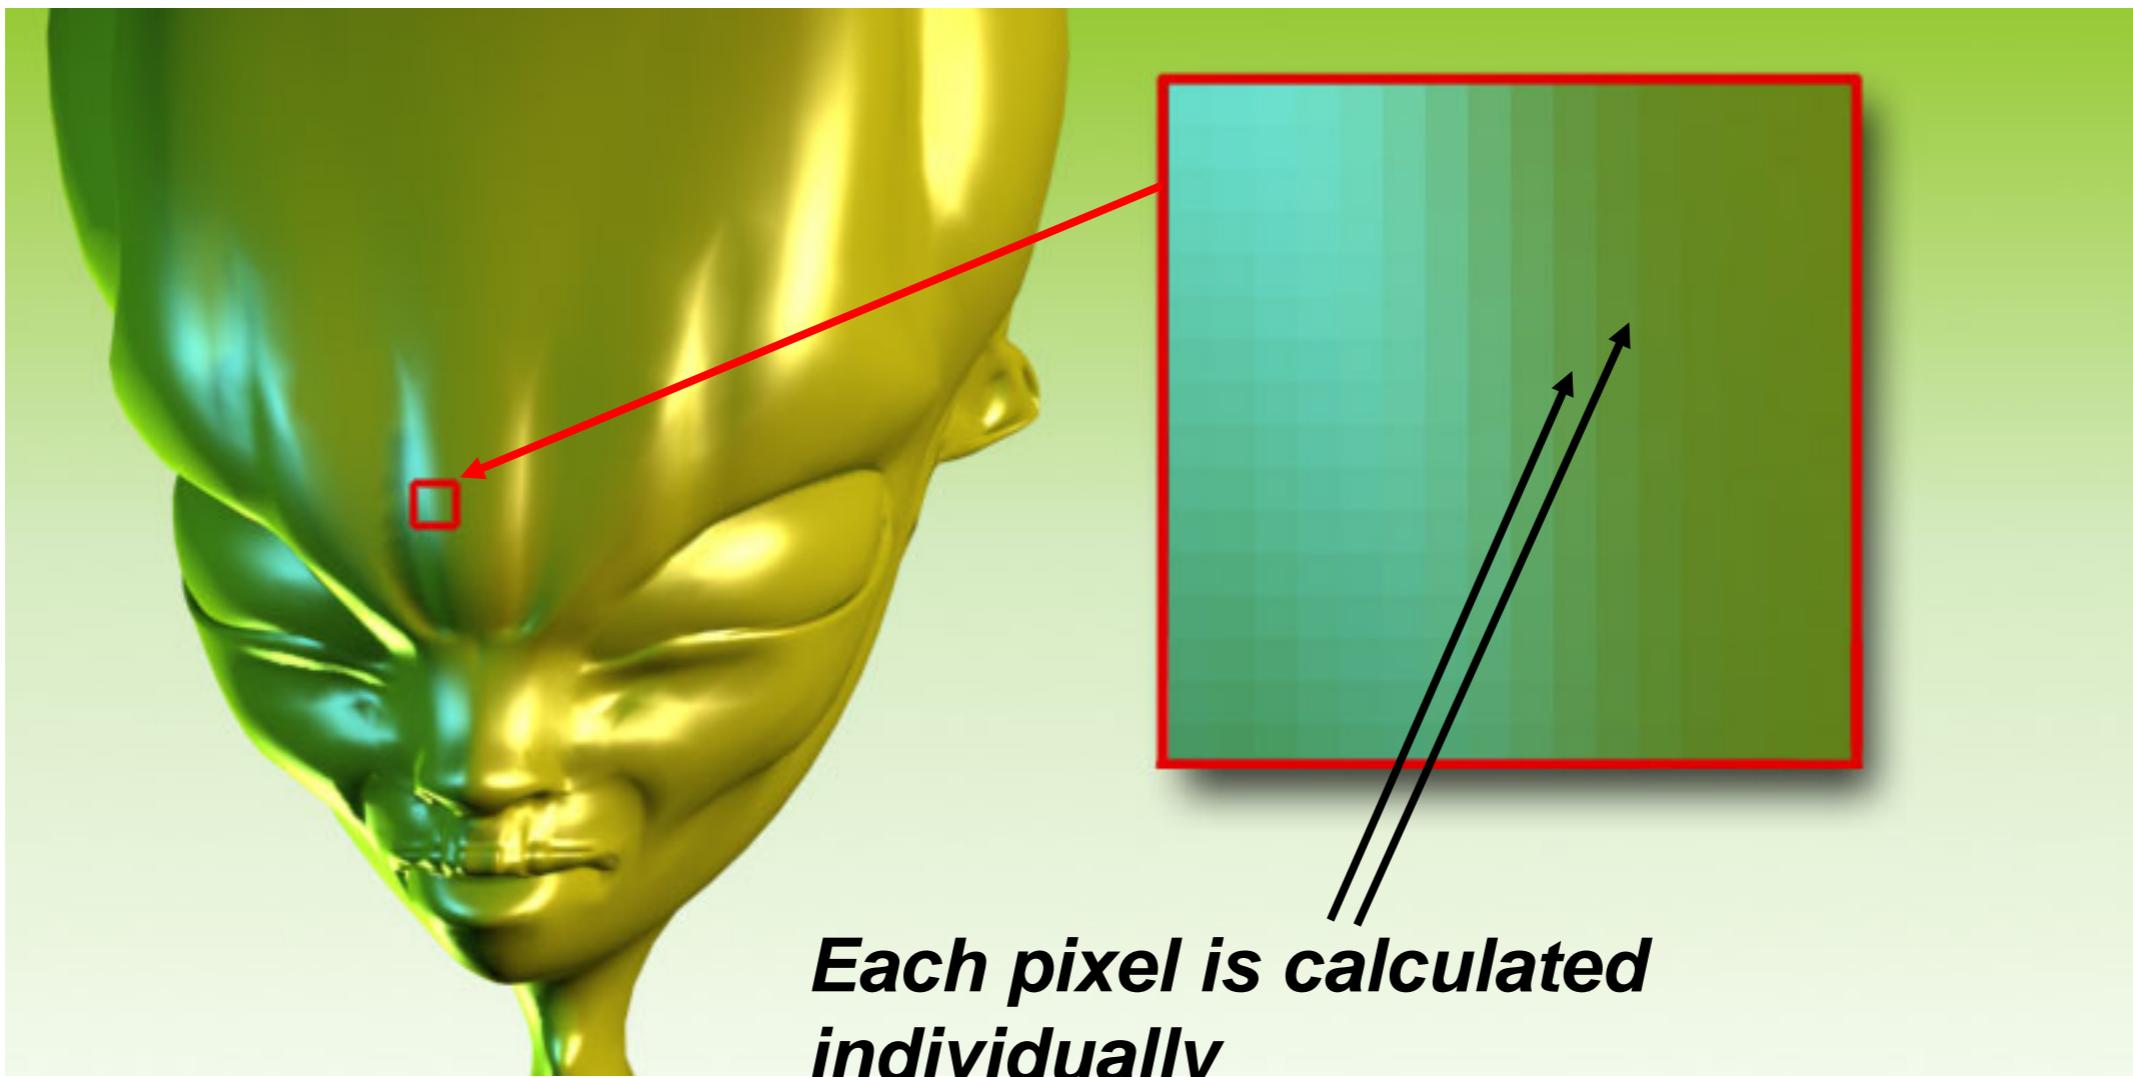
- **Prepare data for pixel shaders**
(e.g. texture coordinates)
 - Stored/computed at vertices
 - Linearly interpolated per pixel
- Modern graphics hardware provides tons of ***interpolants*** (32x4 floats)
- **Project vertices to the screen:**

```
float4 transform(float4 worldPos,  
                uniform float4x4 modelViewProjection) {  
    return mul(modelViewProjection, worldPos)  
}
```

Pixel Shaders

$f(\text{interpolants}) \longrightarrow \text{color, [depth]}$

purely functional (no side-effects)



Pixel shaders have little or no knowledge of neighboring pixels

Figure from NVIDIA

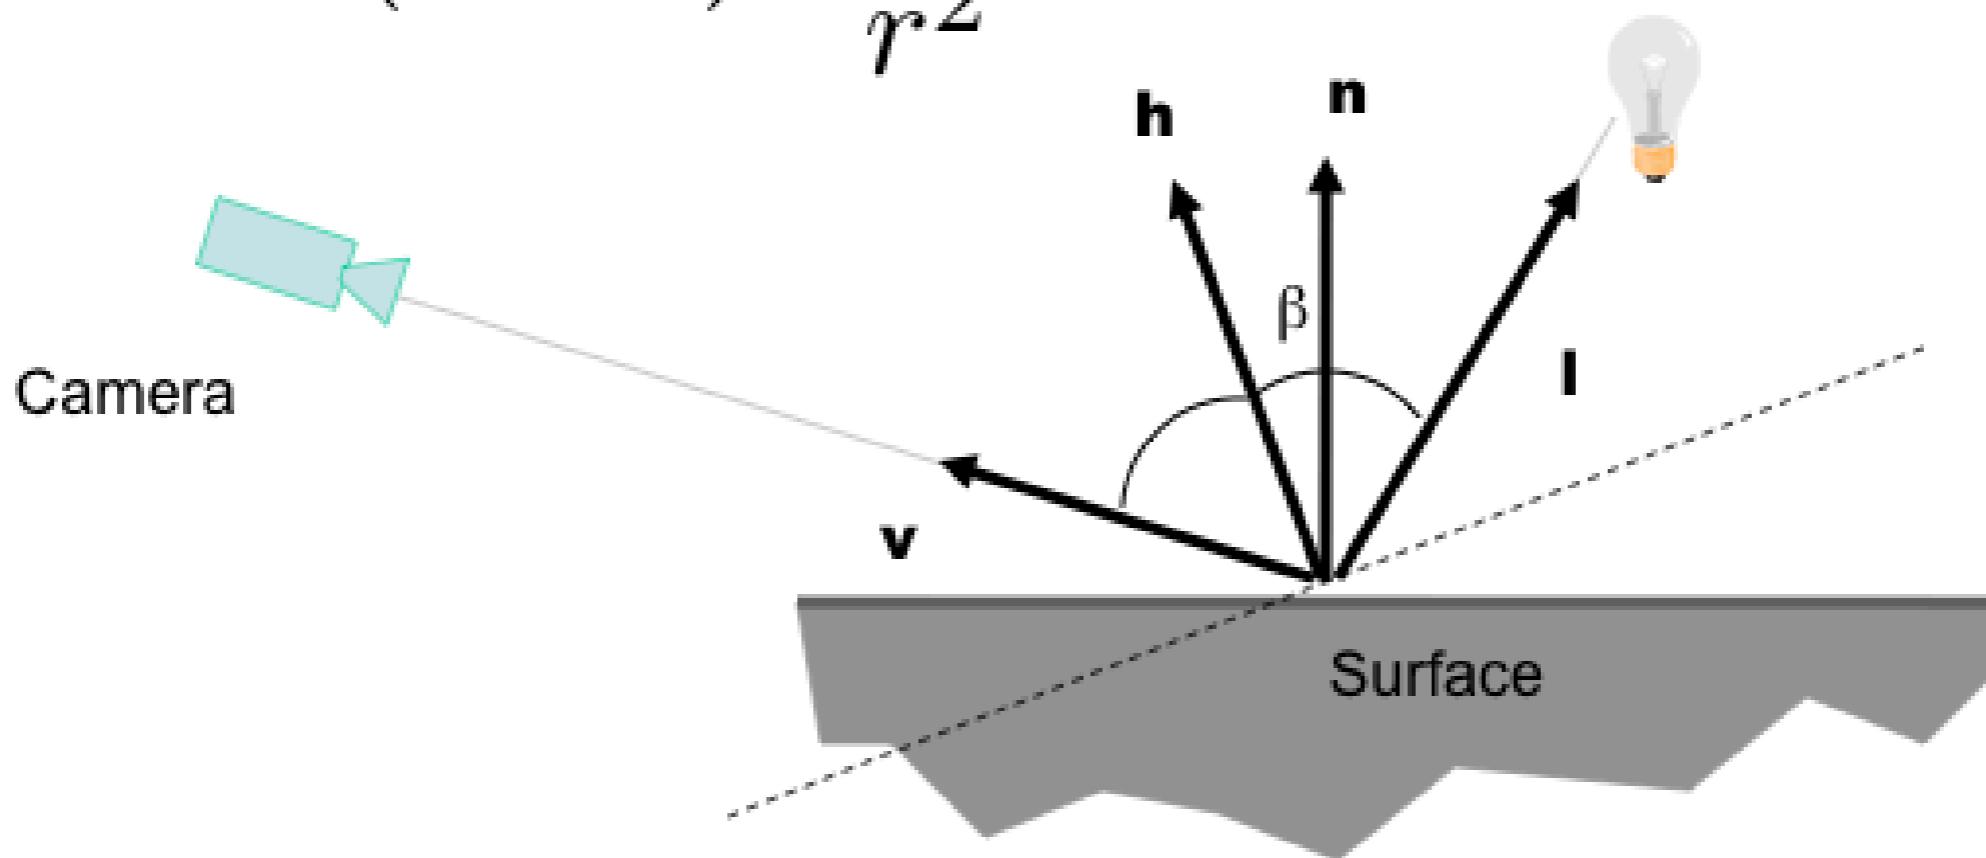
Pixel Shader: Blinn-Torrance Phong

Uses the “halfway vector” \mathbf{h} between \mathbf{l} and \mathbf{v} .

$$L_o = k_s \cos(\beta)^q \frac{L_i}{r^2}$$

$$\mathbf{h} = \frac{\mathbf{l} + \mathbf{v}}{\|\mathbf{l} + \mathbf{v}\|}$$

$$= k_s (\mathbf{n} \cdot \mathbf{h})^q \frac{L_i}{r^2}$$



Pixel Shader: Blinn-Torrance Phong

```
struct interpolants { float4 p, n, v }
```

```
struct light { float4 pos, float Li }
```

```
float4 phong(interpolants in, uniform light lgt,  
              uniform float q, uniform float Ks)
```

```
{
```

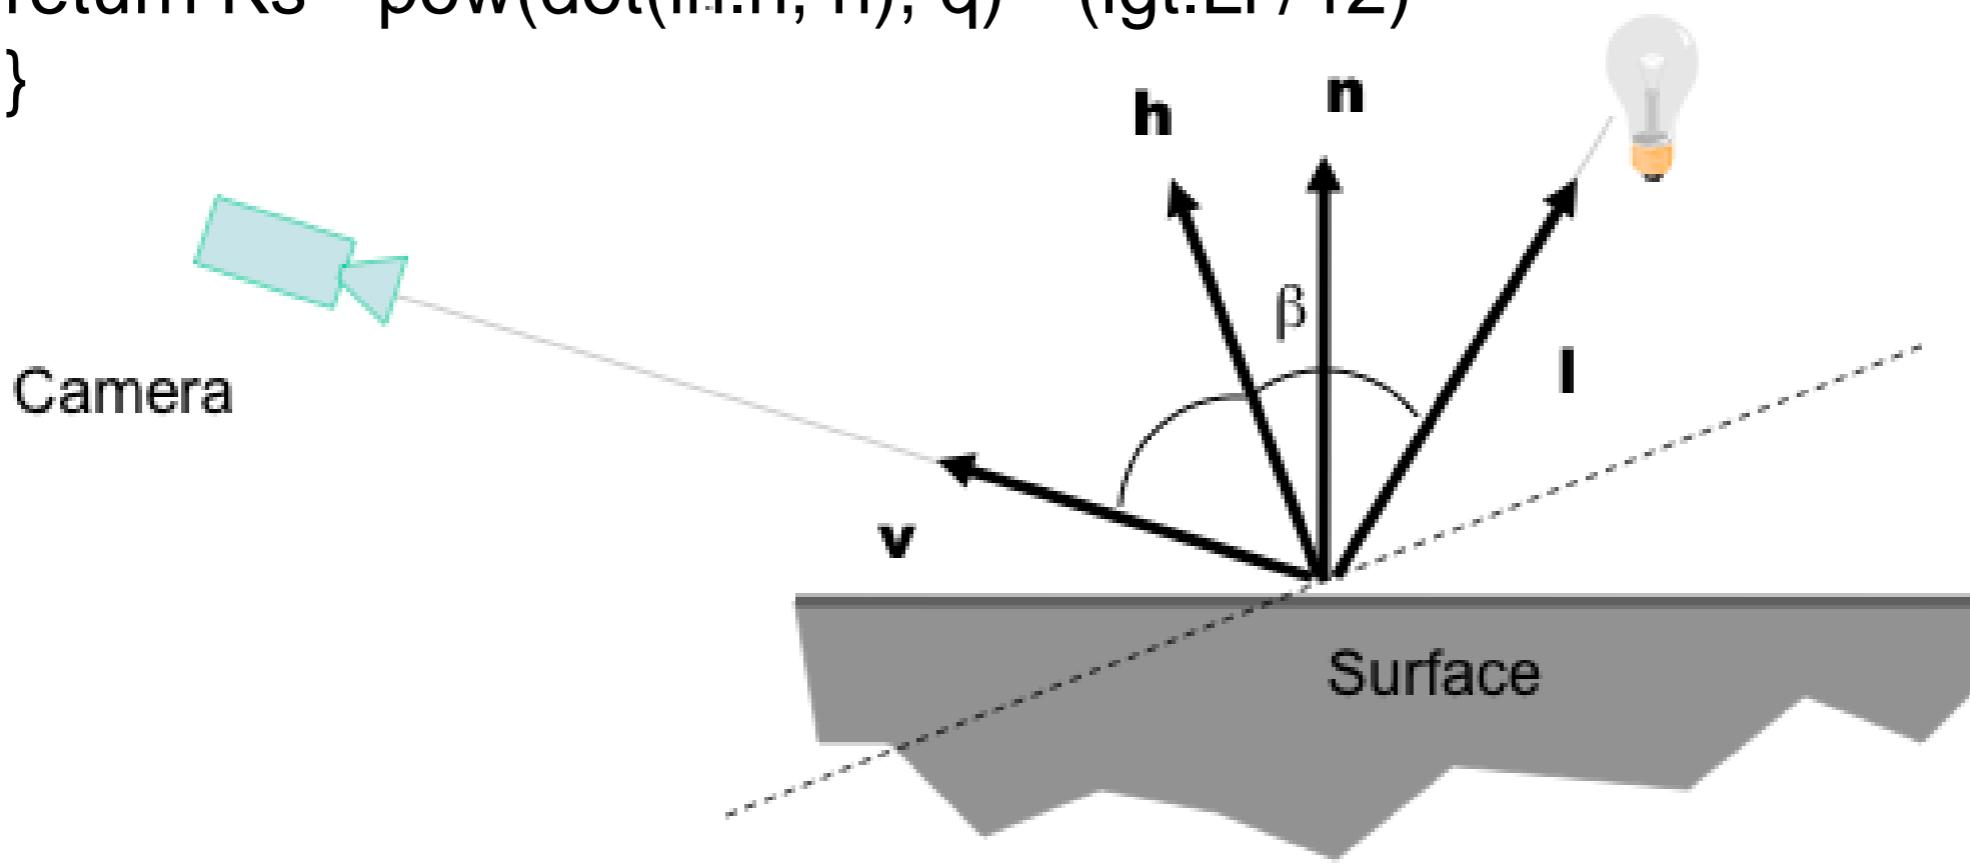
```
    float4 I = lgt.pos - in.p
```

```
    float r2 = length(I)*length(I)
```

```
    float4 h = normalize(I+in.v) // useful built-ins
```

```
    return Ks * pow(dot(in.n, h), q) * (lgt.Li / r2)
```

```
}
```



Brushed Metal

- Procedural texture
- Anisotropic lighting

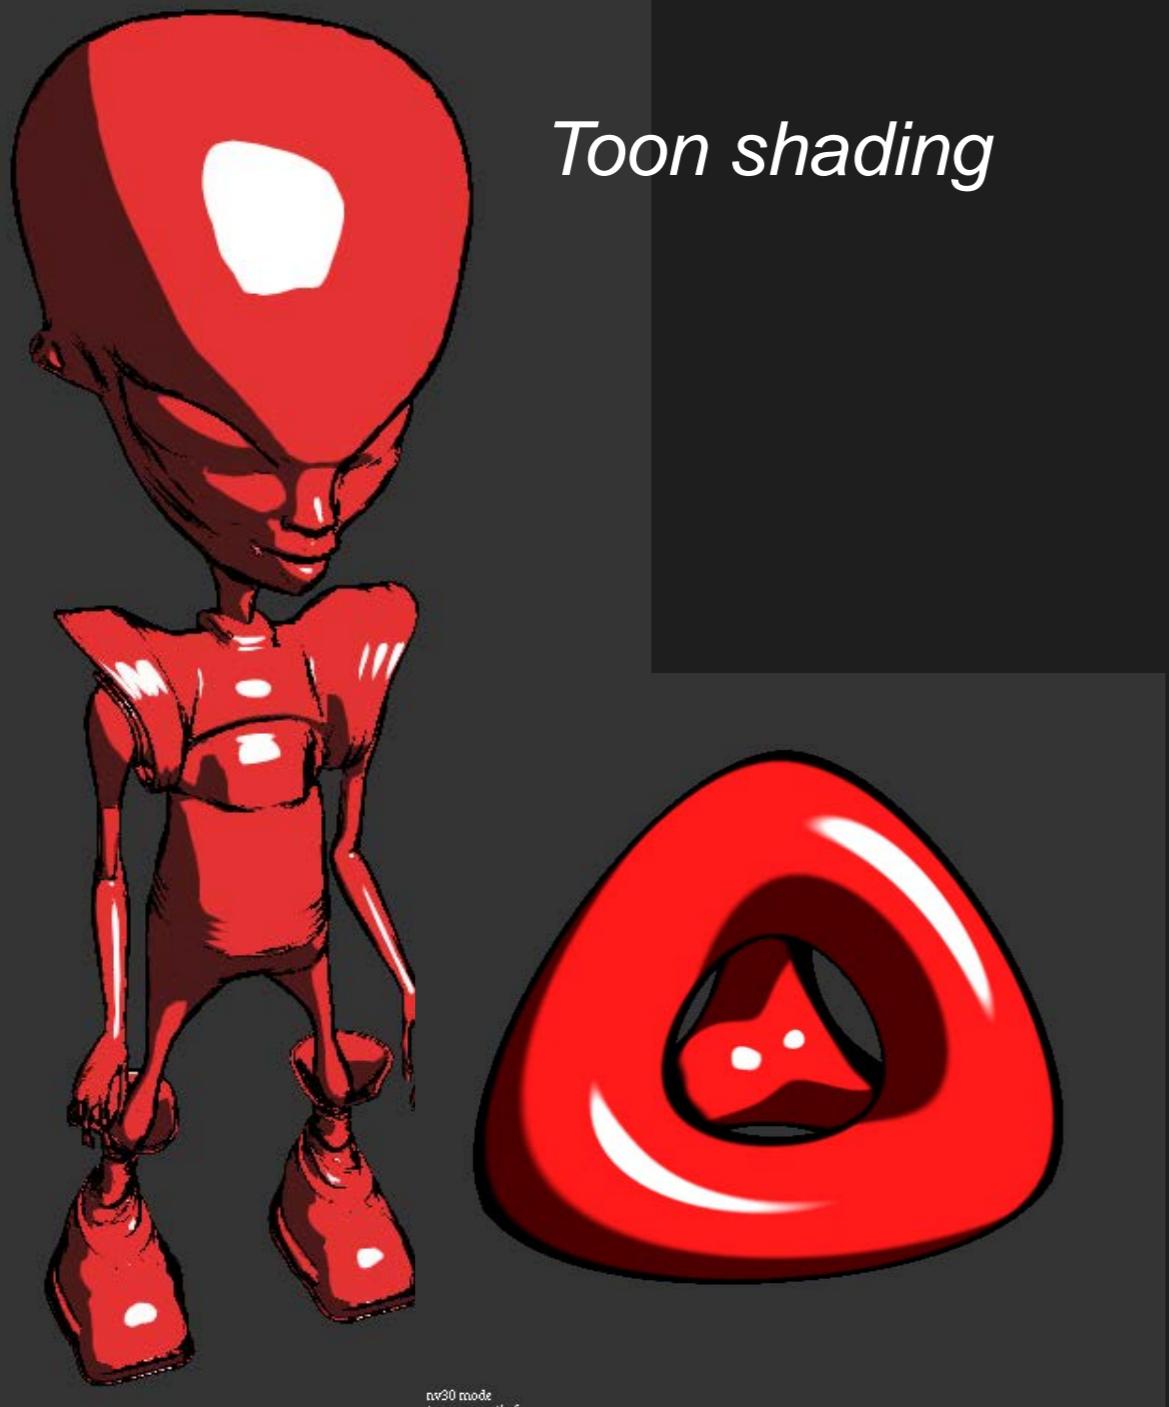


Melting Ice

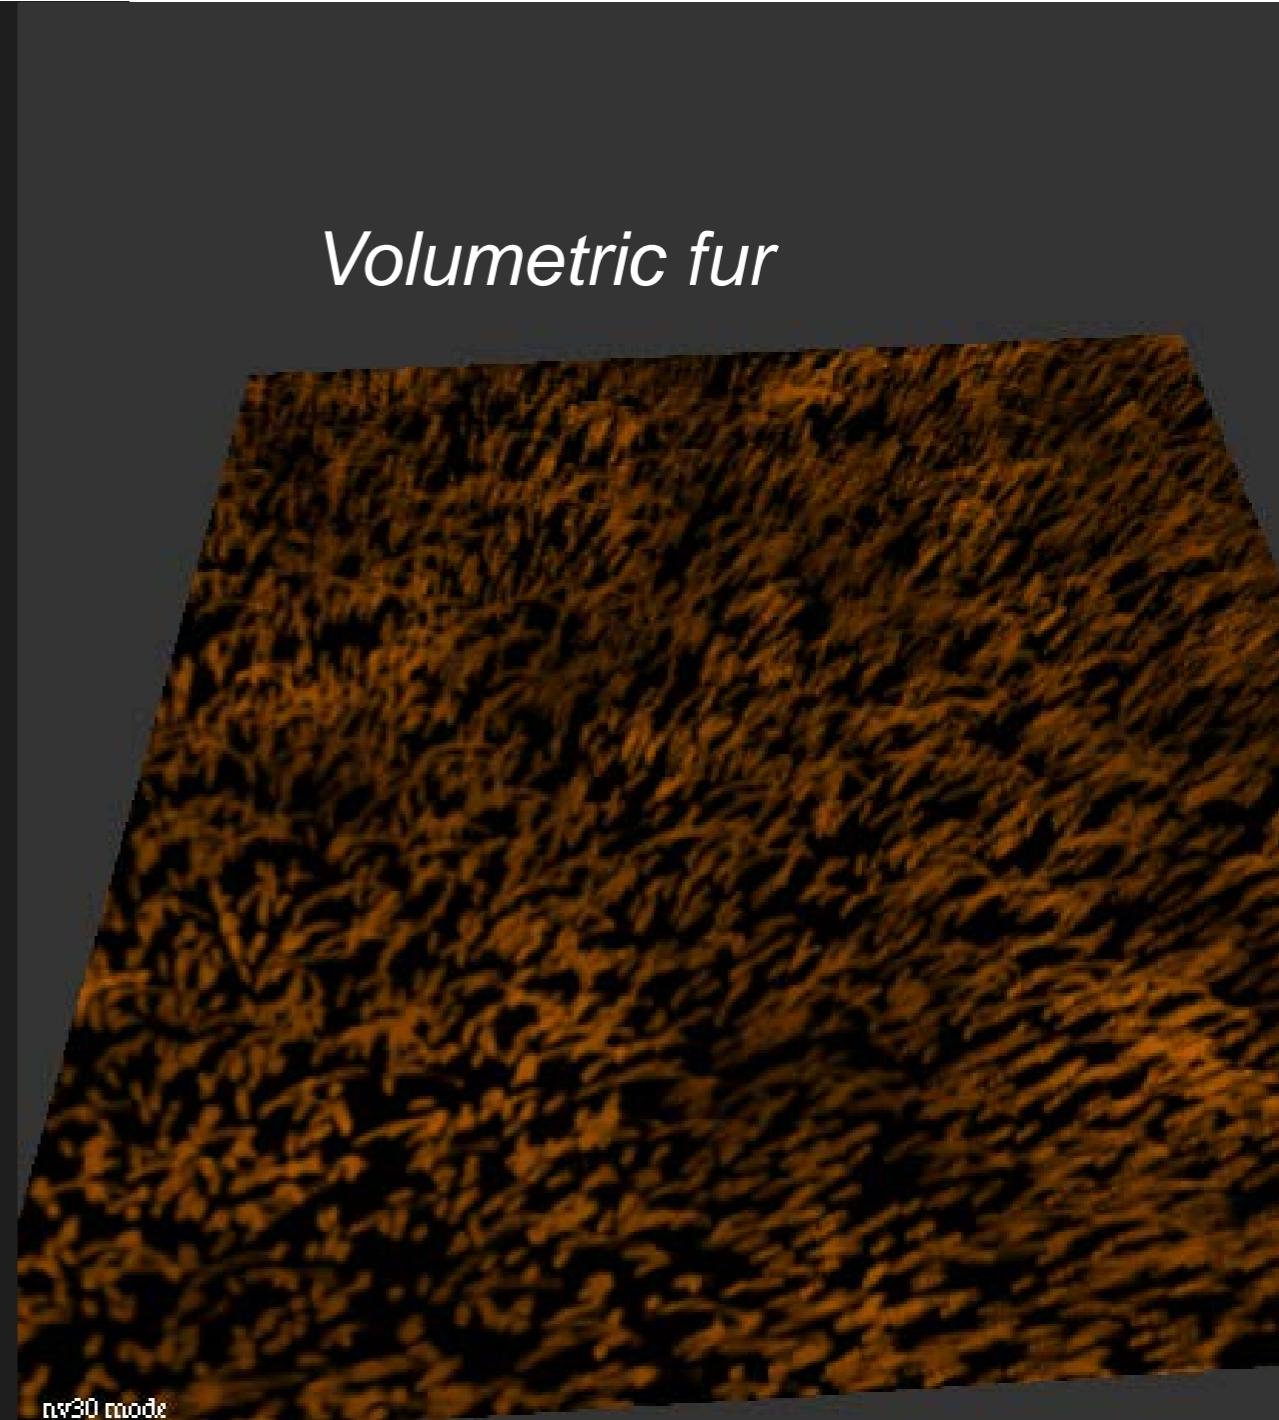


- Procedural, animating texture
- Bumped environment map

Toon & Fur



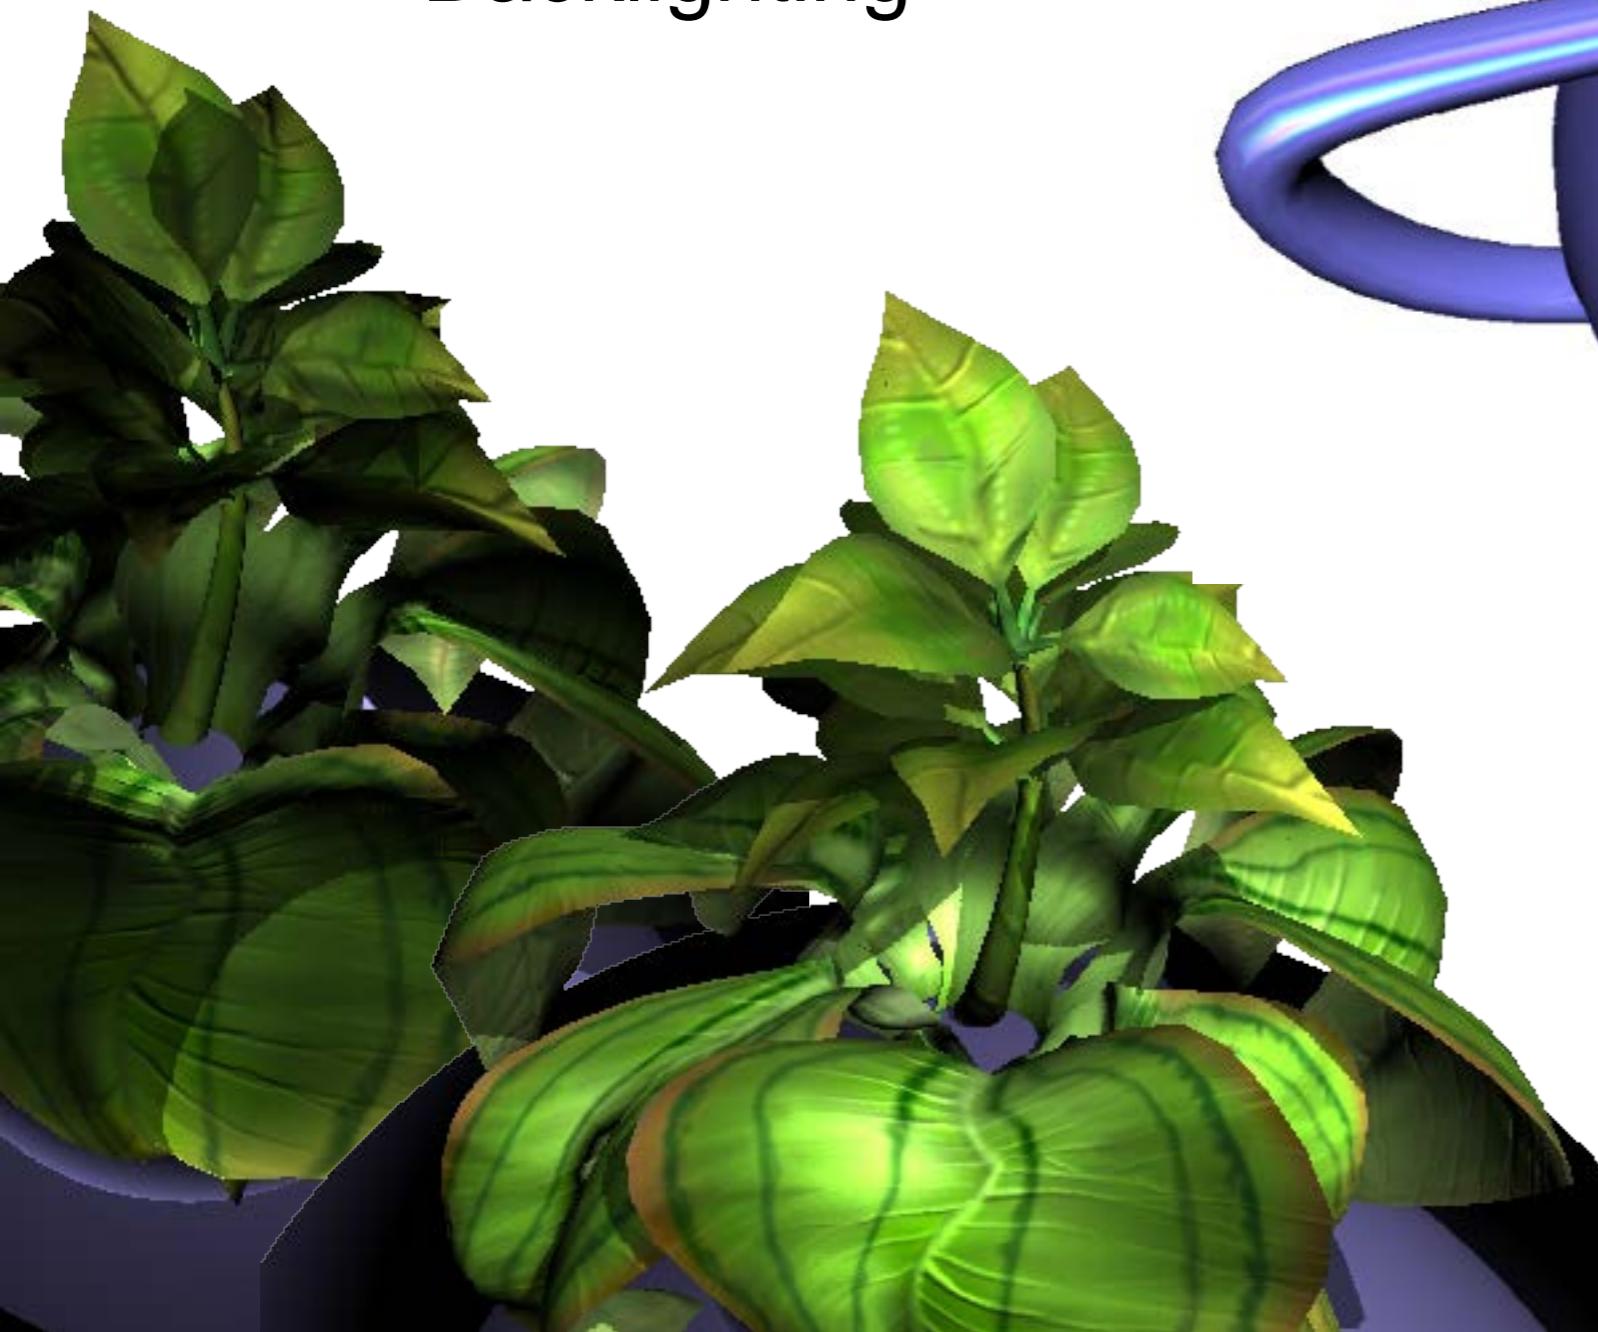
Toon shading



Volumetric fur

Vegetation & Thin Film

Translucence
Backlighting



Custom lighting
Simulates iridescence

Allows for amazing quality





Earth, 4.5×10^9 B.C.E.

Crysis, 2007

Rich scene appearance

- Vertex shader
 - Geometry (skinning, displacement)
 - Setup interpolants for pixel shaders
- Pixel shader
 - Visual appearance
 - Also used for image processing, GPGPU abuses
- Multi-pass
 - Render the scene or part of the geometry multiple times
 - Different viewpoints (e.g. shadow map, shadow volume)
 - More complex shading (e.g. pass per-light)
 - Image processing (e.g. HDR tone mapping, bloom)

Questions?

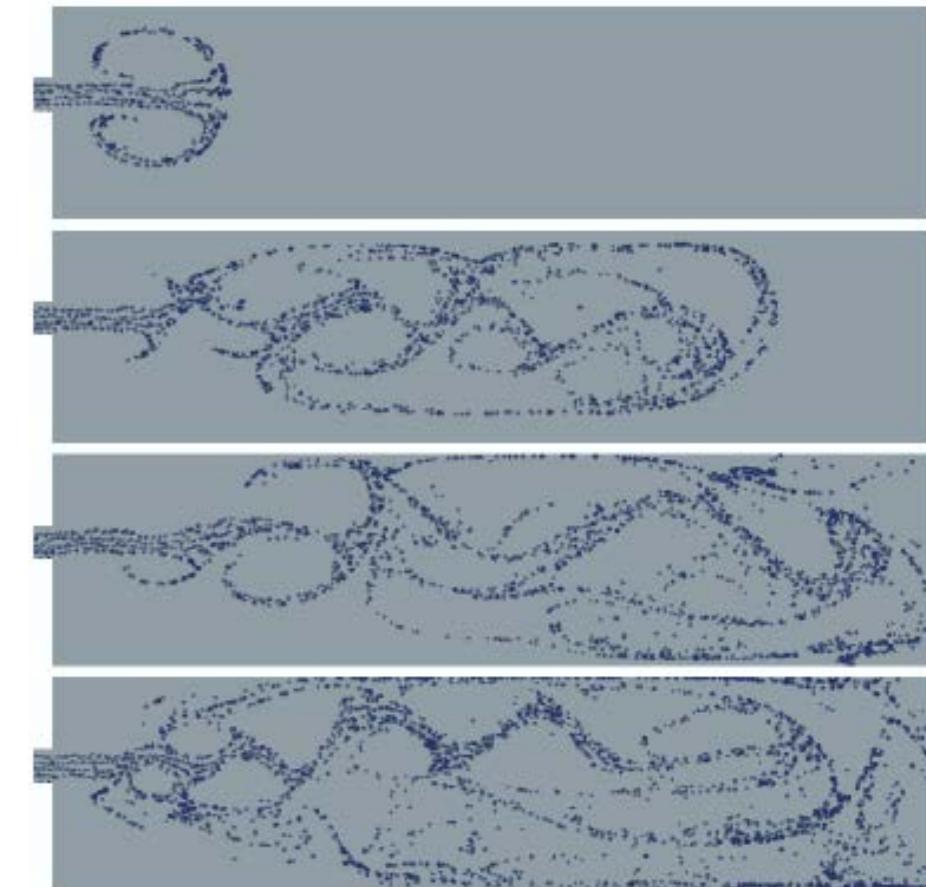
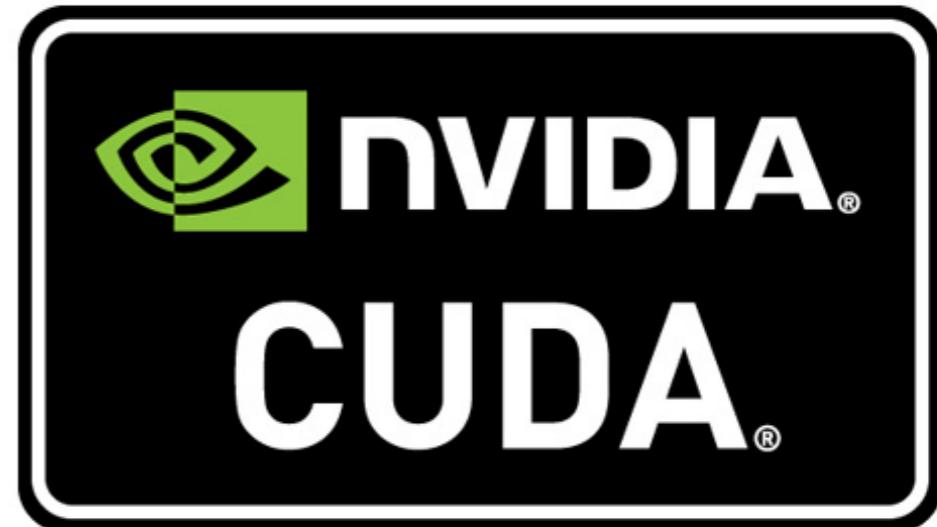


How to program shaders?

- Assembly code, or Higher-level language and compiler (e.g. HLSL, GLSL, Cg)
- Write separate **vertex** and **fragment** programs
- Send to the GPU via **graphics API** (Direct3D, OpenGL) (much like a texture or other piece of state)
- (Heavily) optimized at runtime by the driver's **JIT compiler**

General-purpose computation on GPUs

- Dense data-parallel processor (TeraFLOPS)
- Increasingly programmable (Code executed for each vertex or each pixel)
- C & C-like languages
- Use “shaders” for general-purpose computation
- Peak performance difficult to reach
- Parallelism = hard



Navier-Stokes on GPU [Bolz et al.]

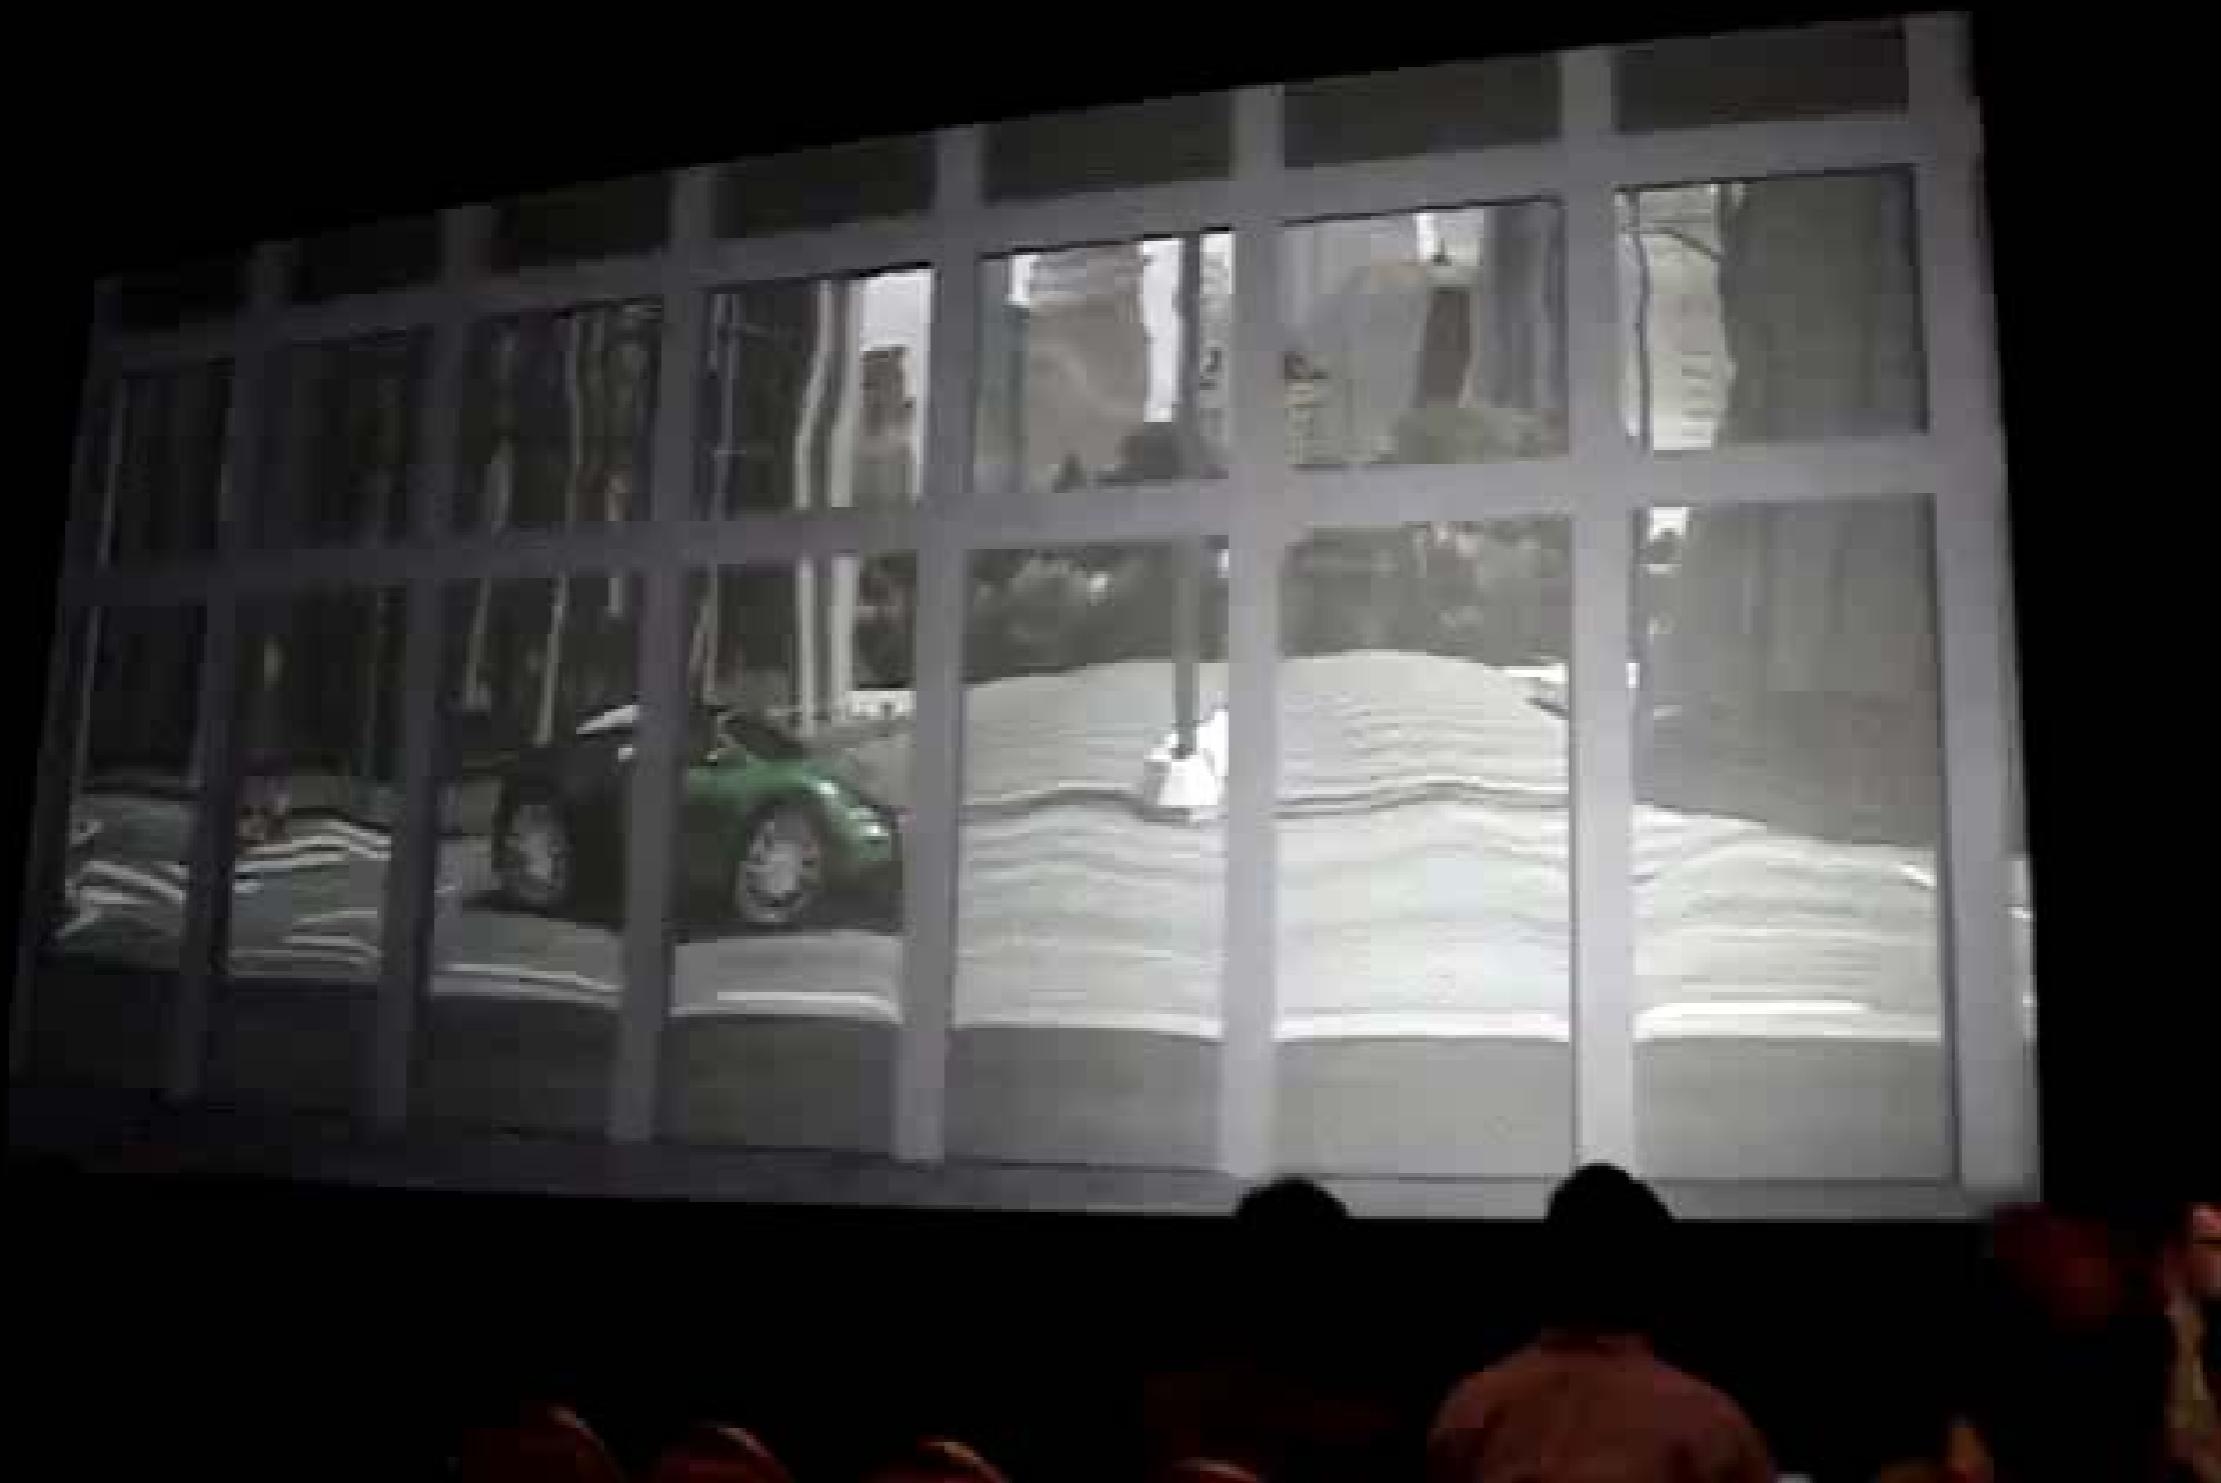
Real-Time Gravity Demo (NVIDIA)

Ray Tracing as GPGPU

- cool video!



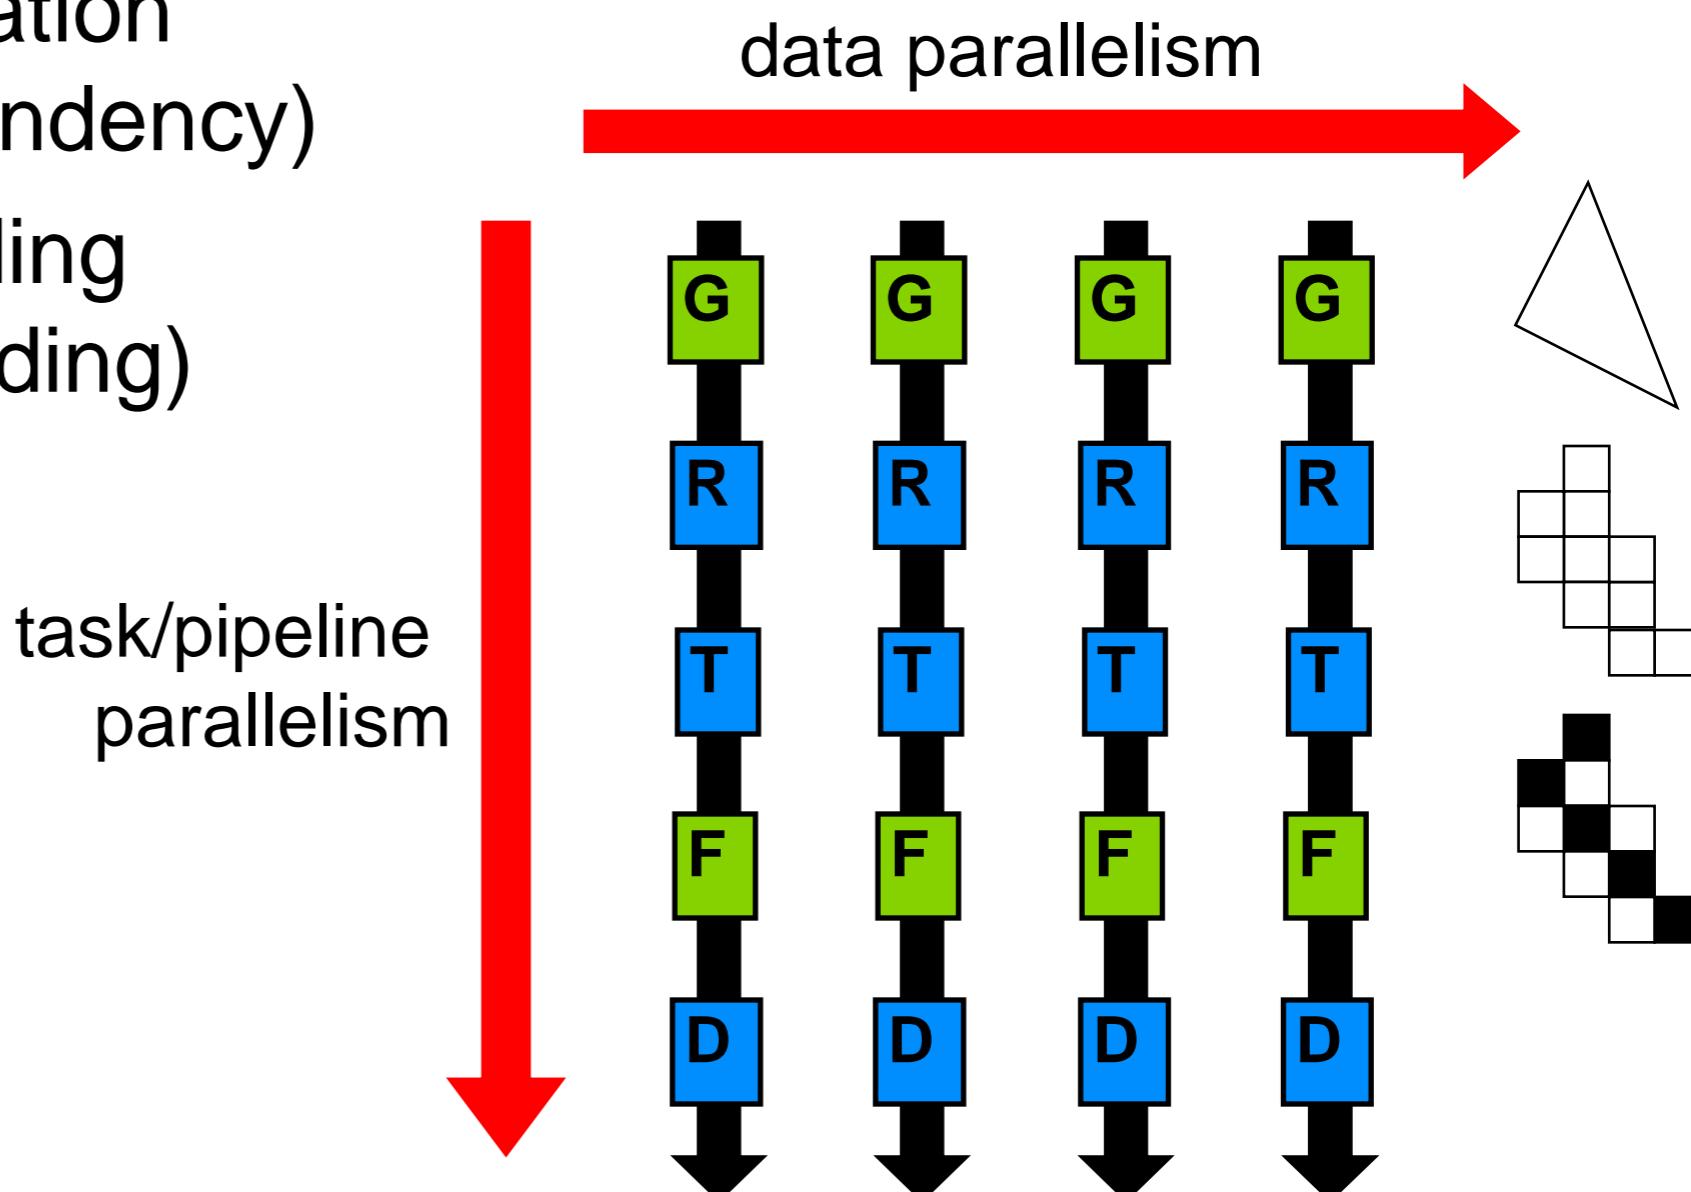
Ray Tracing as GPGPU



Questions?

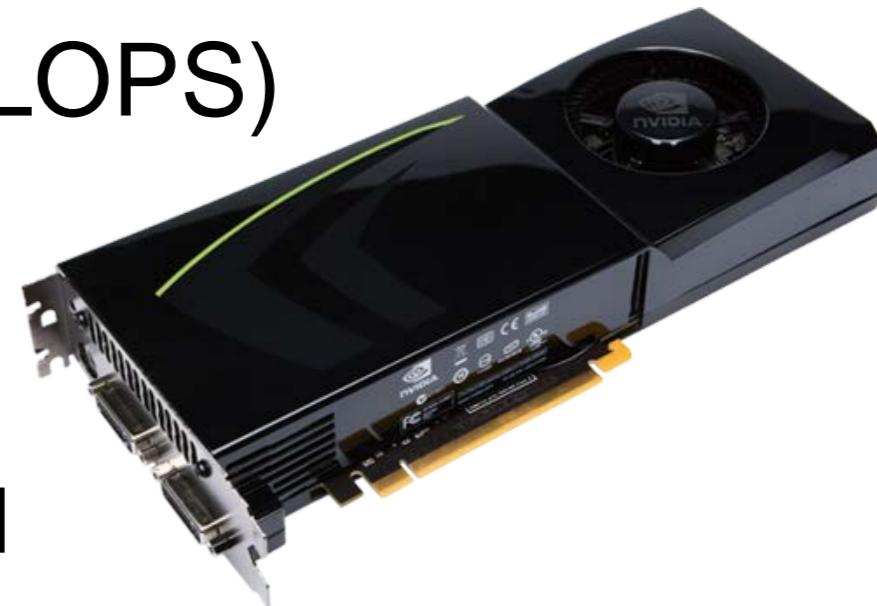
Graphics Hardware

- High performance through:
 - Parallelism
 - Specialization
 - Limited synchronization
(minimal data dependency)
 - Efficient latency hiding
(pre-fetching, threading)



Modern Graphics Hardware

- 100s of parallel shader pipelines (TFLOPS)
- Deep pipeline (~1k stages)
- Hierarchical tiling of screen (~4x4)
 - Early Z rejection if entire tile is occluded
- Pixels processed as quads (2x2 pixel neighborhoods)
 - Allows *derivatives* in pixel shader (for filtering)
- Massive multithreading to hide texture latency
 - And smart memory layout
 - Throughput (bandwidth)-optimized memory controllers (at the expense of latency)



Why is graphics hardware so fast?

GPU: throughput

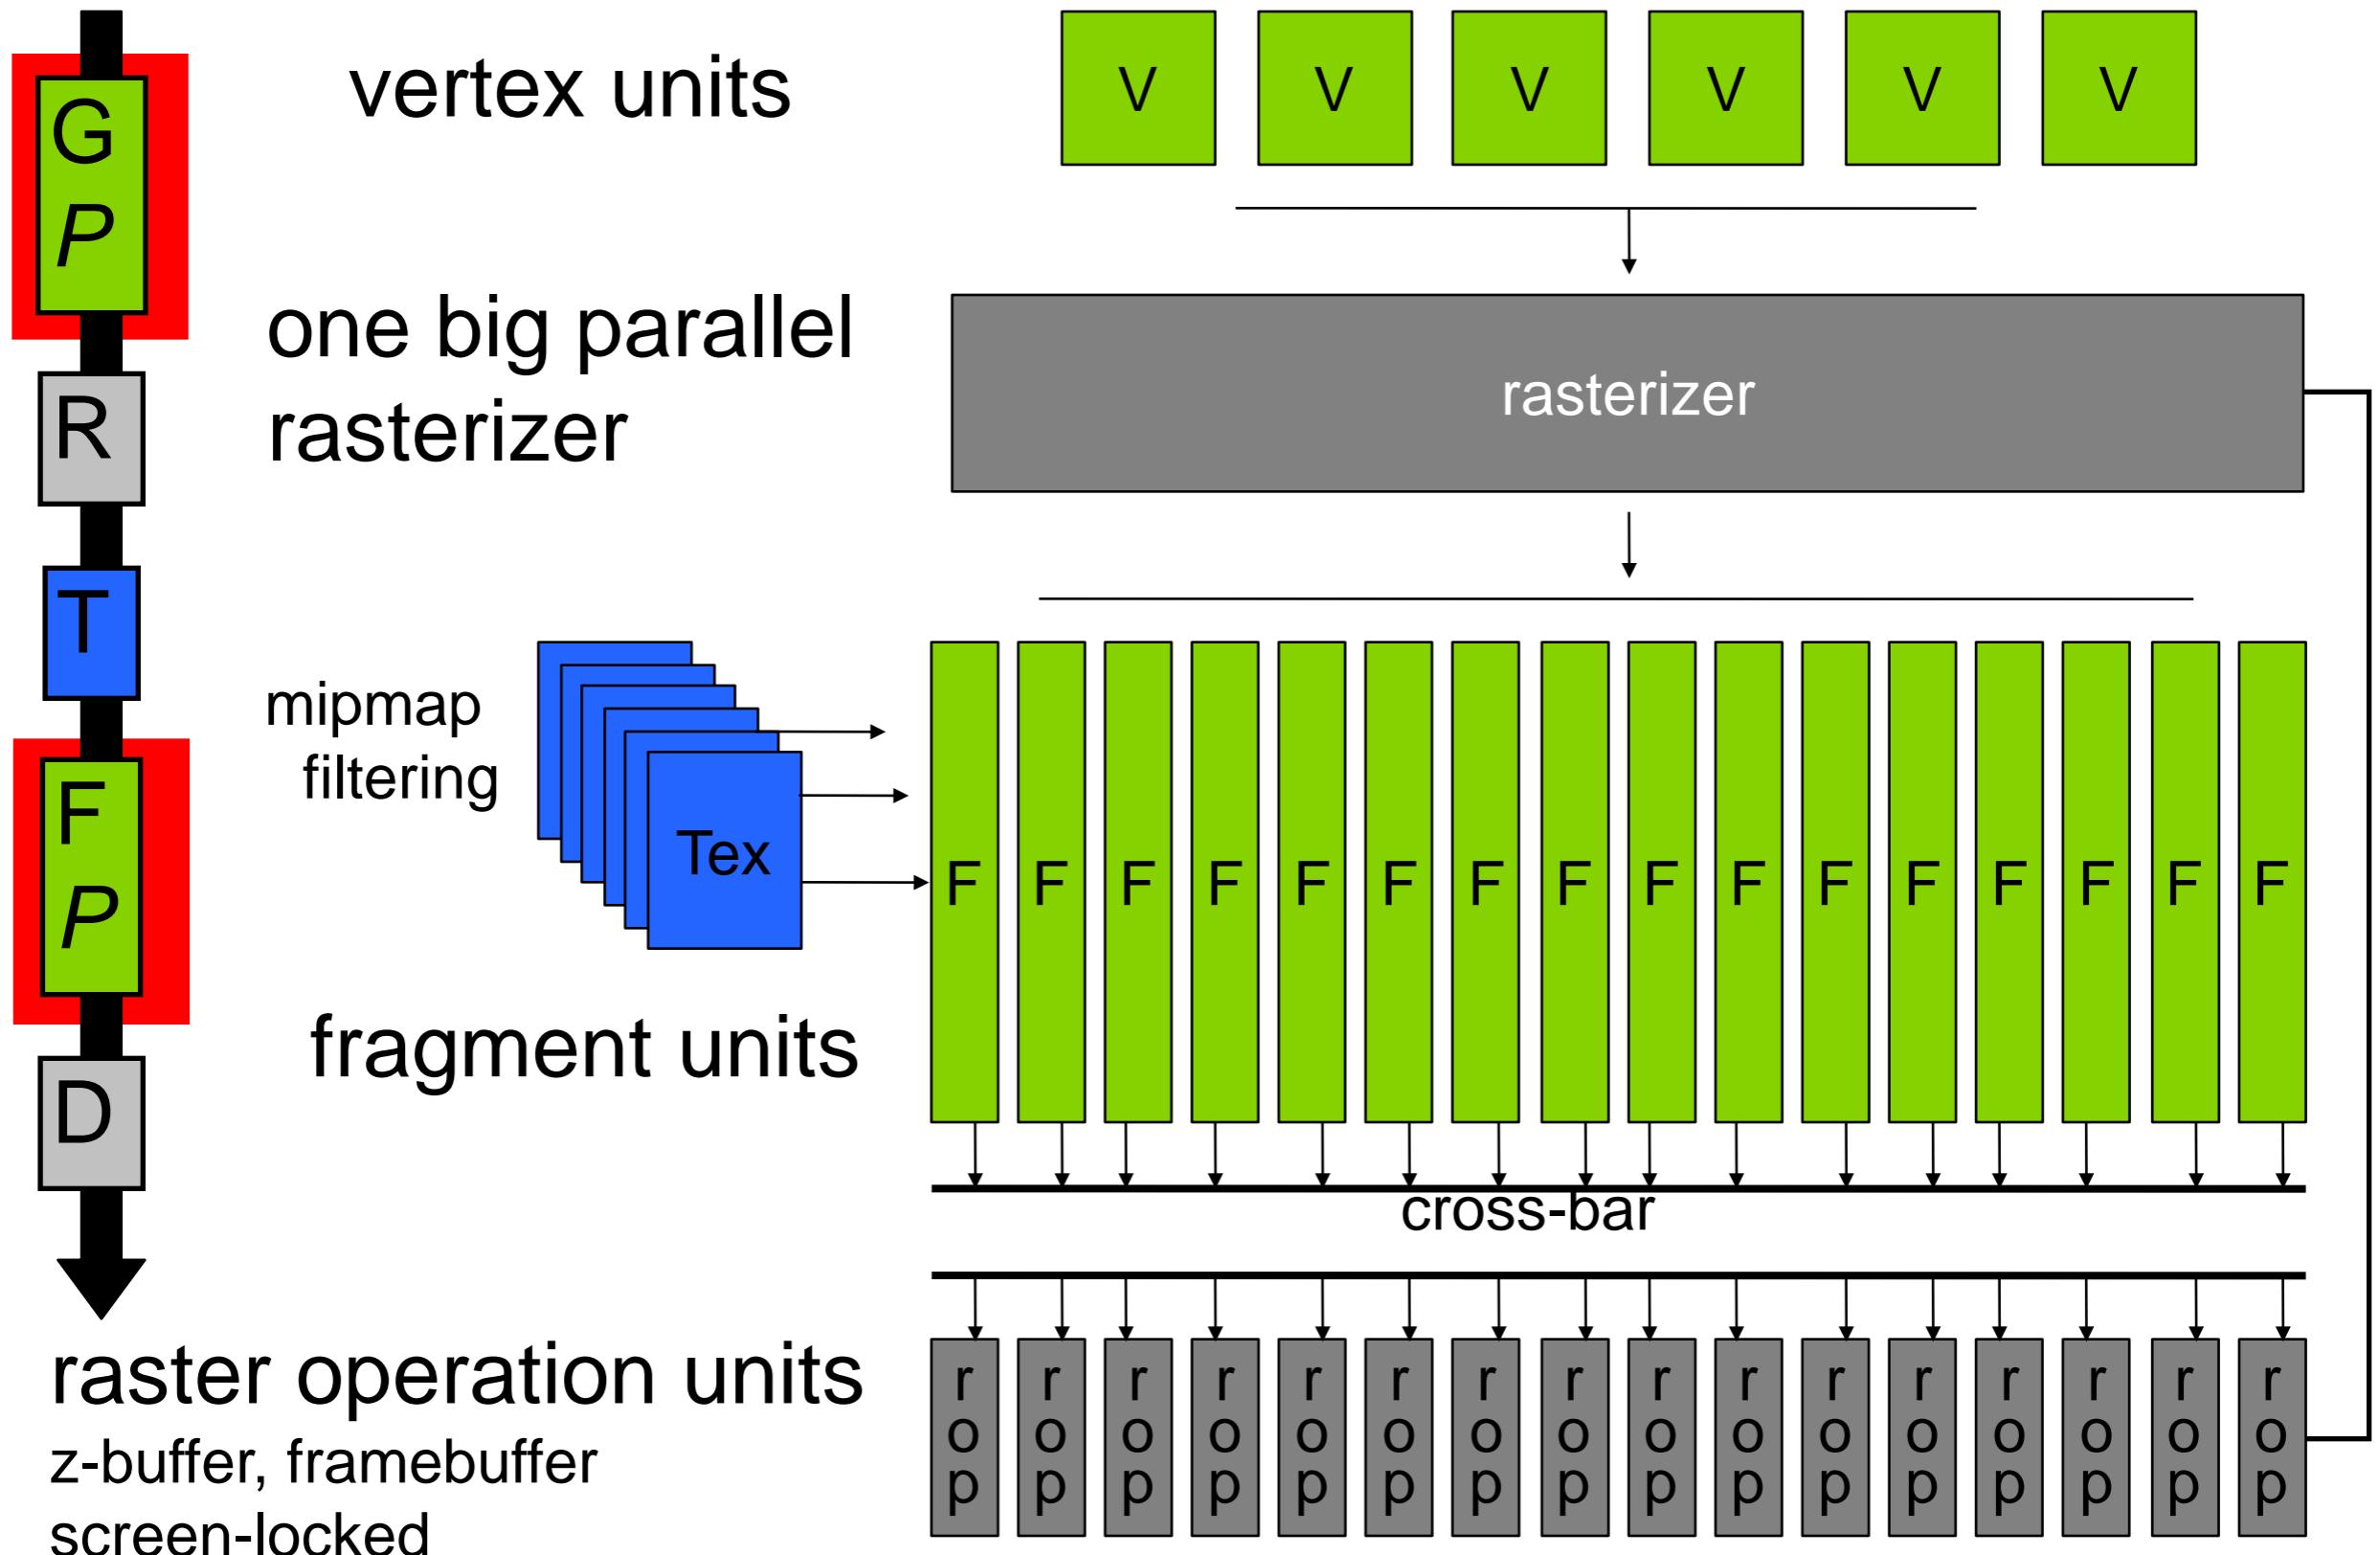
- Massive, explicit **parallelism**
- More transistors do **compute**, less cache, bookkeeping
- **Specialization** (rasterizer, texture filtering)
- **High** arithmetic intensity (math ops » memory ops)
- Deep pipeline, latency hiding, prefetching
- Little data dependence
- Controlled memory access patterns

CPU: latency

- Predominantly **scalar**
- Most transistors are **bookkeeping, cache**
- Highly **general-purpose** (x86 ISA)
- **Low** arithmetic intensity (math ops ≈ memory ops)
- Short pipeline, latency intolerance, out-of-order exec
- Arbitrary data dependence
- Arbitrary, dynamic memory access patterns

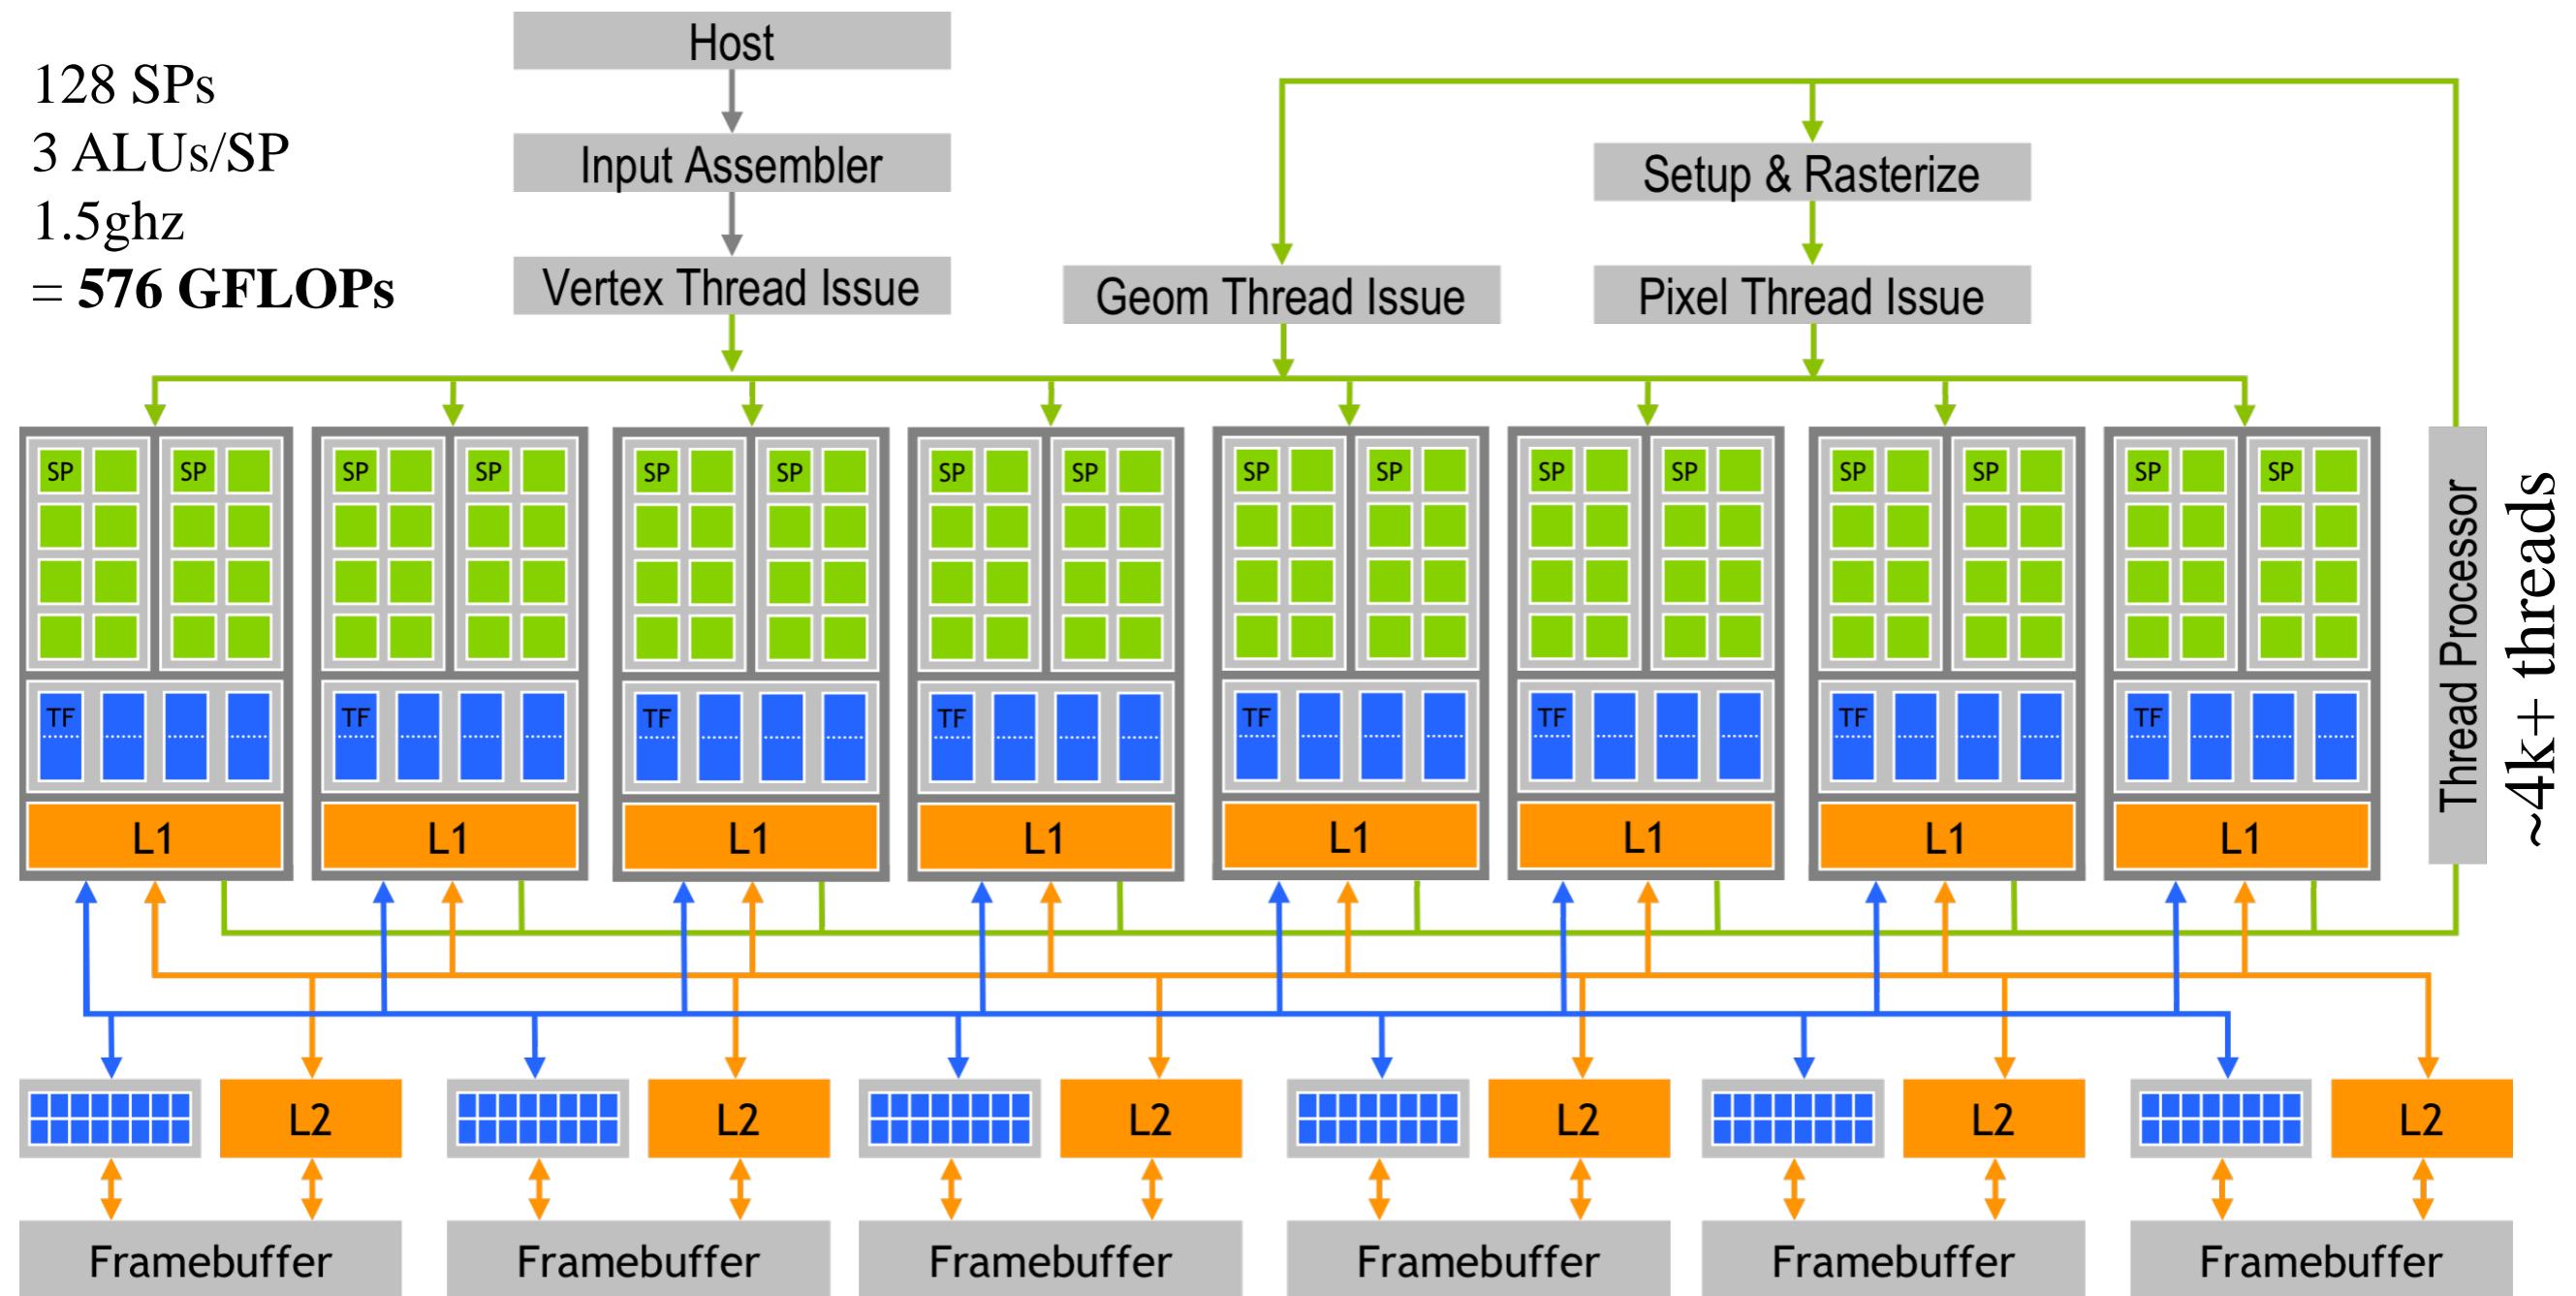
Questions?

Traditional hardware architecture



Unified shading architecture

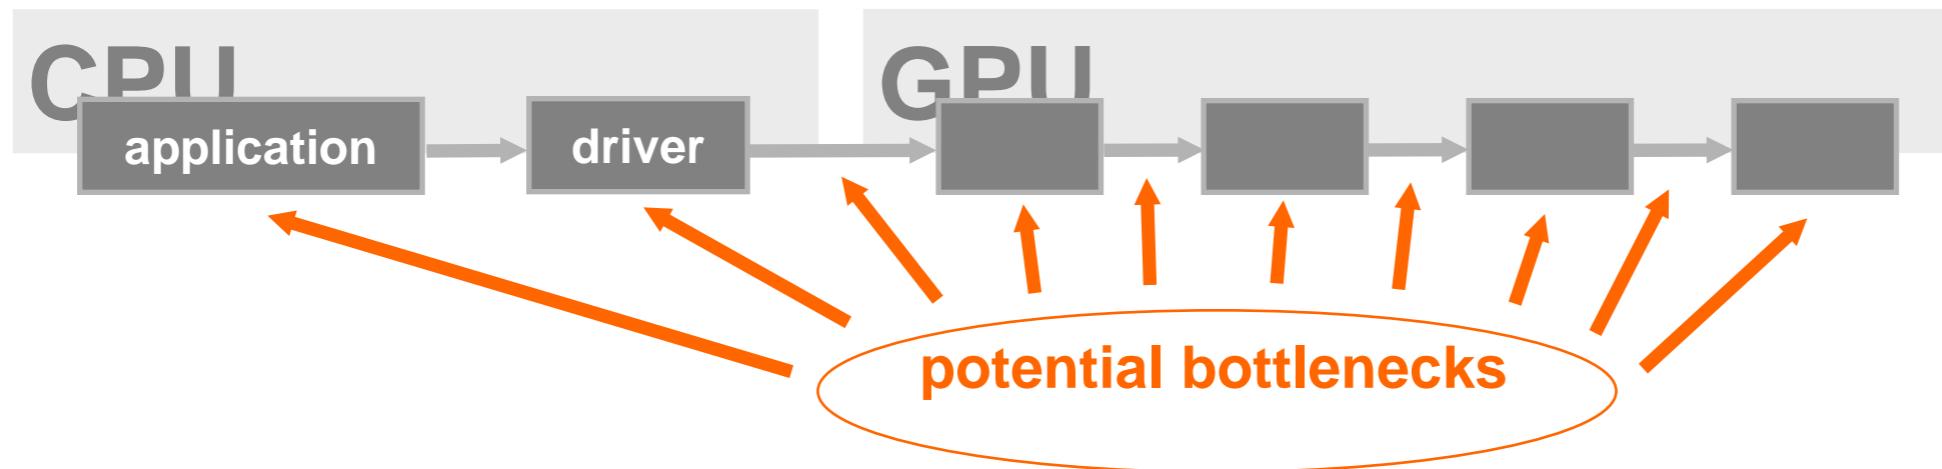
(NVIDIA G80)



85 GB/sec

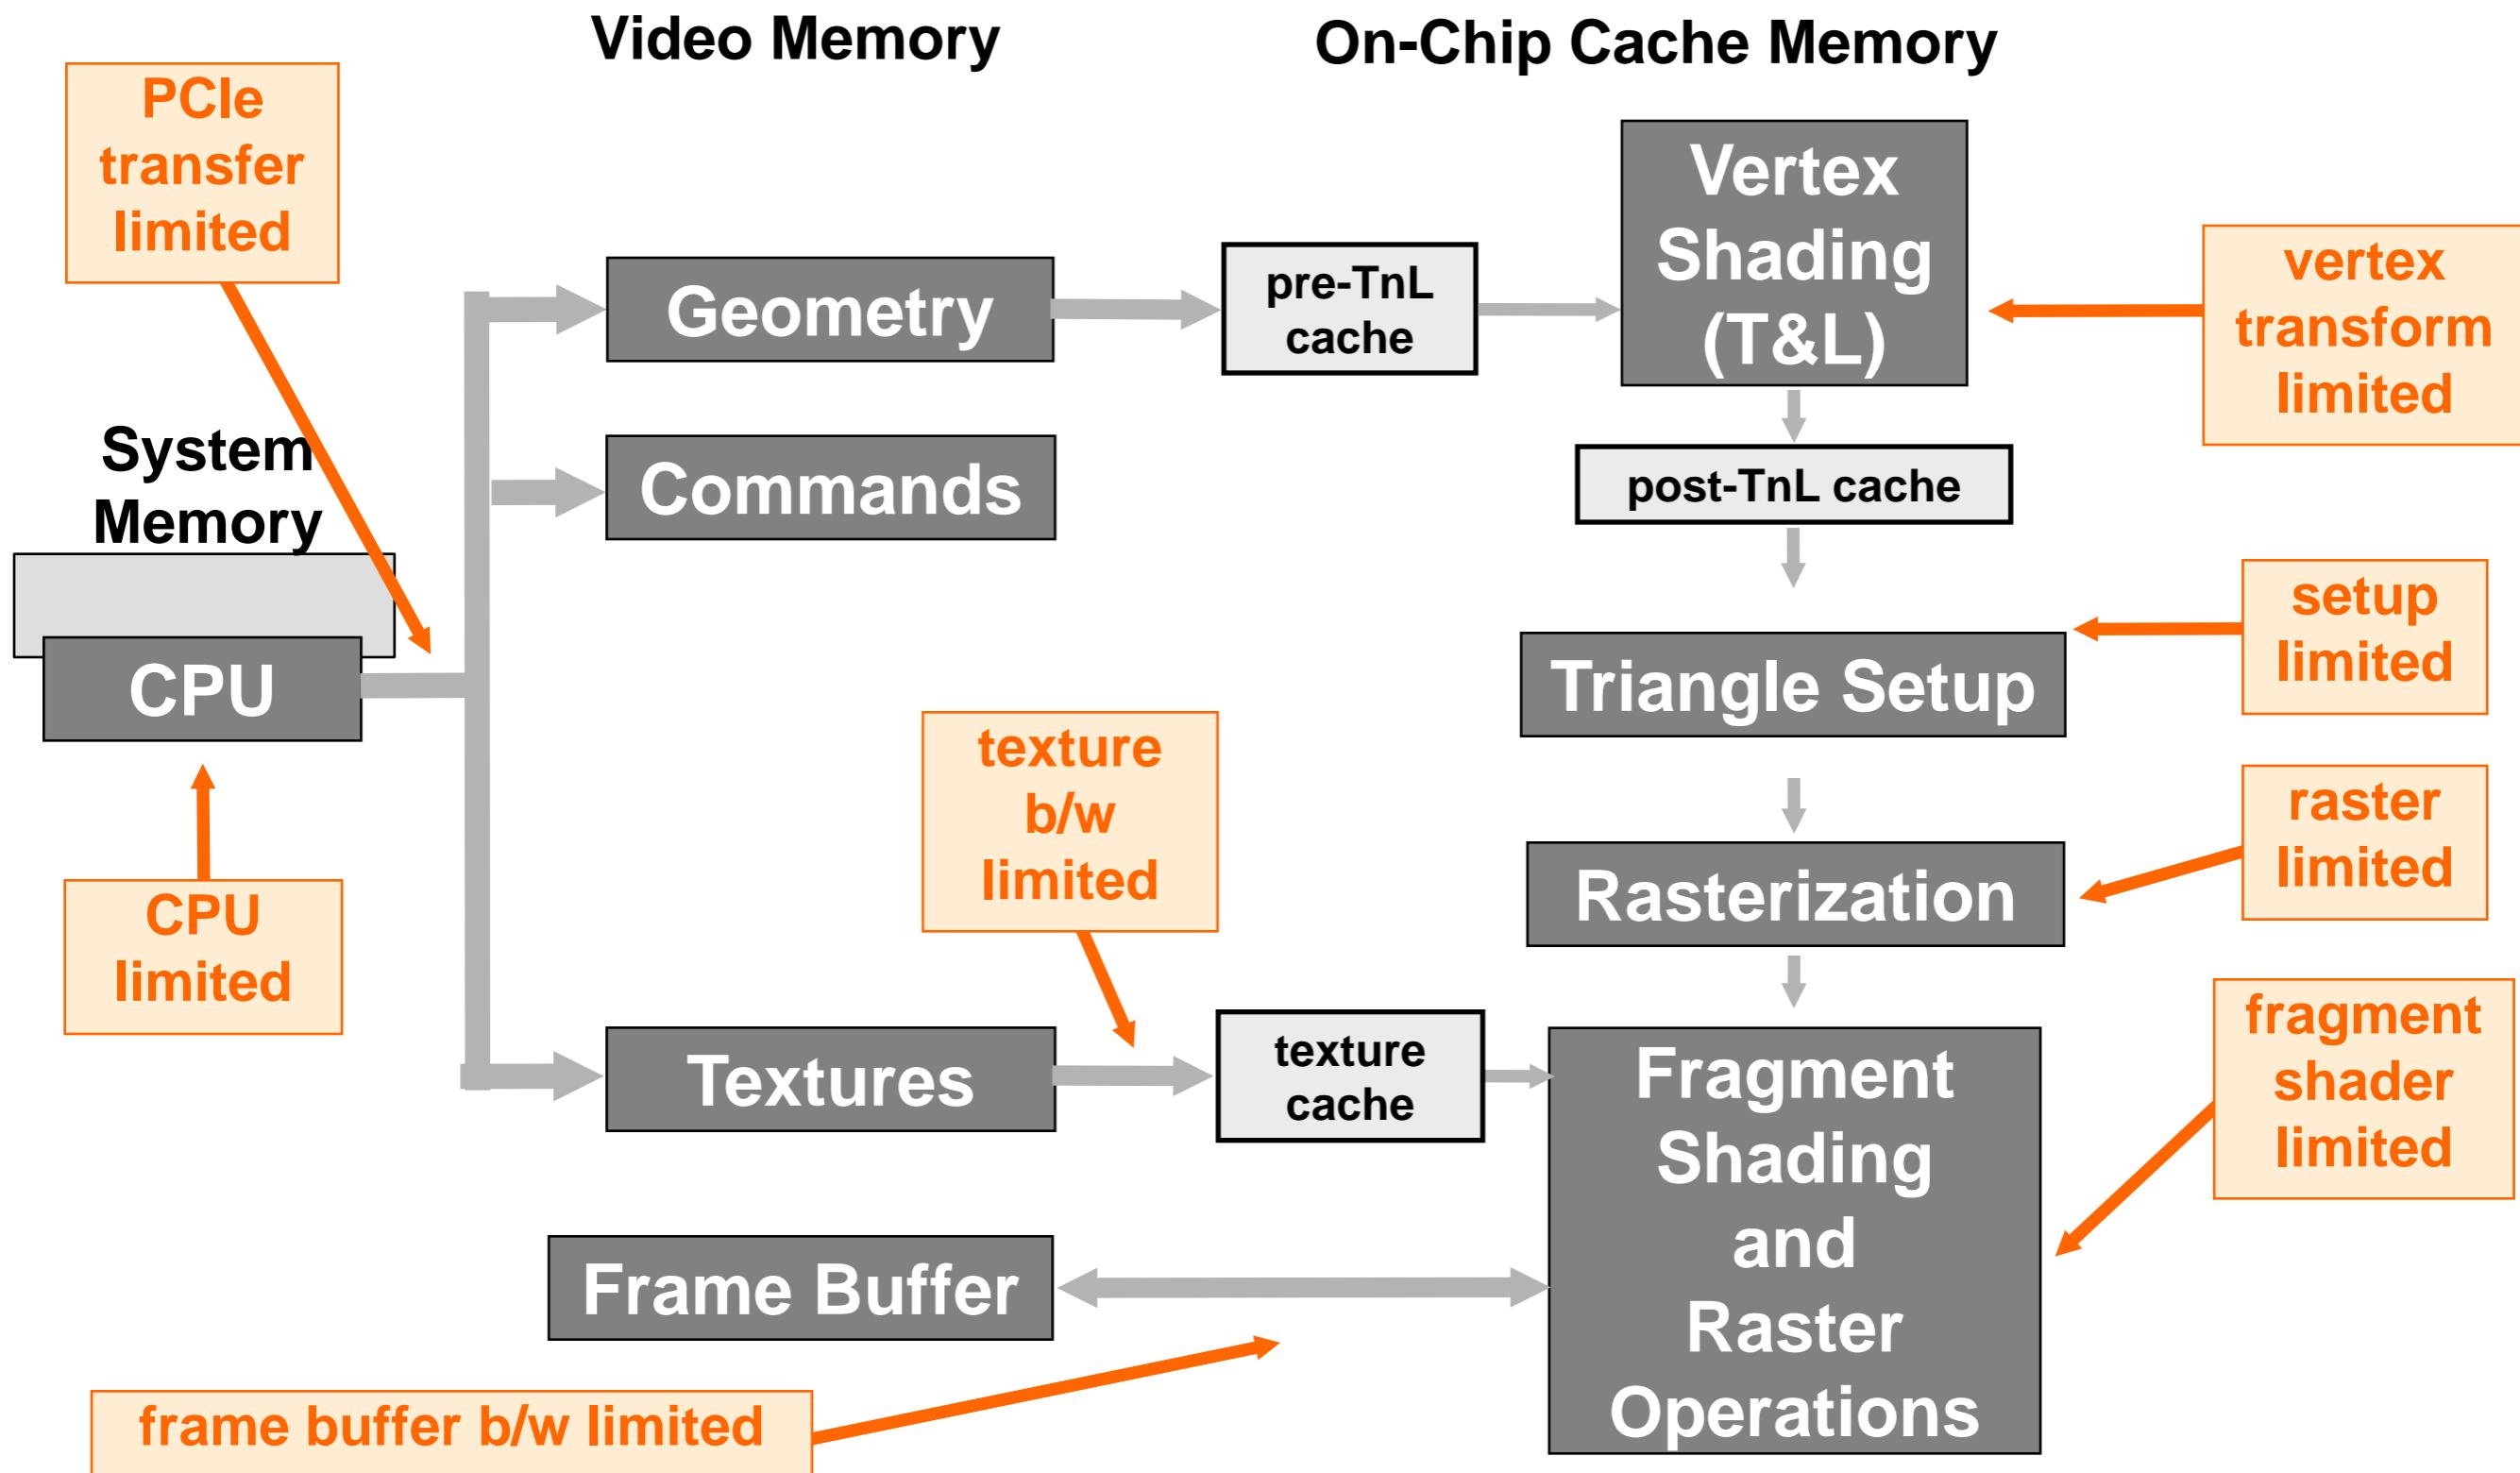
Questions?

Bottlenecks?



- The bottleneck determines overall throughput
- In general, the bottleneck varies over the course of an application and even over a frame
- Getting good performance is all about finding and eliminating bottlenecks in the pipeline

Potential Bottlenecks



Rendering pipeline bottlenecks

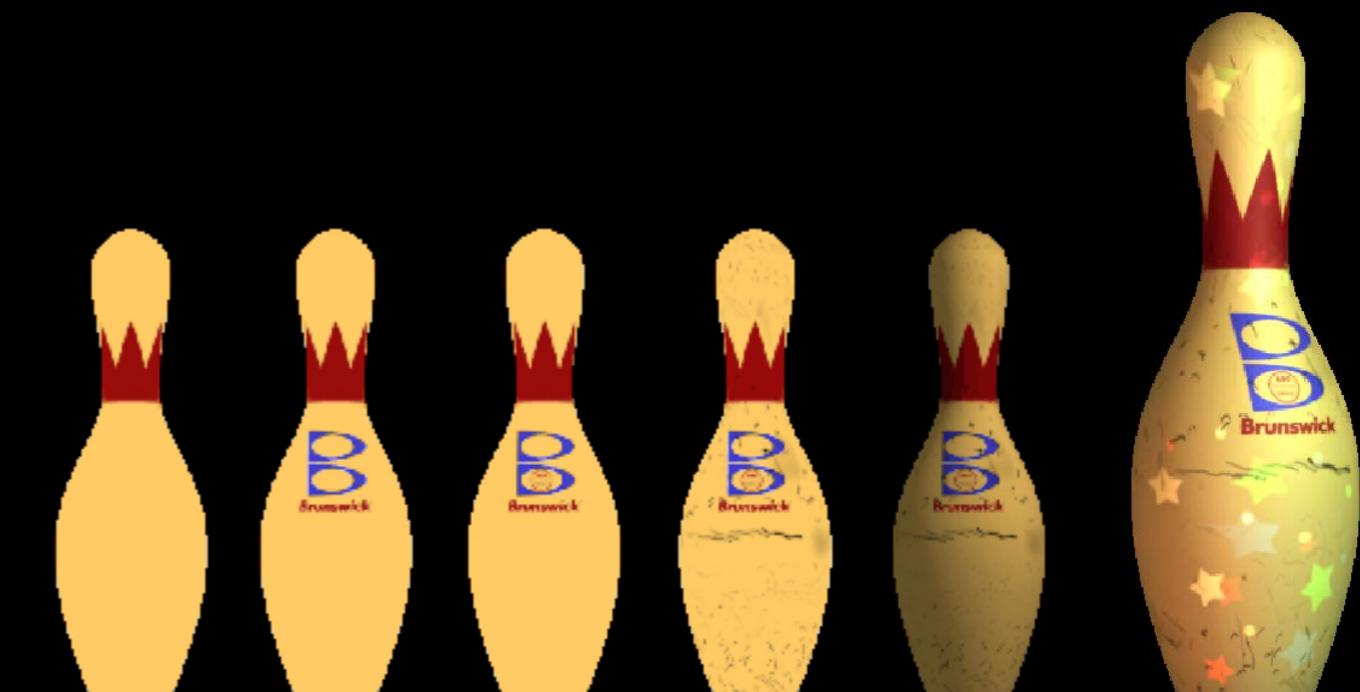
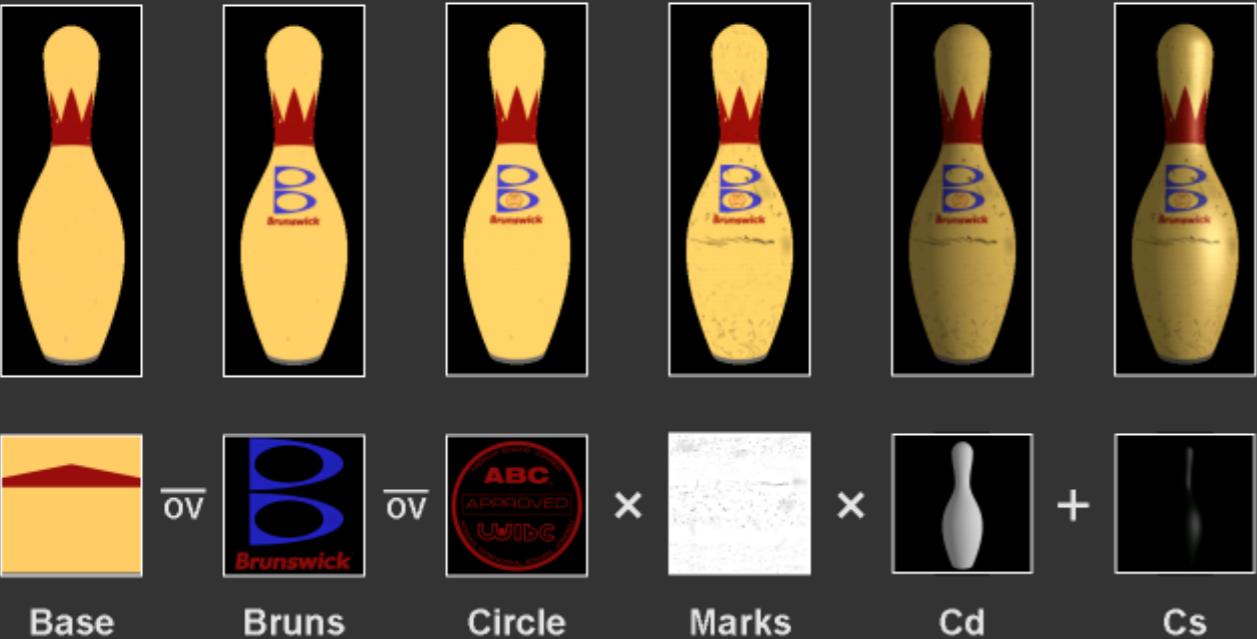
- The term “transform/vertex/geometry bound” often means the bottleneck is “anywhere before the rasterizer”
- The term “fill/raster bound” often means the bottleneck is “anywhere after setup for rasterization” (computation of edge equations)
- Can be both transform and fill bound over the course of a single frame!

Questions?

Shader zoo

Layering

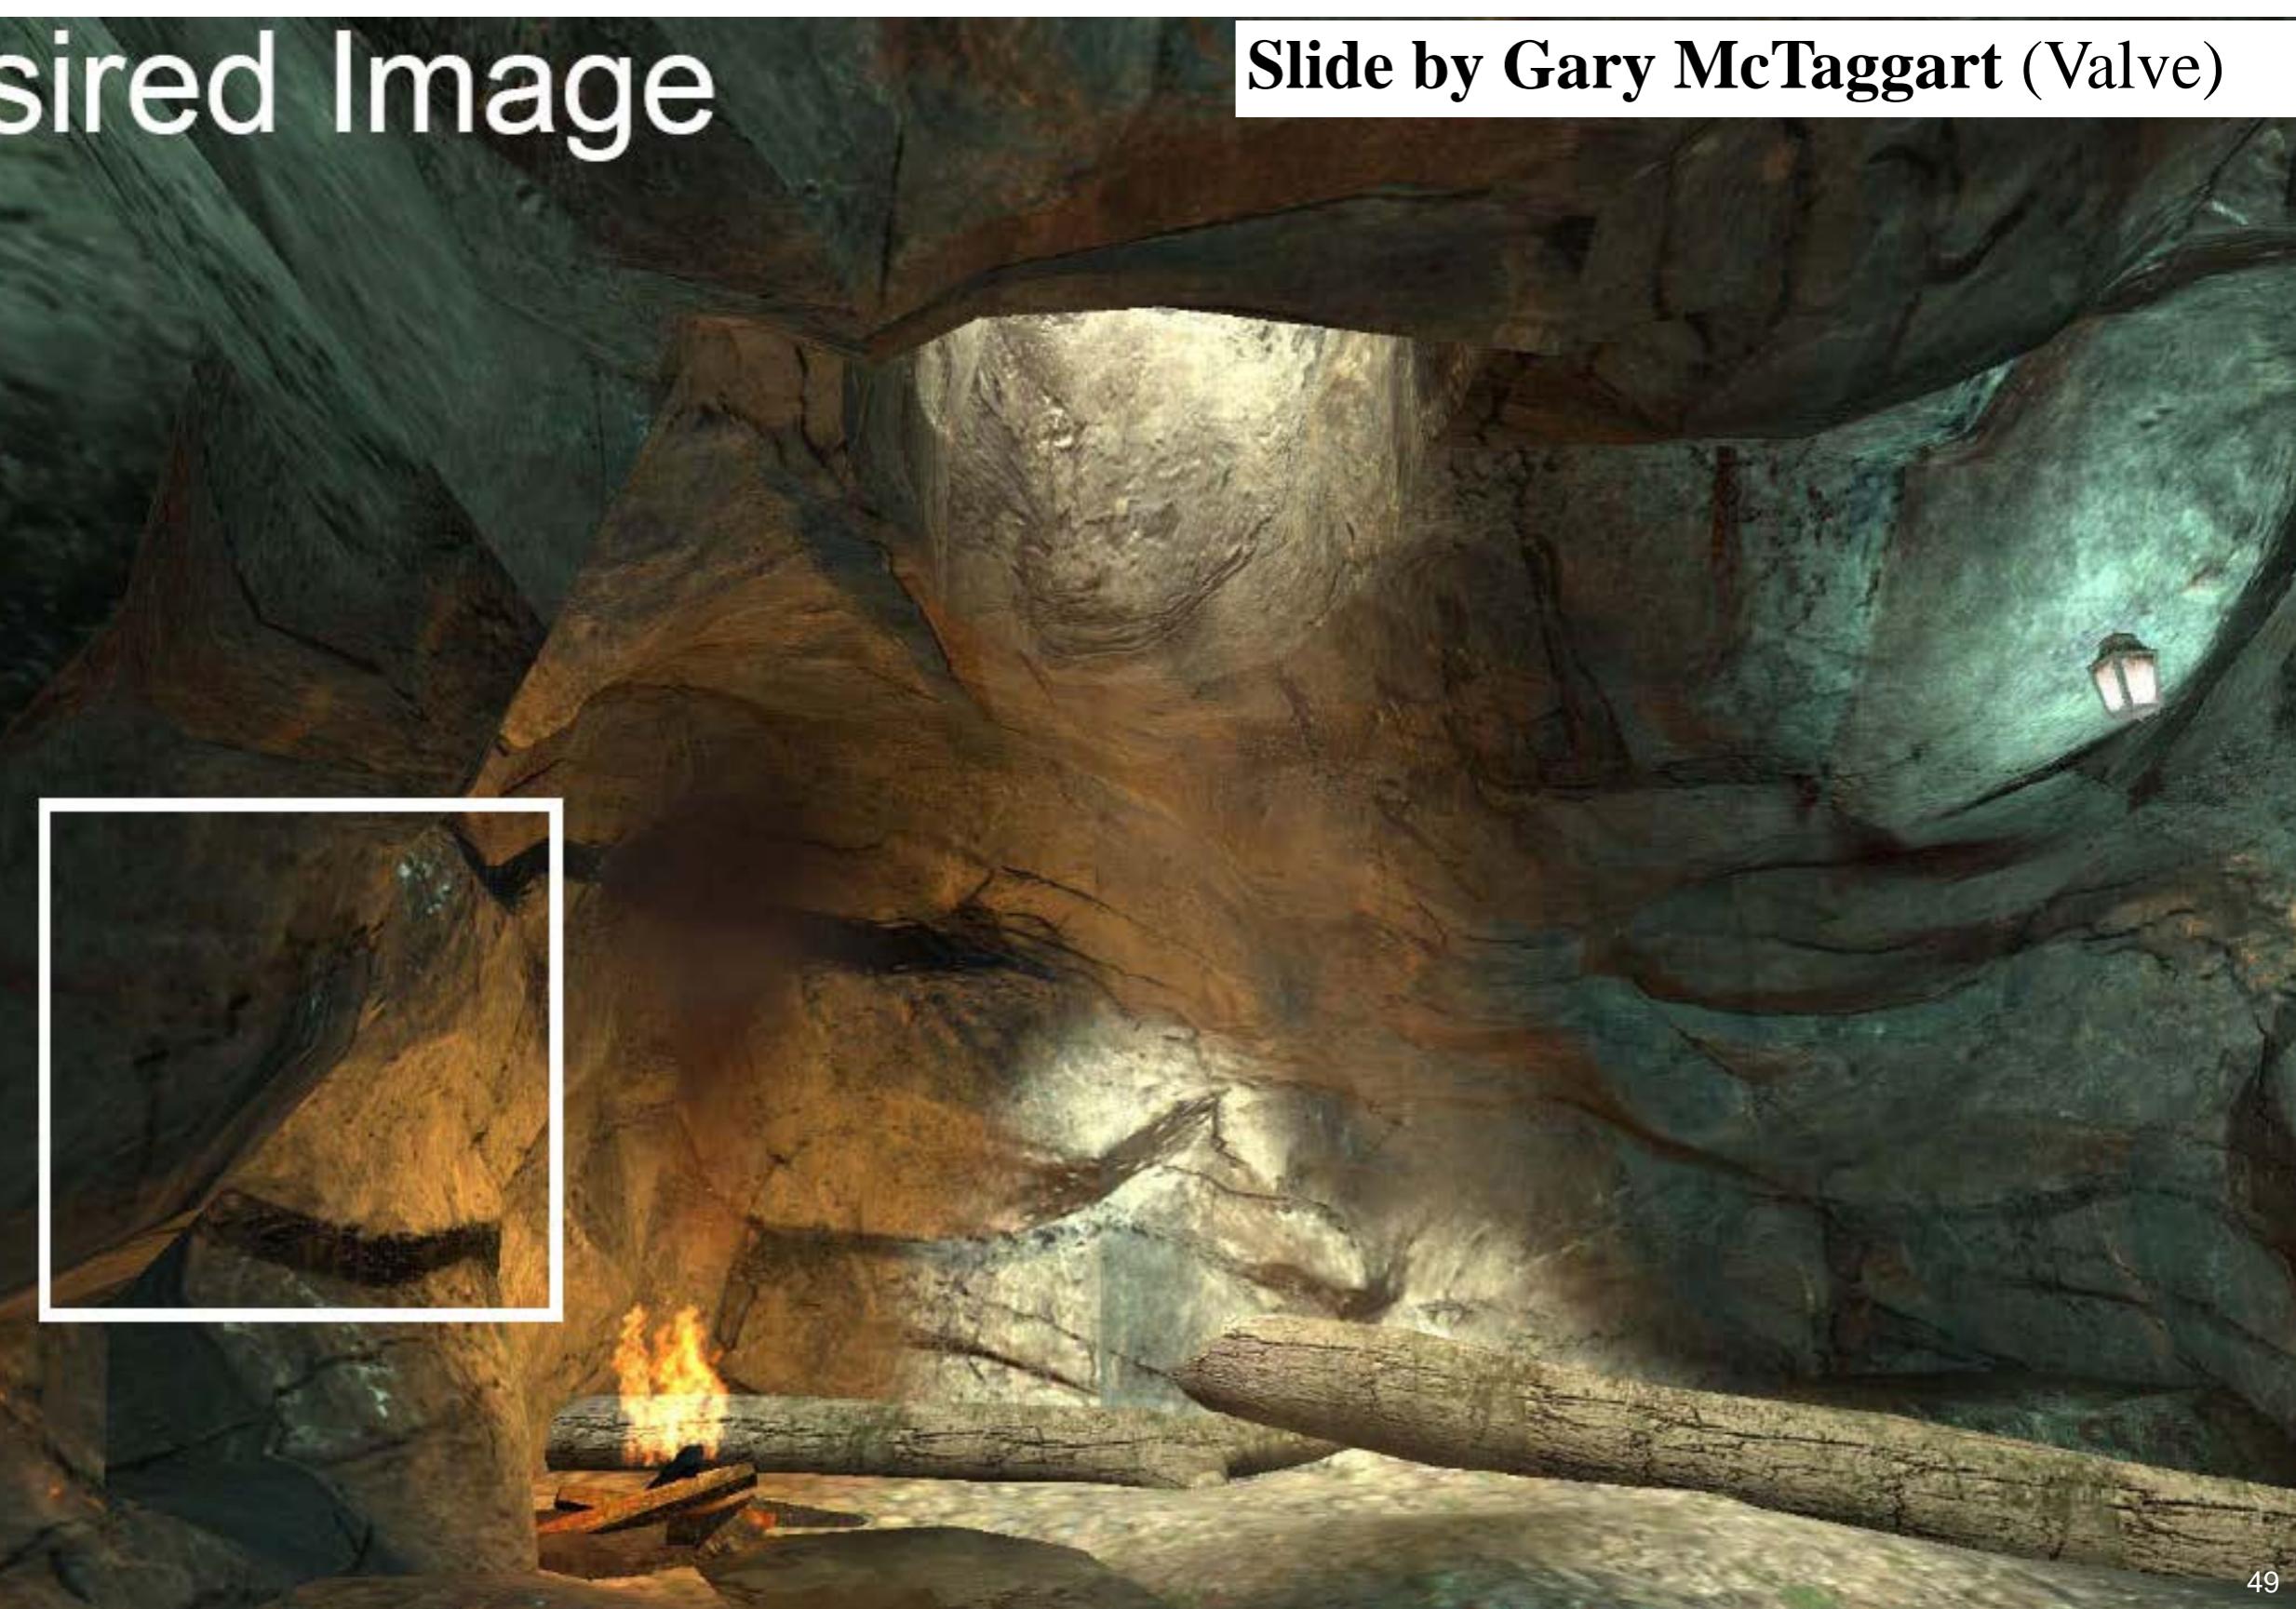
```
// bowling pin, based on RenderMan example  
(CIRCLE over (BRUNS over BASE)) * MARKS * Cd + Cs
```



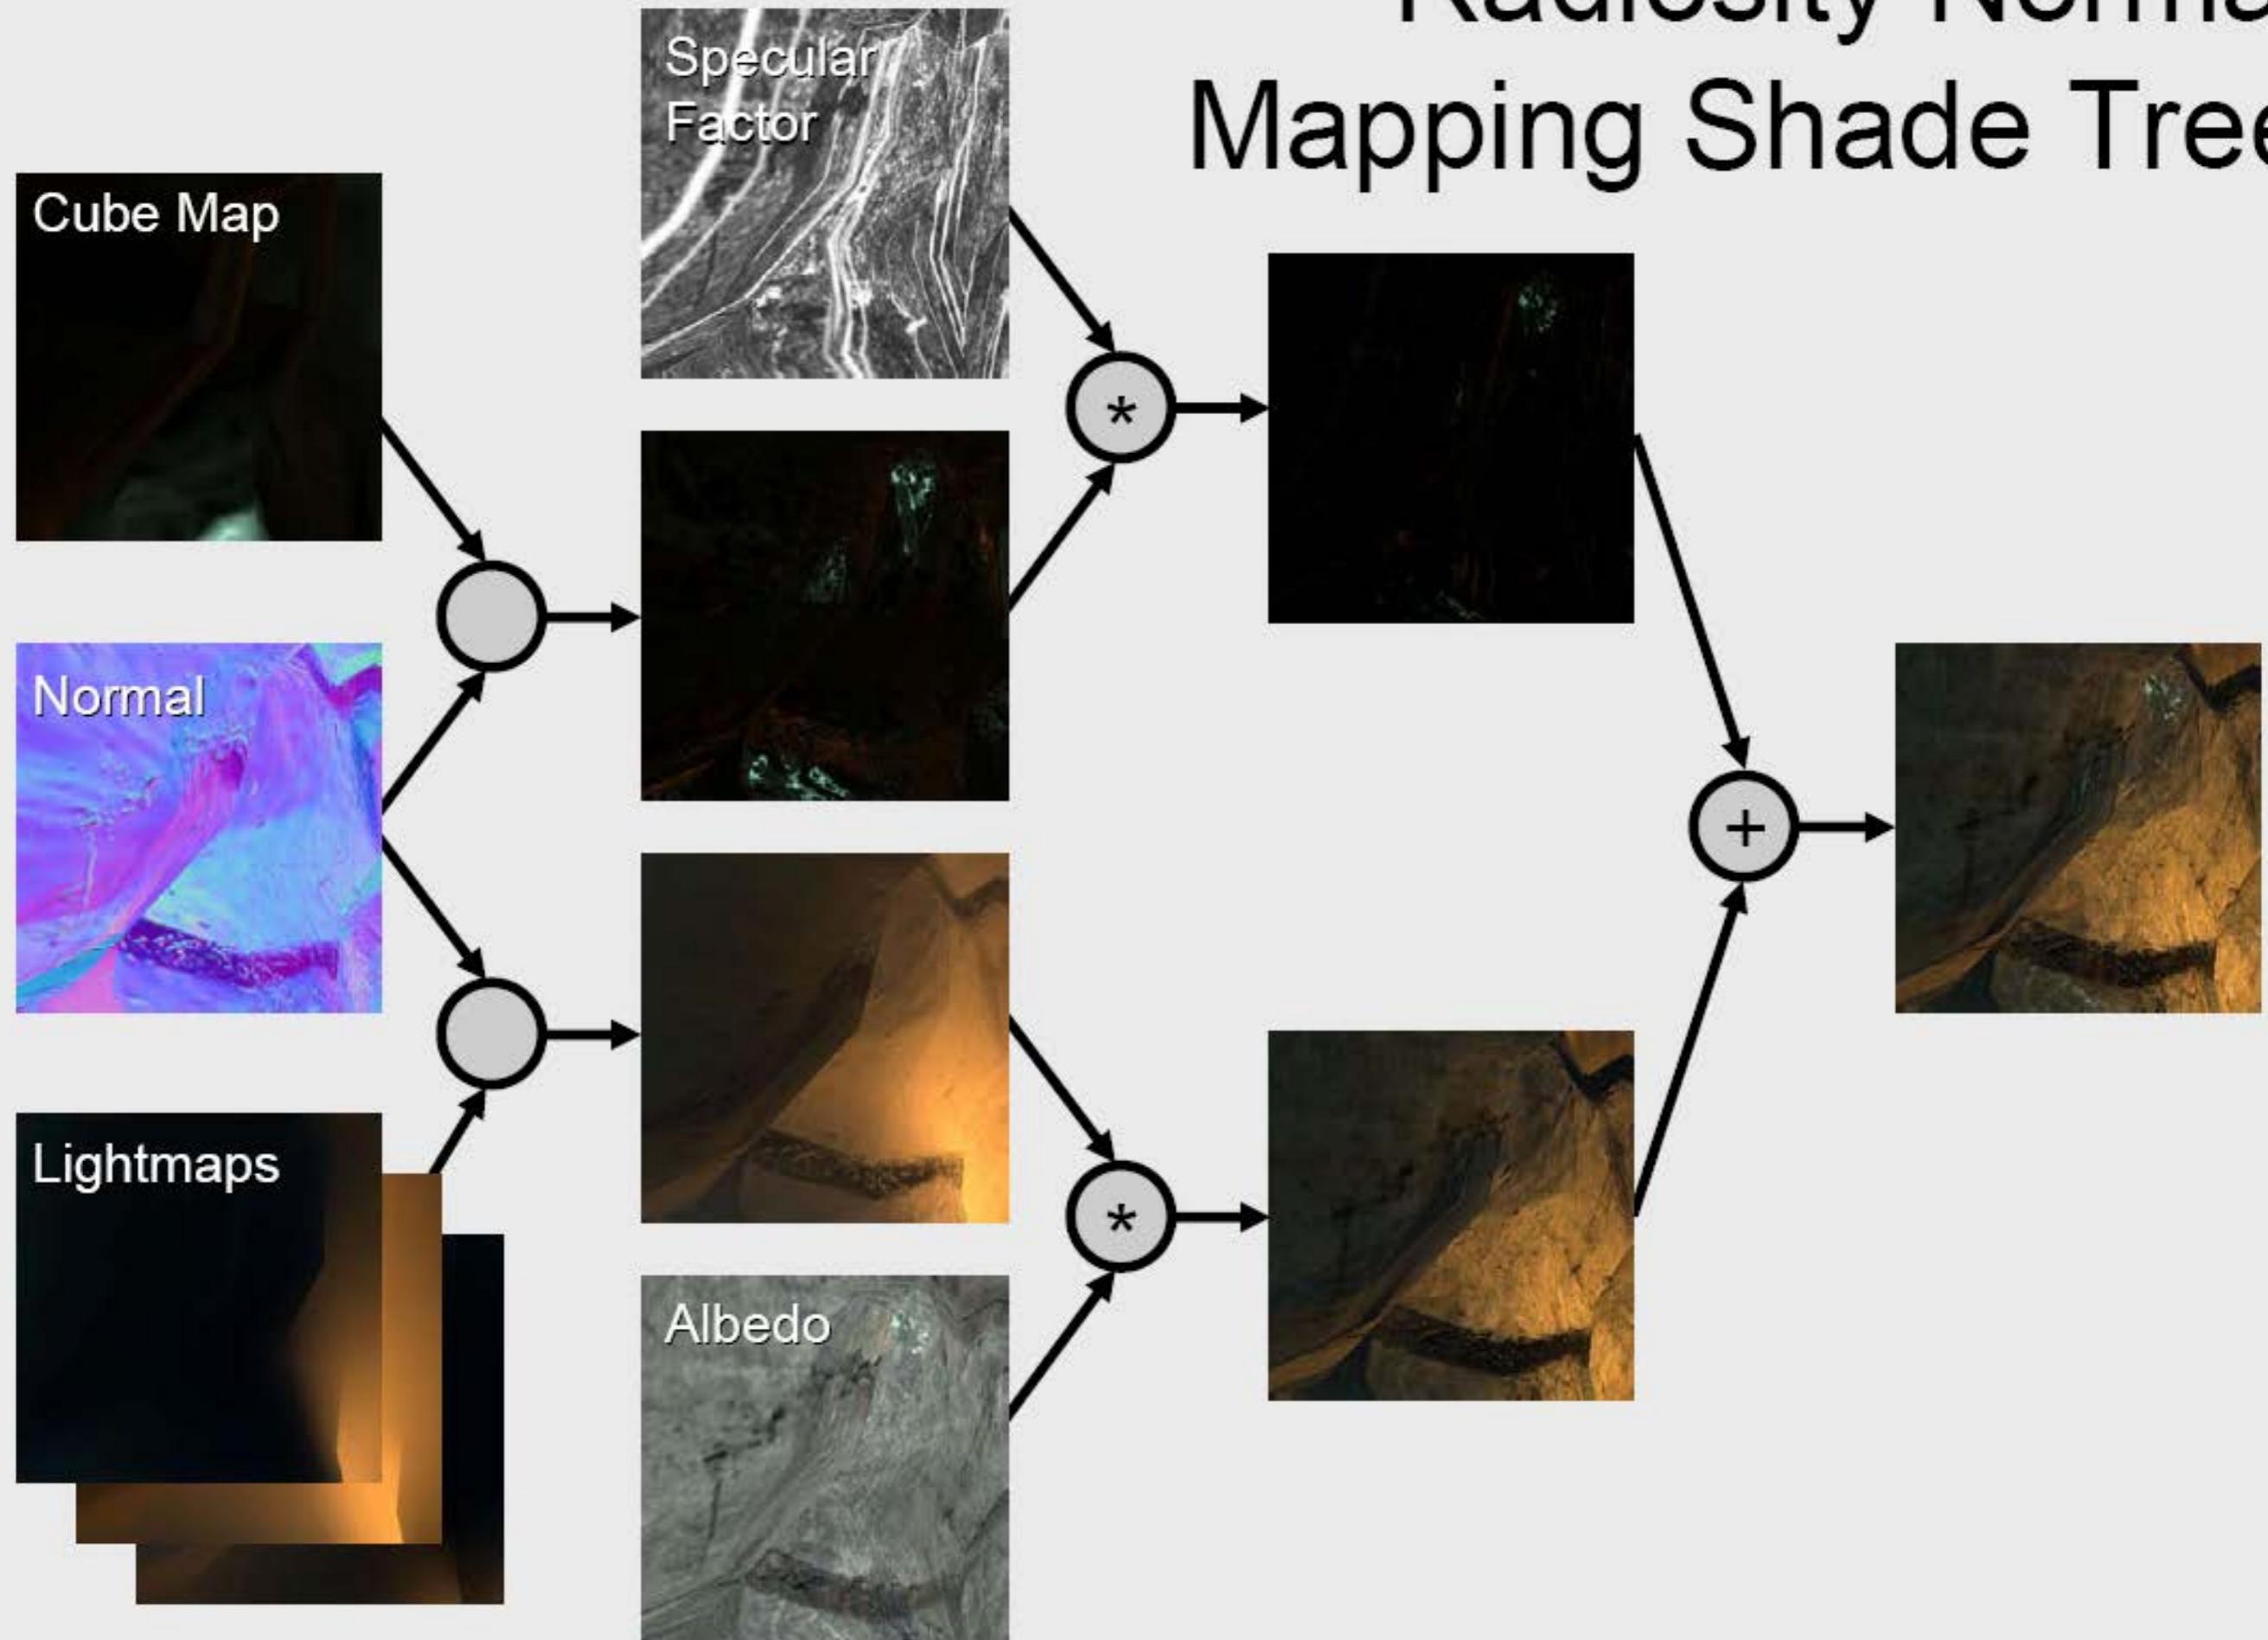
From Half Life 2 (Valve)

Desired Image

Slide by Gary McTaggart (Valve)

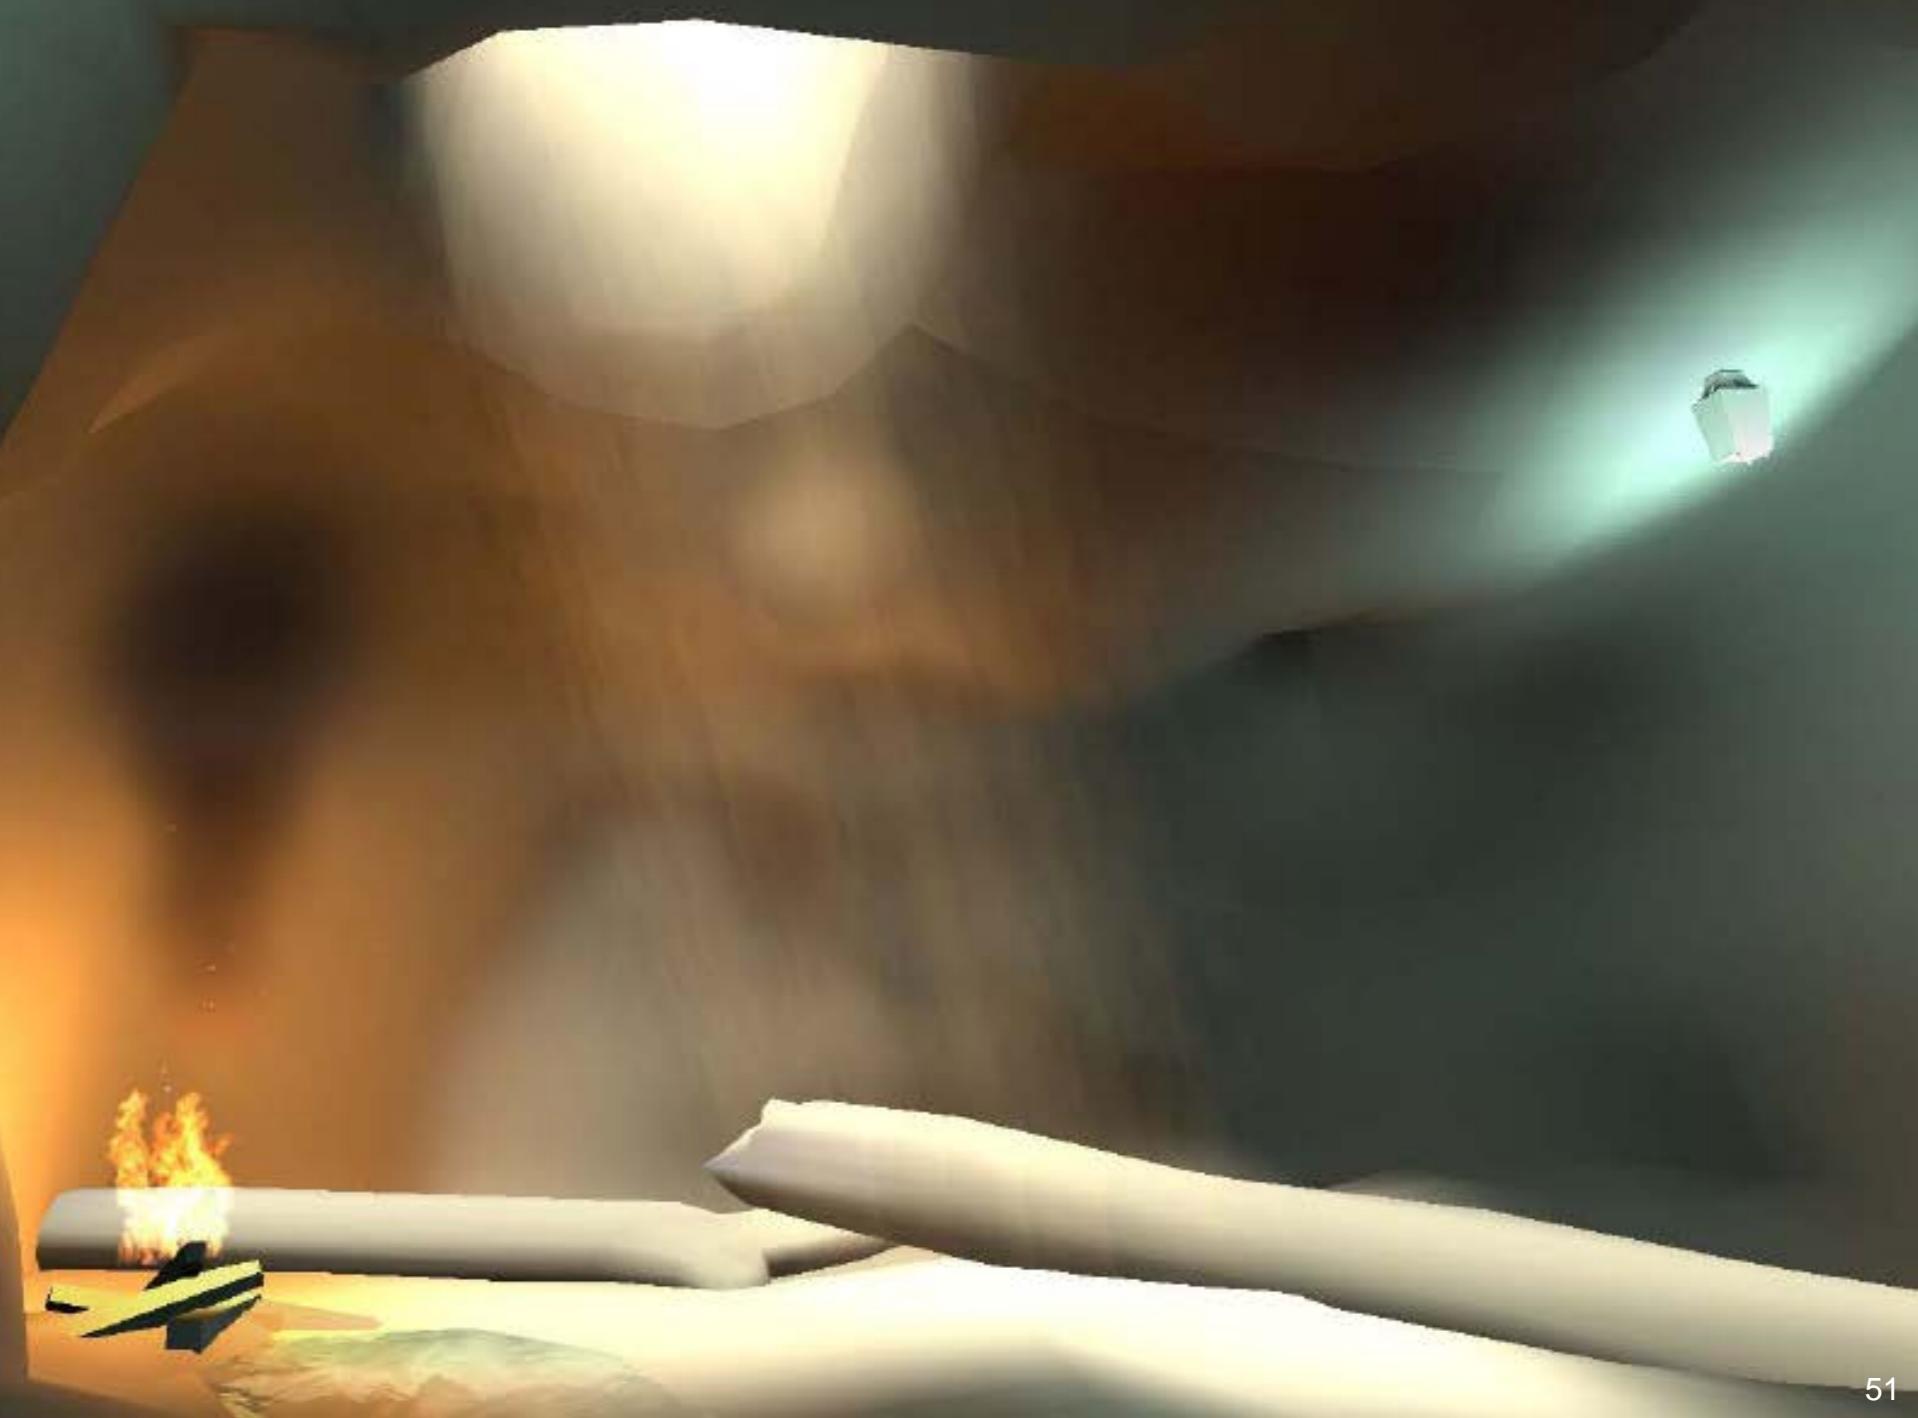


Radiosity Normal Mapping Shade Tree



Slide by Gary McTaggart (Valve)

Radiosity



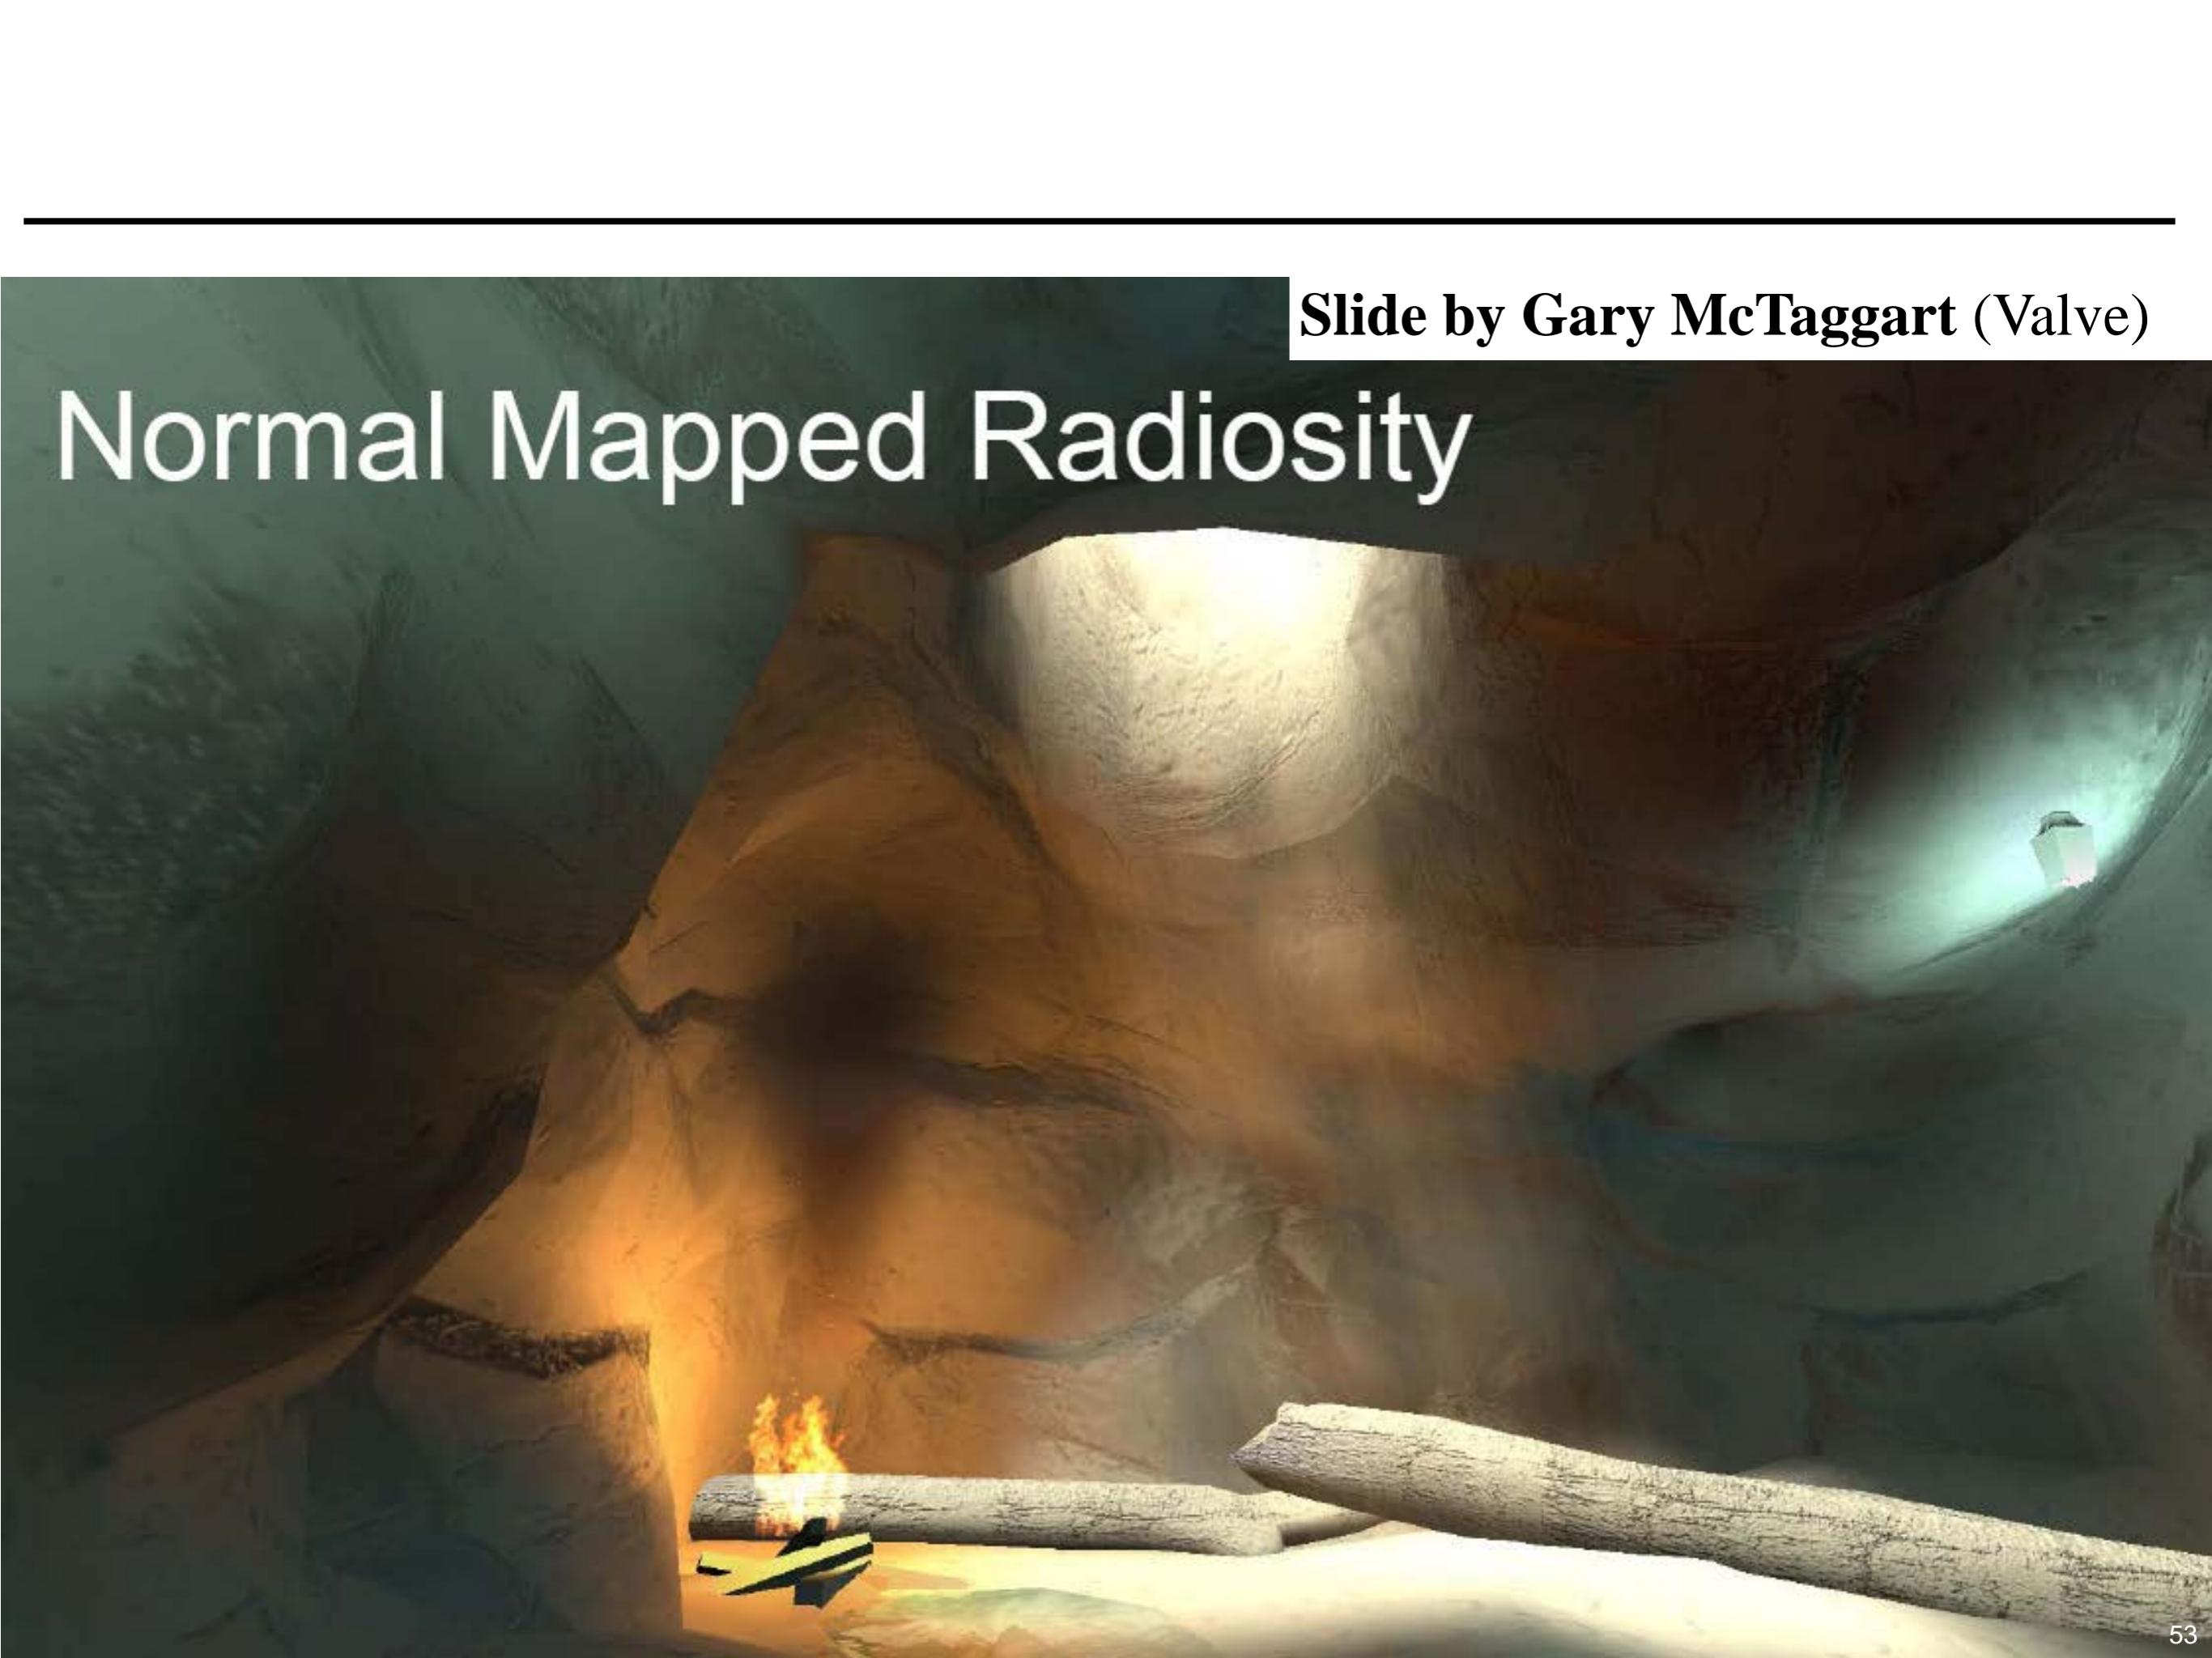
Slide by Gary McTaggart (Valve)

Normal

A 3D rendering of a scene featuring a campfire in the foreground, a path leading towards a building, and a street lamp on the right. The entire scene is covered with a dense, colorful field of normal vectors, represented by short lines in shades of blue, green, and purple, which highlight the surface geometry and lighting.

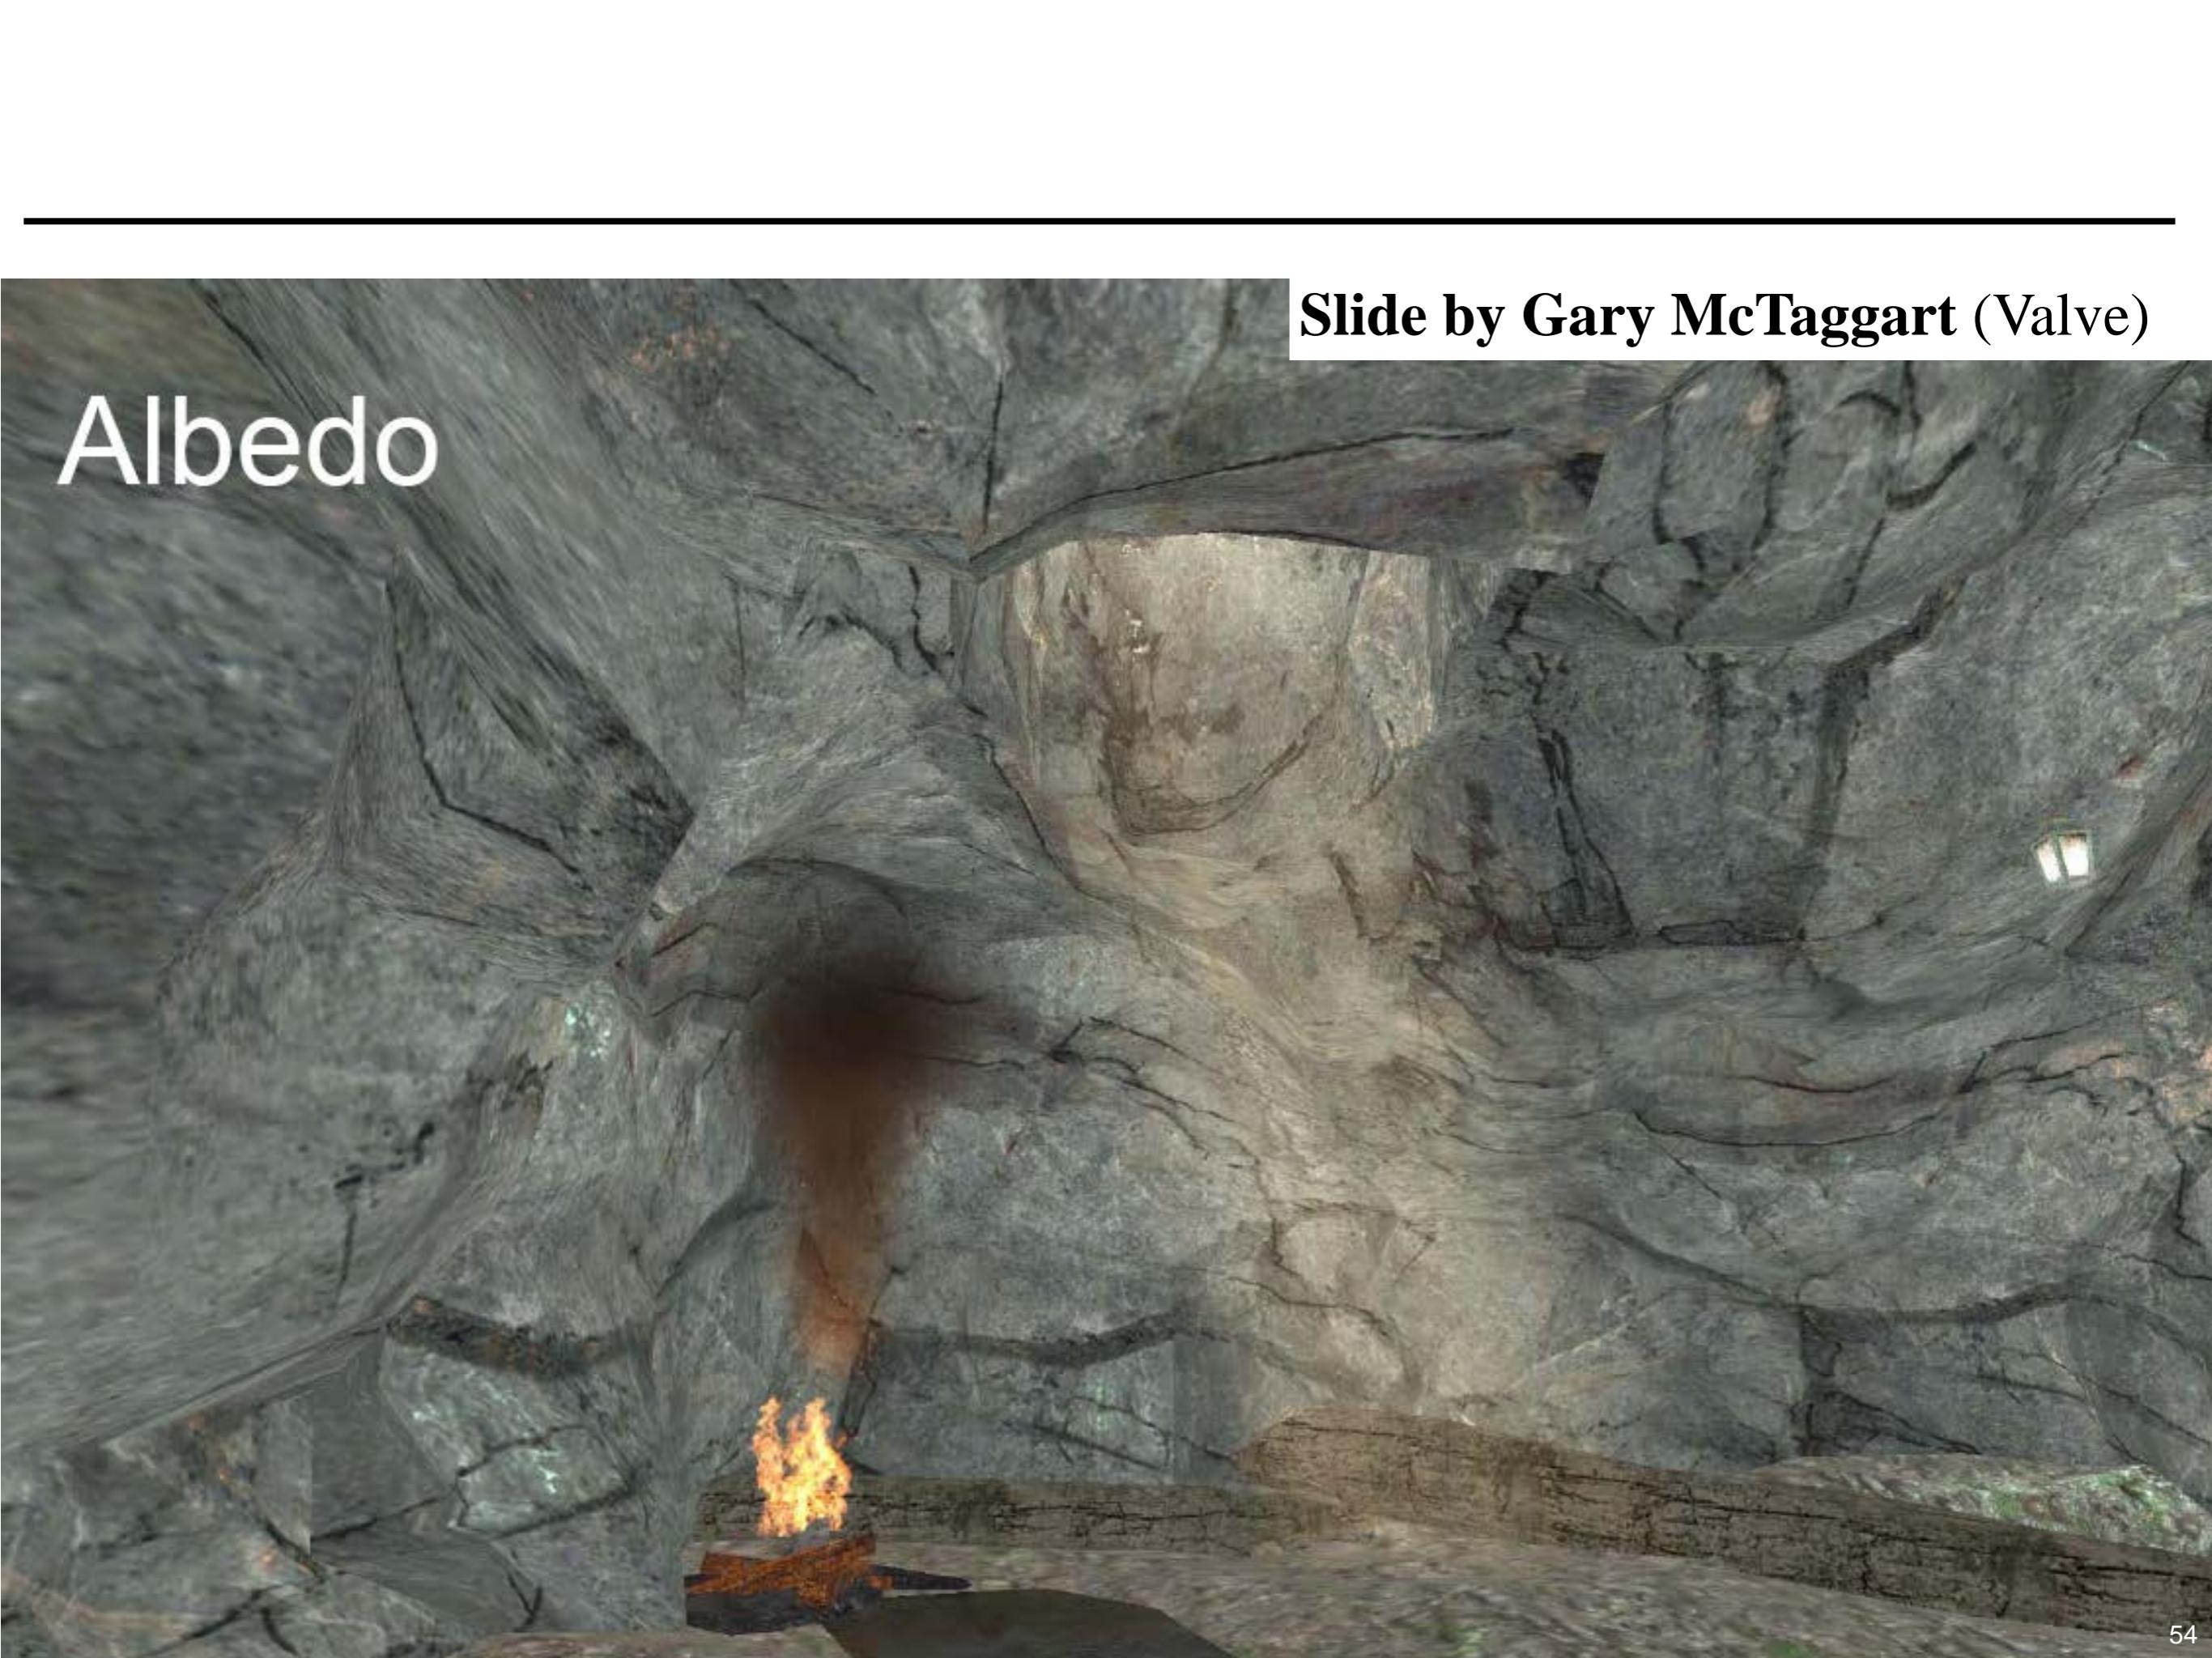
Slide by Gary McTaggart (Valve)

Normal Mapped Radiosity



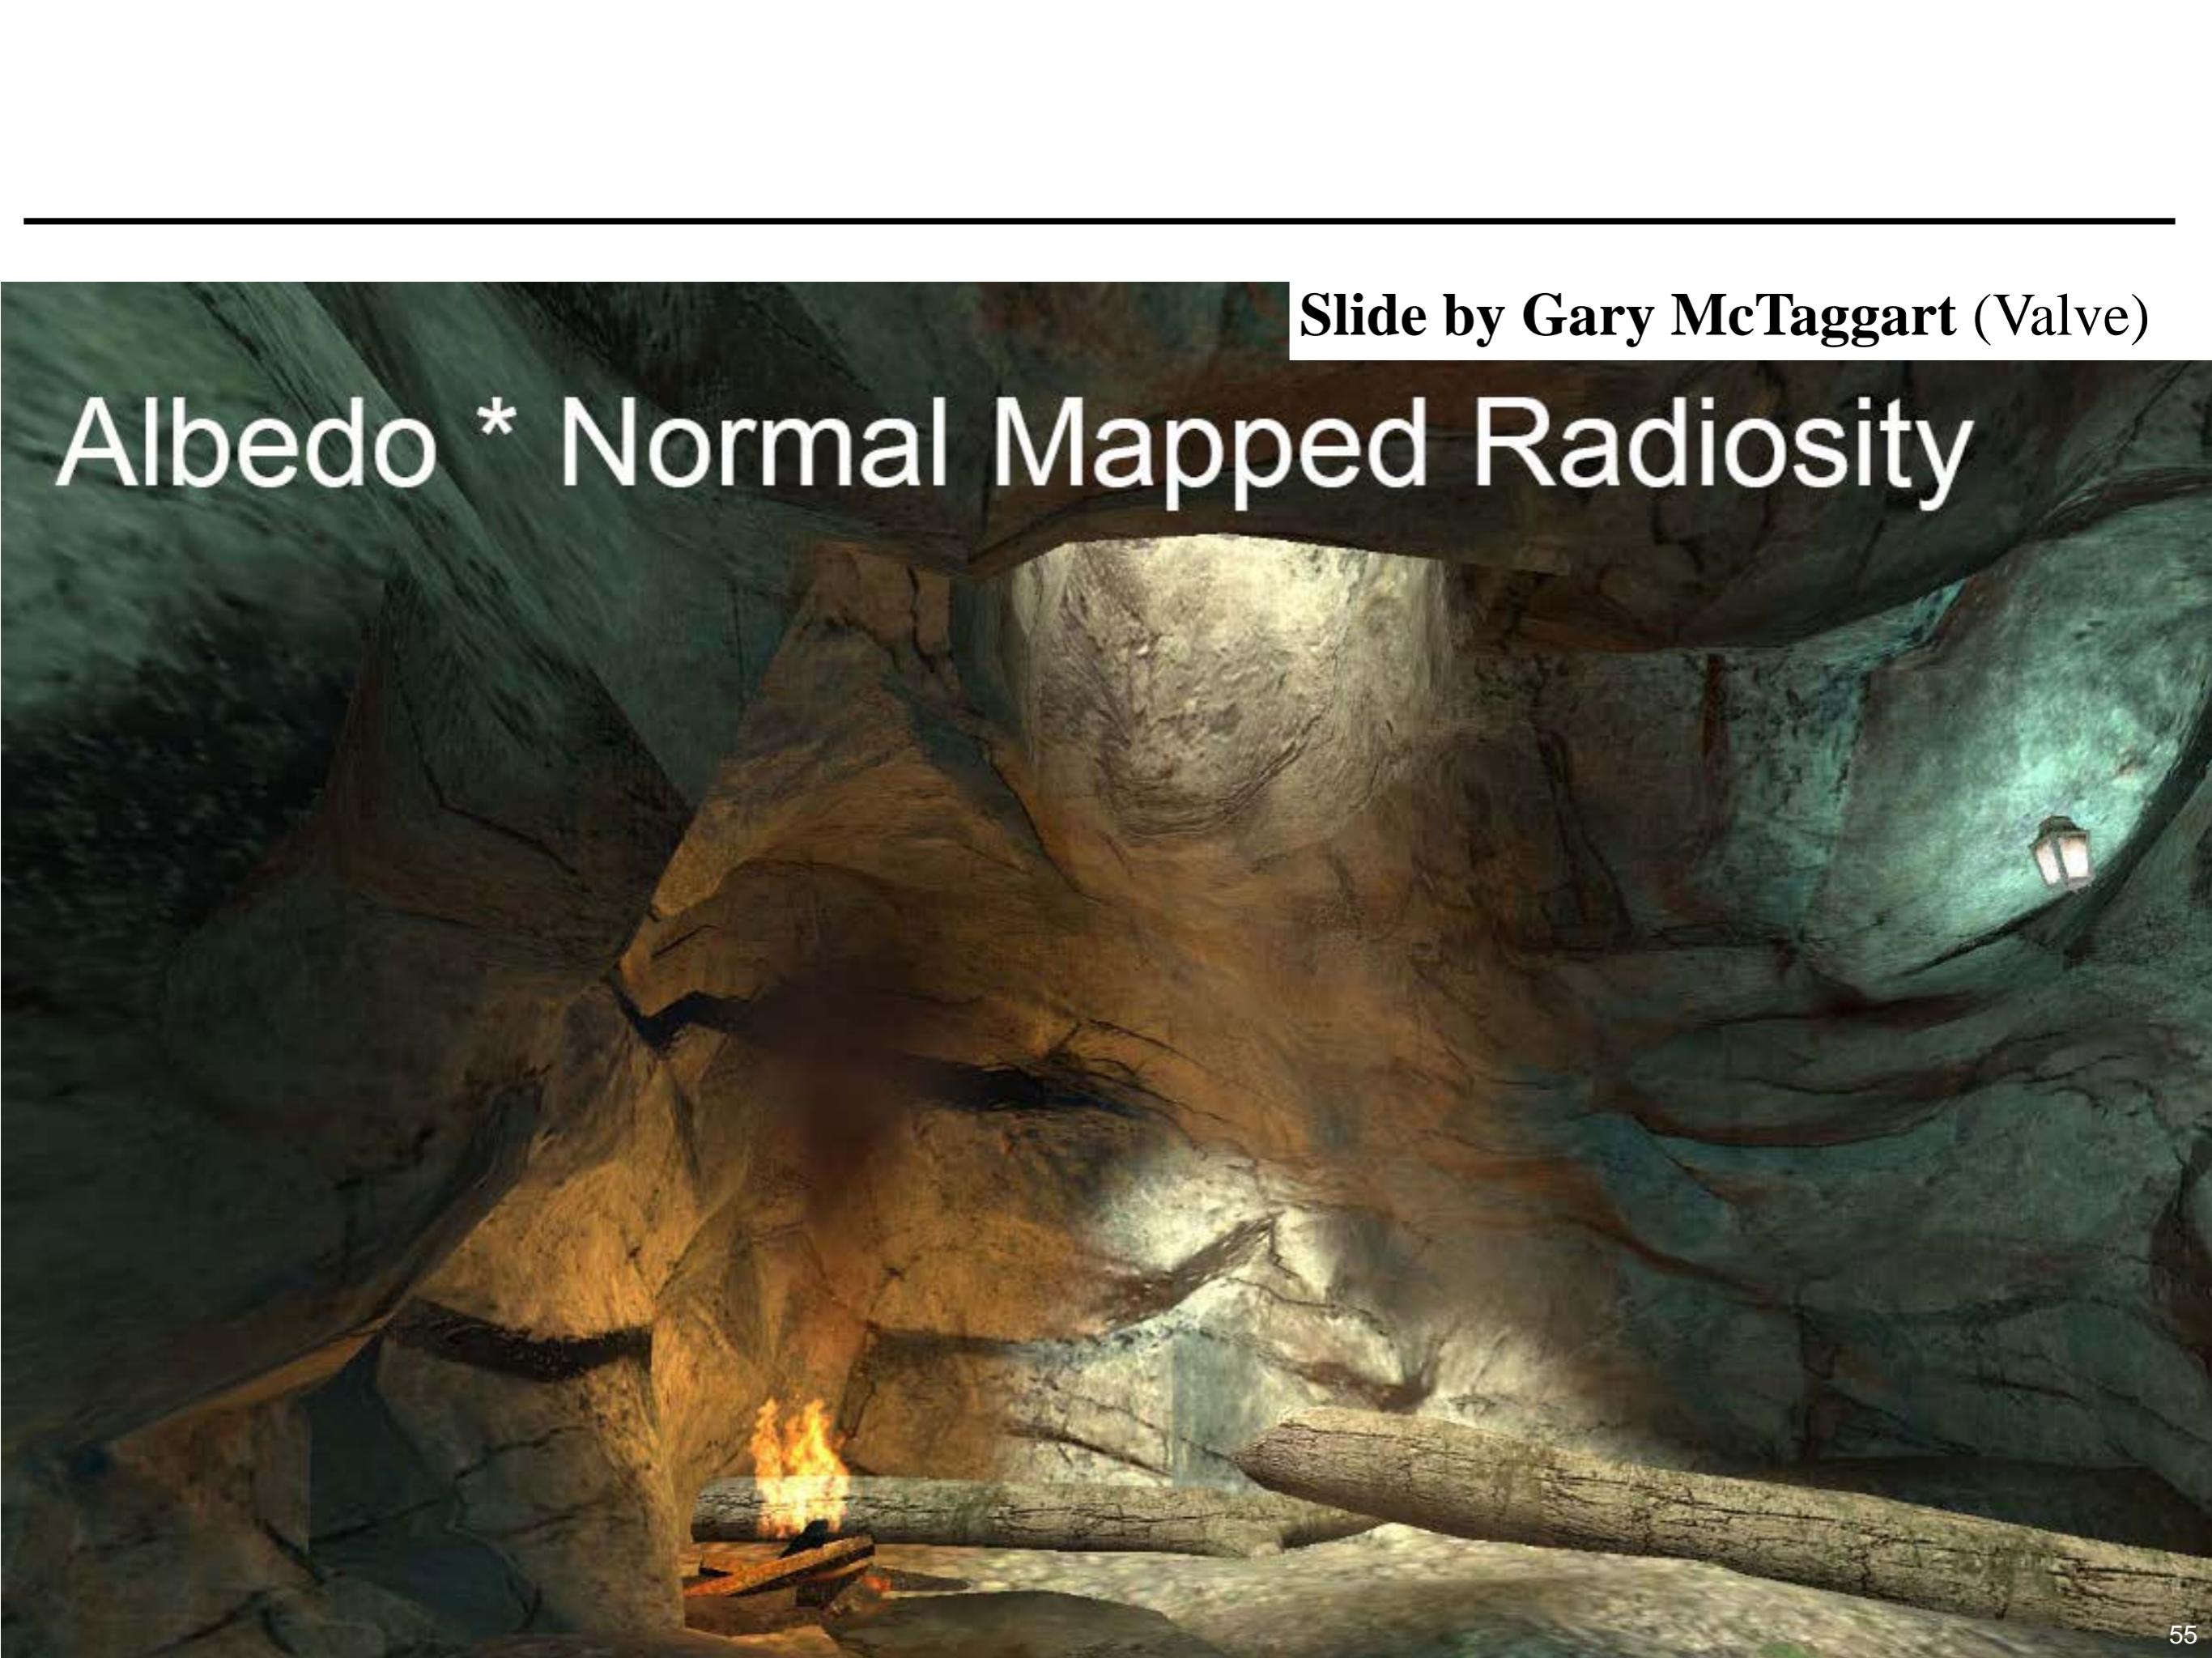
Slide by Gary McTaggart (Valve)

Albedo

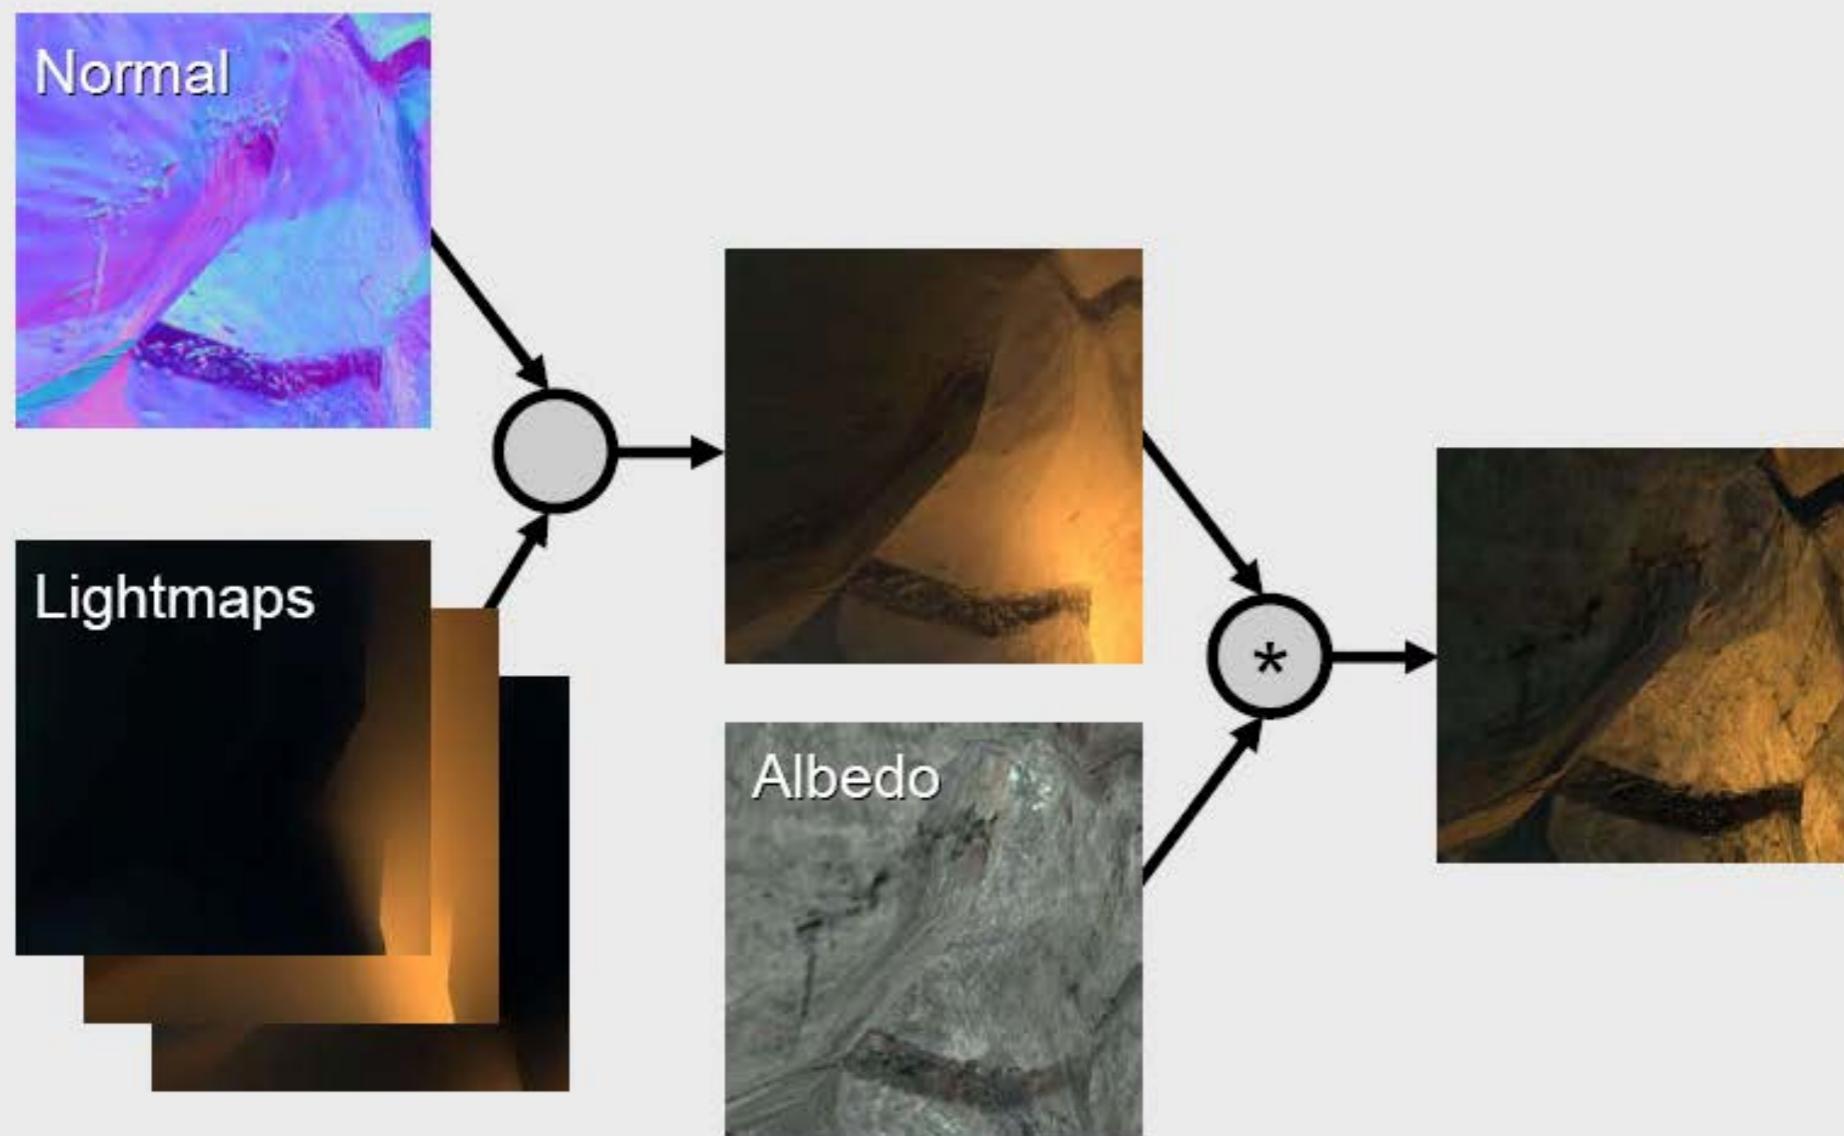
A screenshot from a Valve game, likely Half-Life, showing a dark, rocky tunnel. The walls are made of rough, grey rock with various cracks and textures. In the bottom left corner, there is a small campfire with orange flames. The lighting is dim, coming from the fire and some small, glowing greenish-yellow spots on the ceiling, possibly from a nearby light source or a screen. The overall atmosphere is dark and atmospheric.

Slide by Gary McTaggart (Valve)

Albedo * Normal Mapped Radiosity



Radiosity Normal Mapping Shade Tree



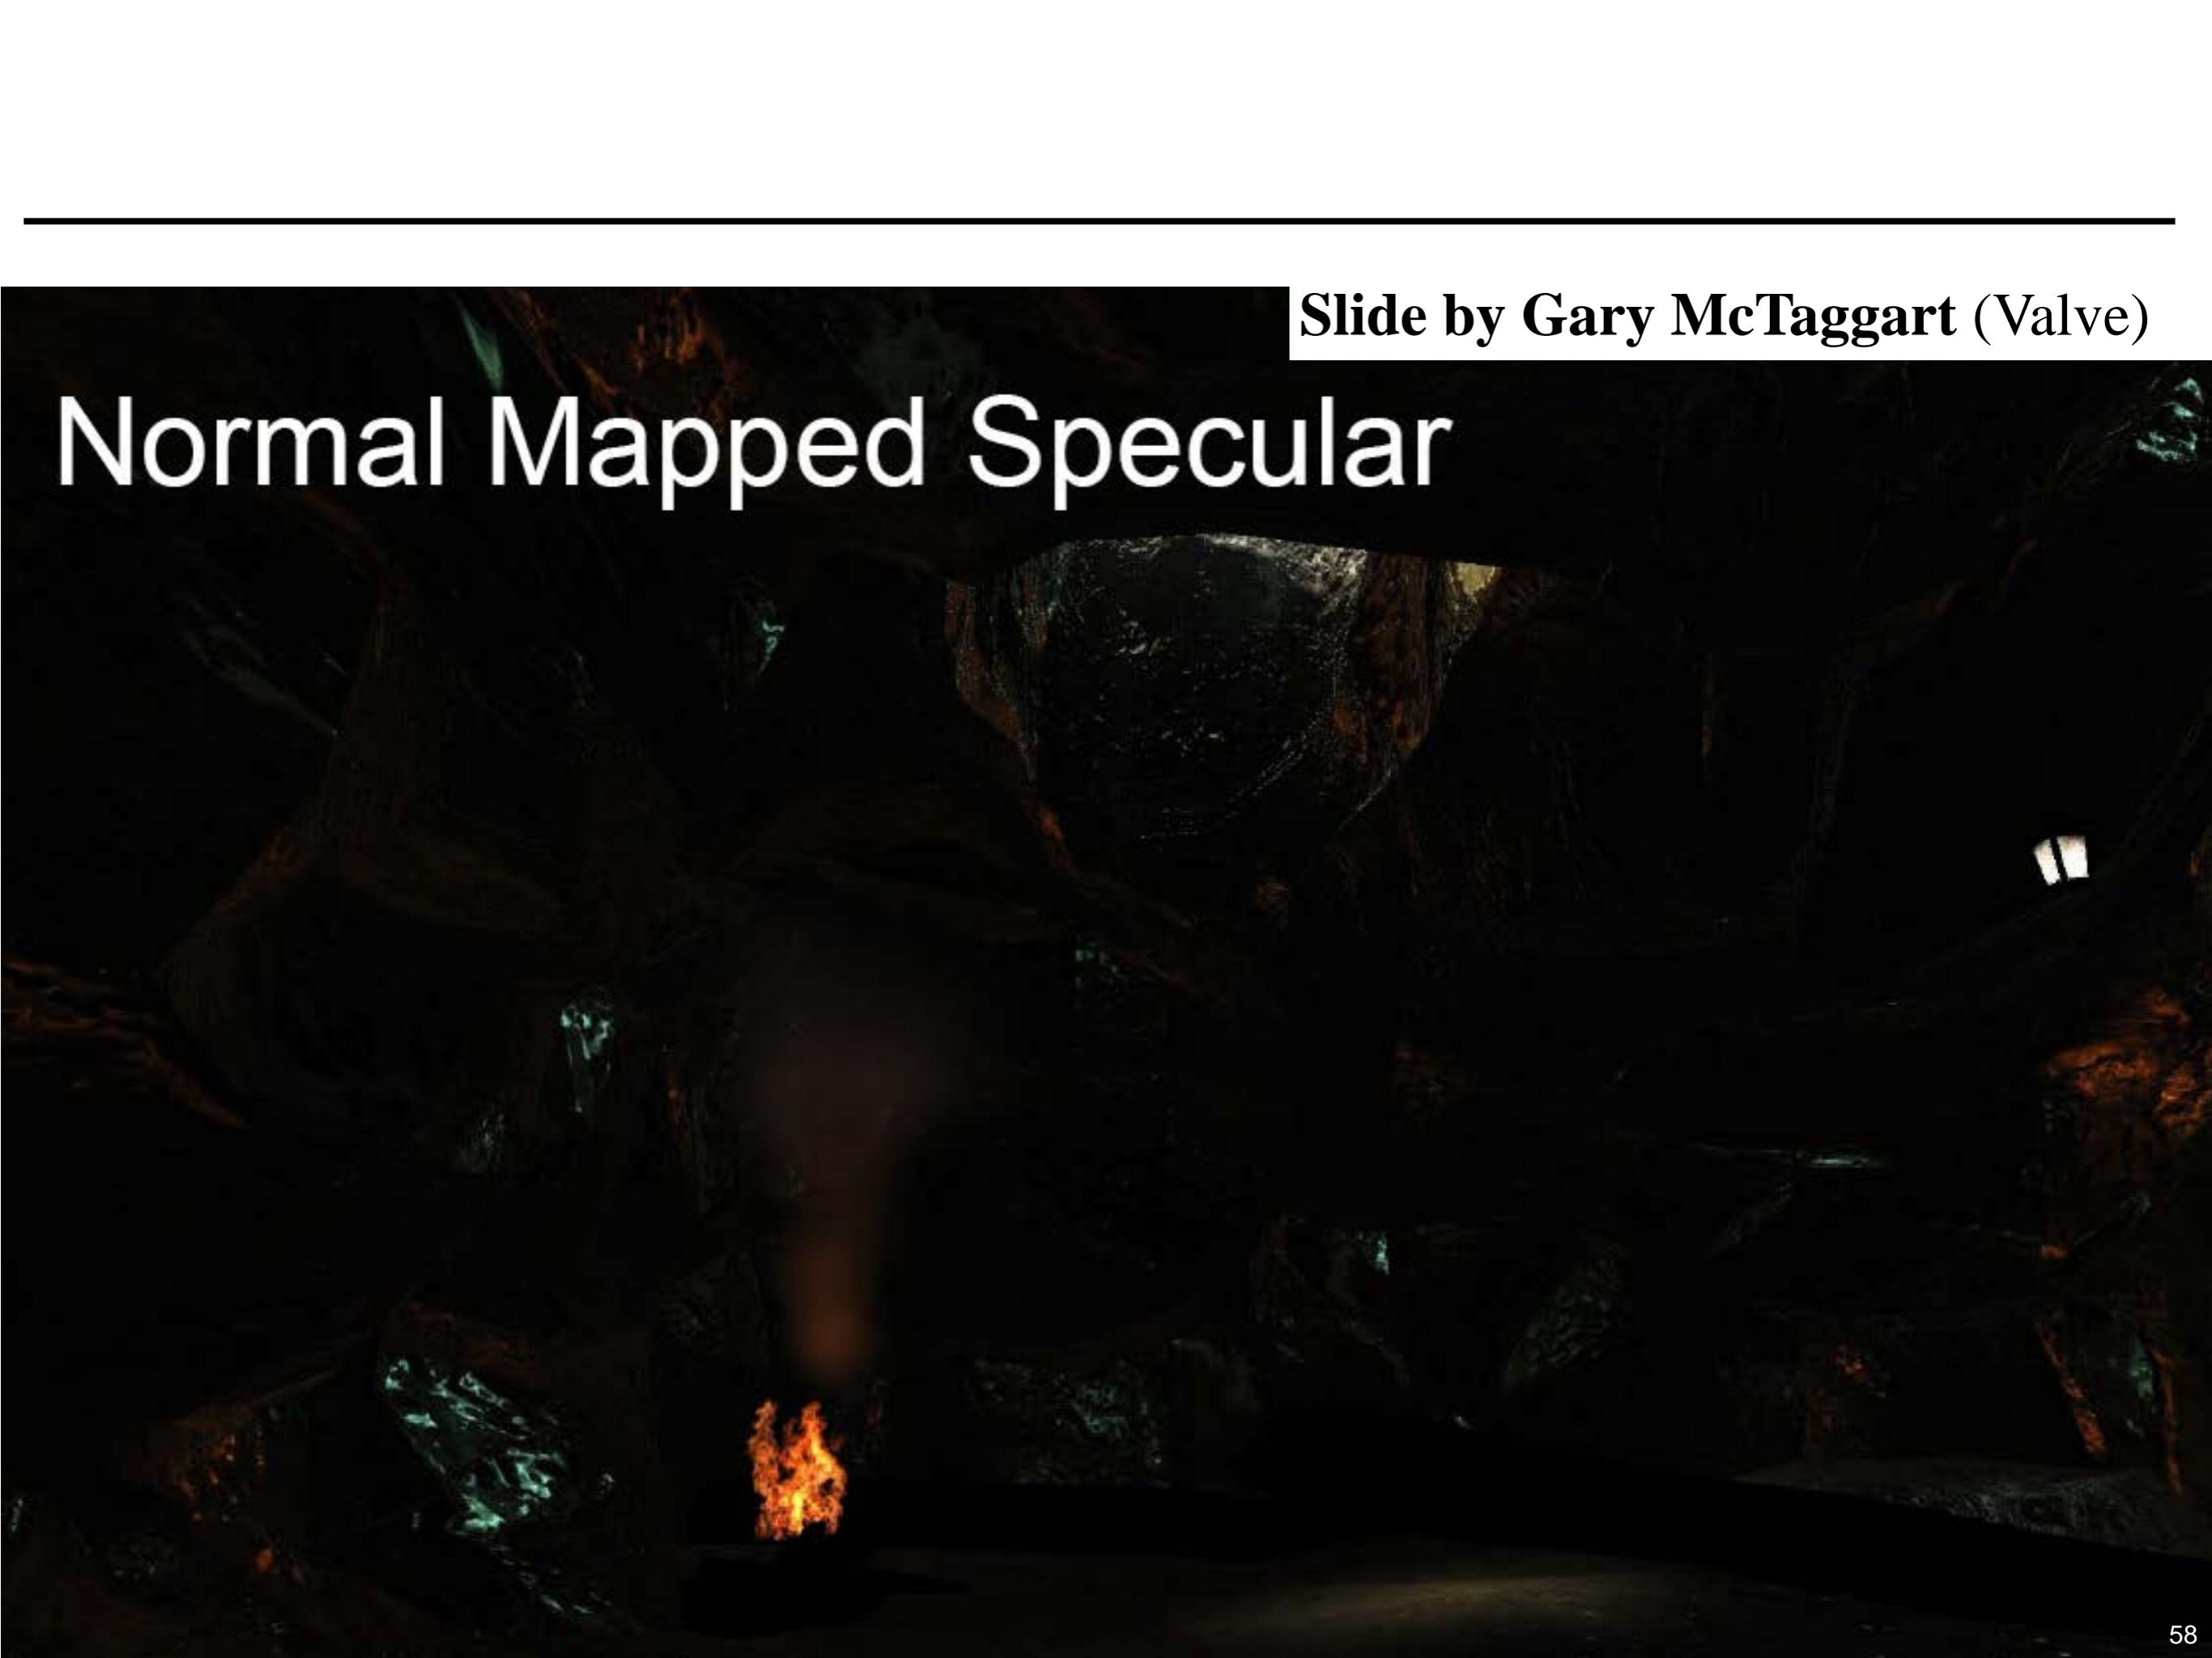
Slide by Gary McTaggart (Valve)

Cube Map Specular



Slide by Gary McTaggart (Valve)

Normal Mapped Specular



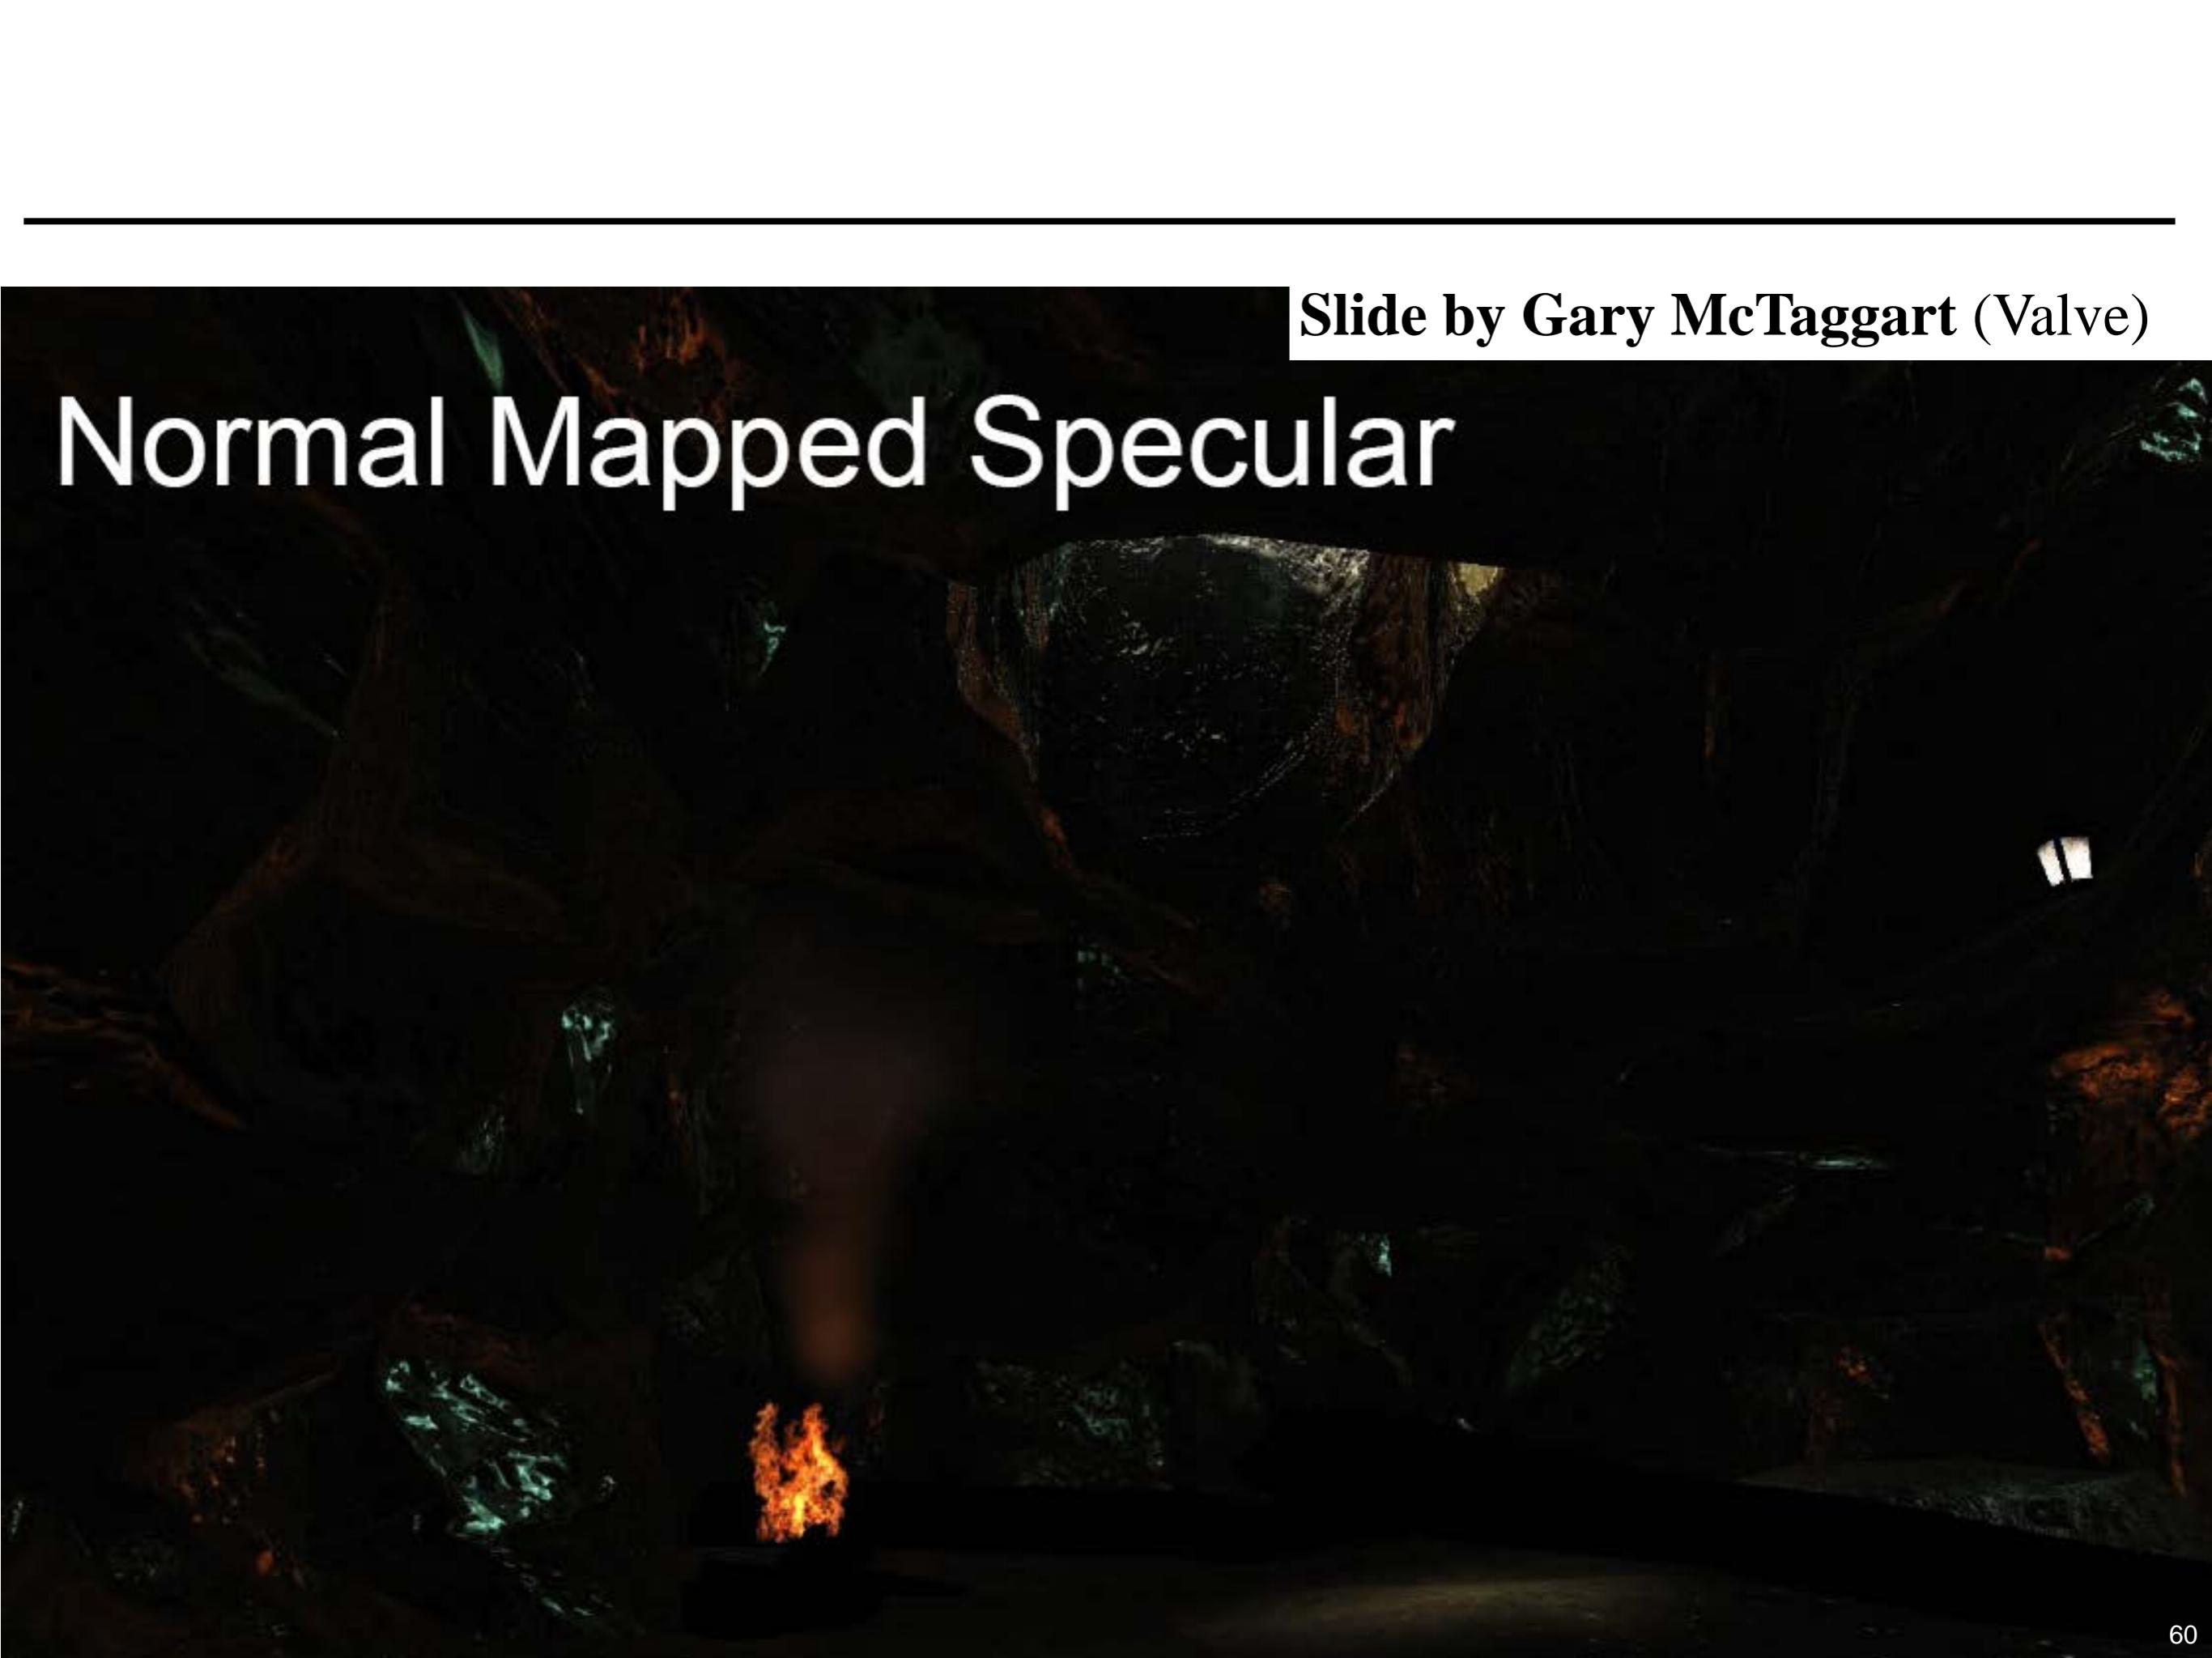
Slide by Gary McTaggart (Valve)

Specular Factor



Slide by Gary McTaggart (Valve)

Normal Mapped Specular

A dark, atmospheric scene from a video game. The environment is a dense forest at night or in low light conditions. In the foreground, there is a small campfire with bright orange and yellow flames. To the left, several glowing mushrooms with green caps and red stems are scattered on the ground. The path ahead is dimly lit, with some distant lights visible through the trees. The overall mood is mysterious and immersive.

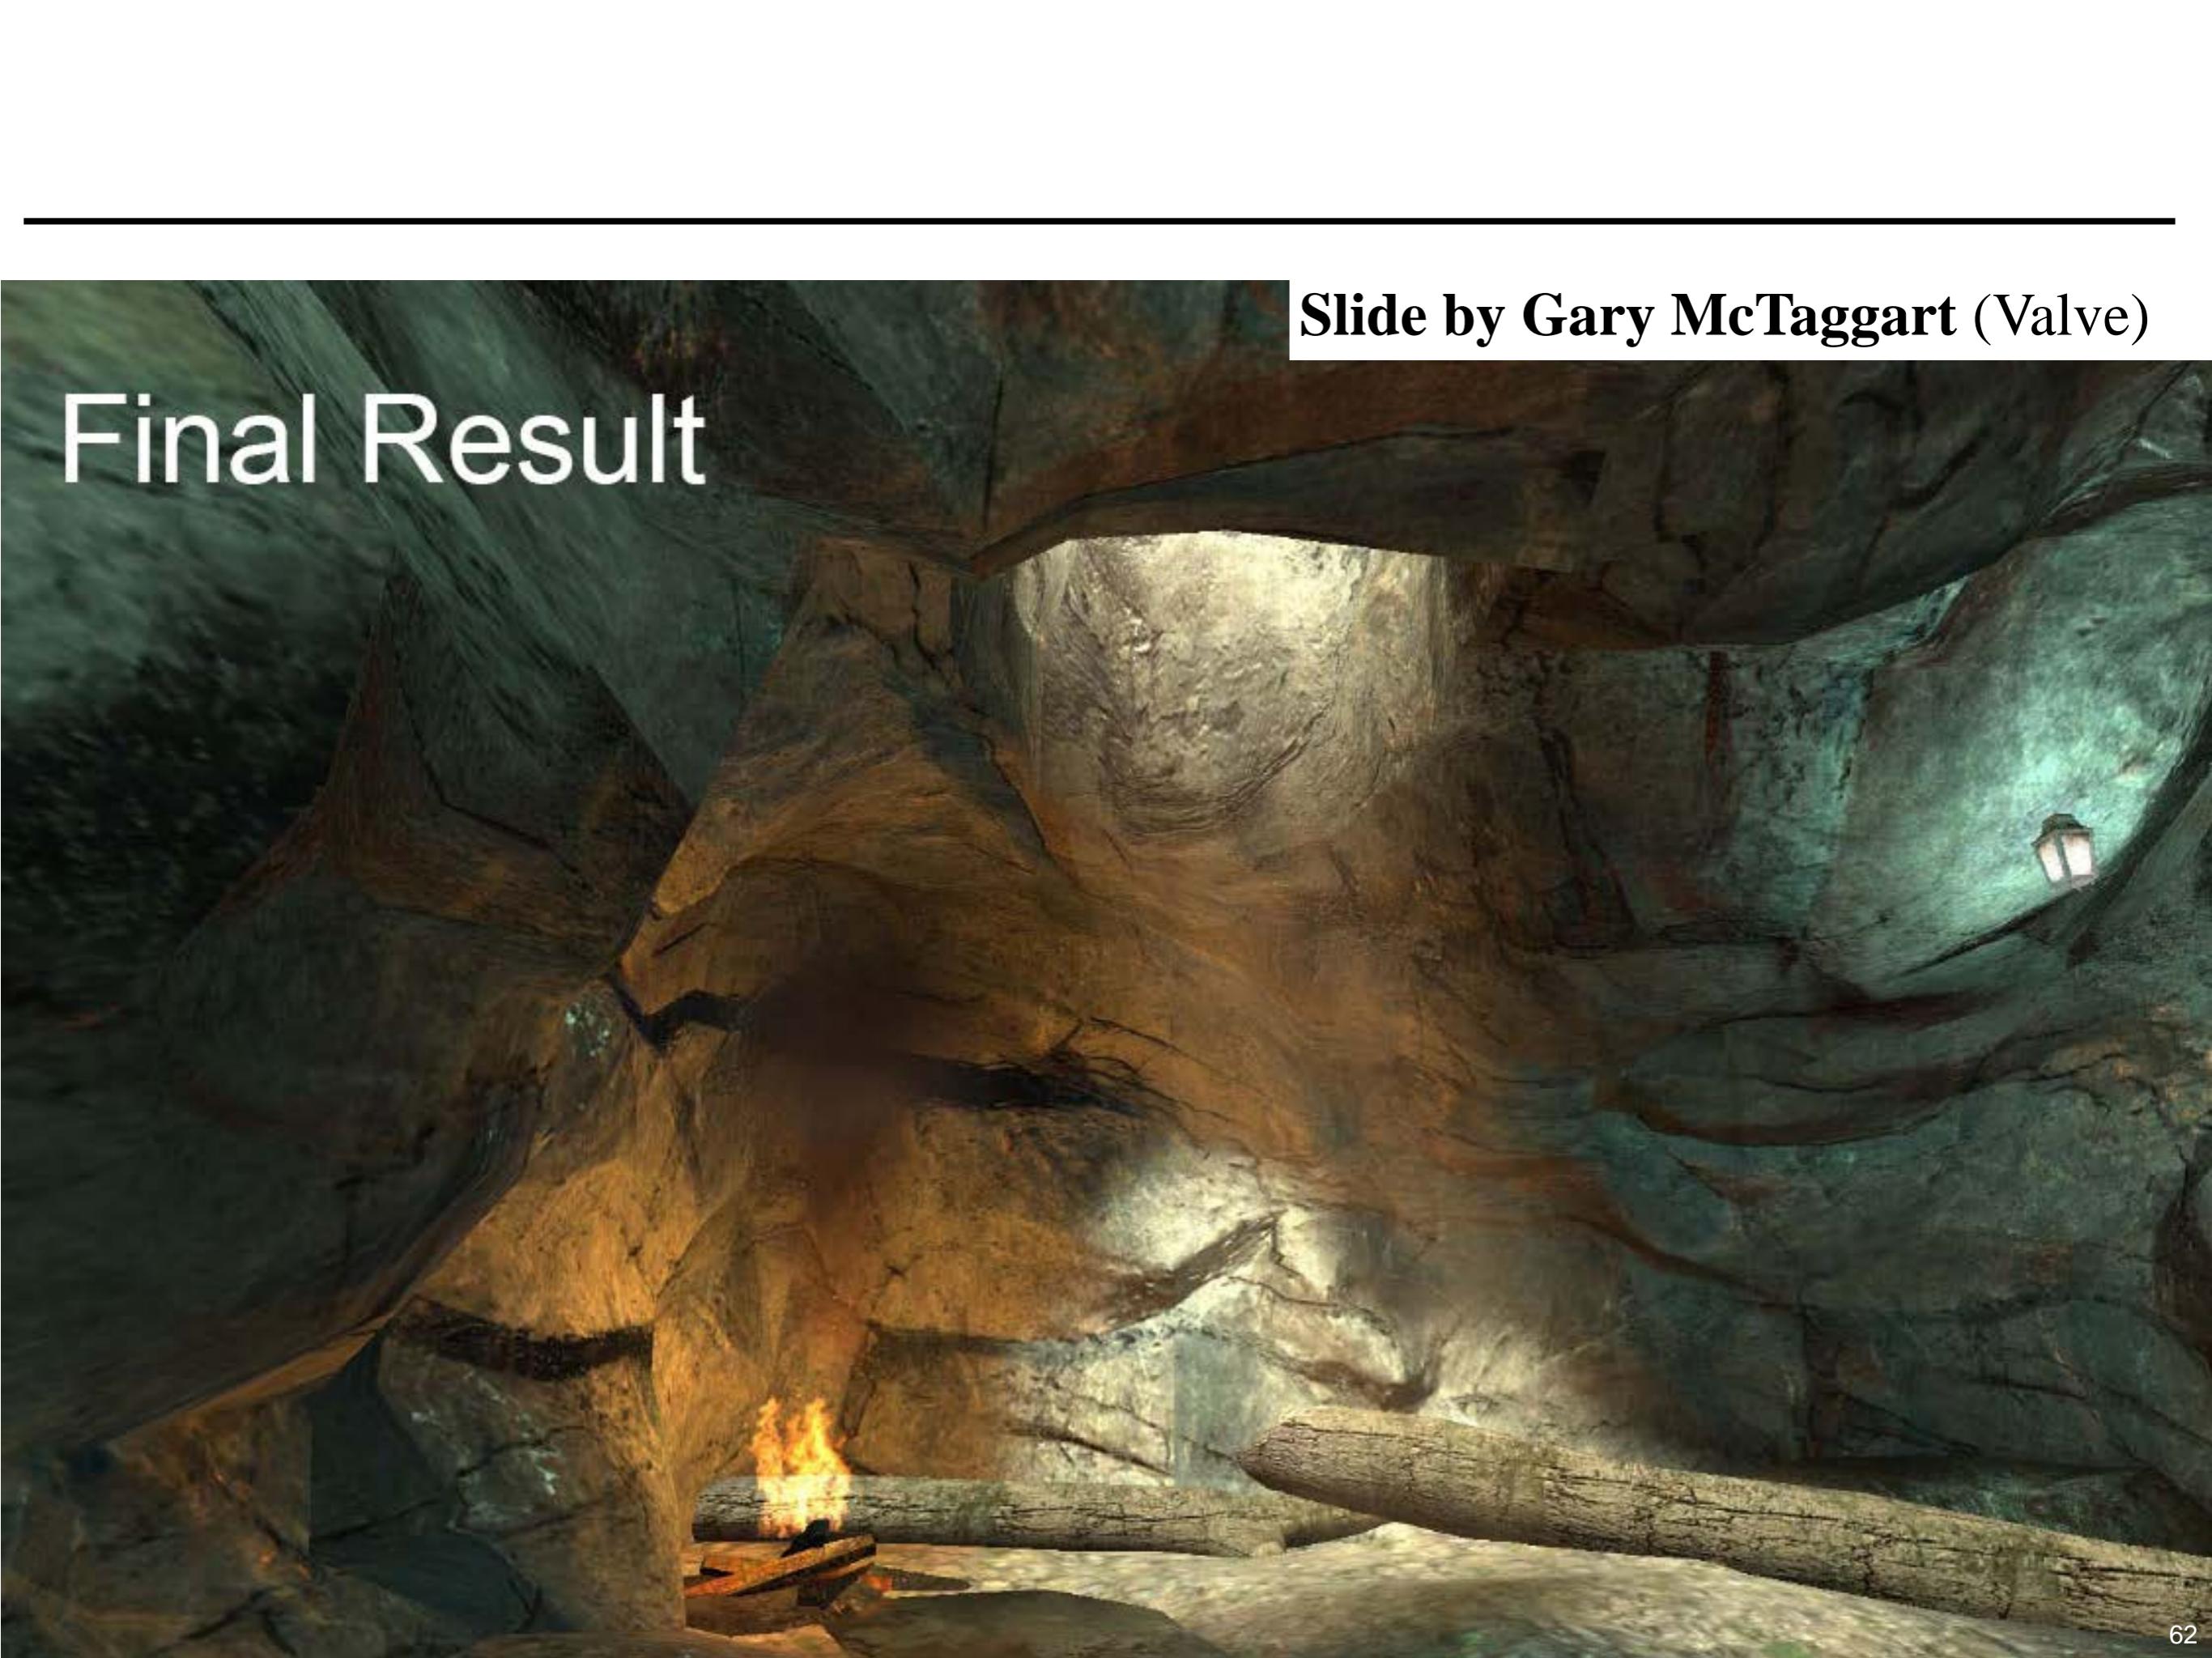
Slide by Gary McTaggart (Valve)

Normal Mapped Specular * Specular Factor

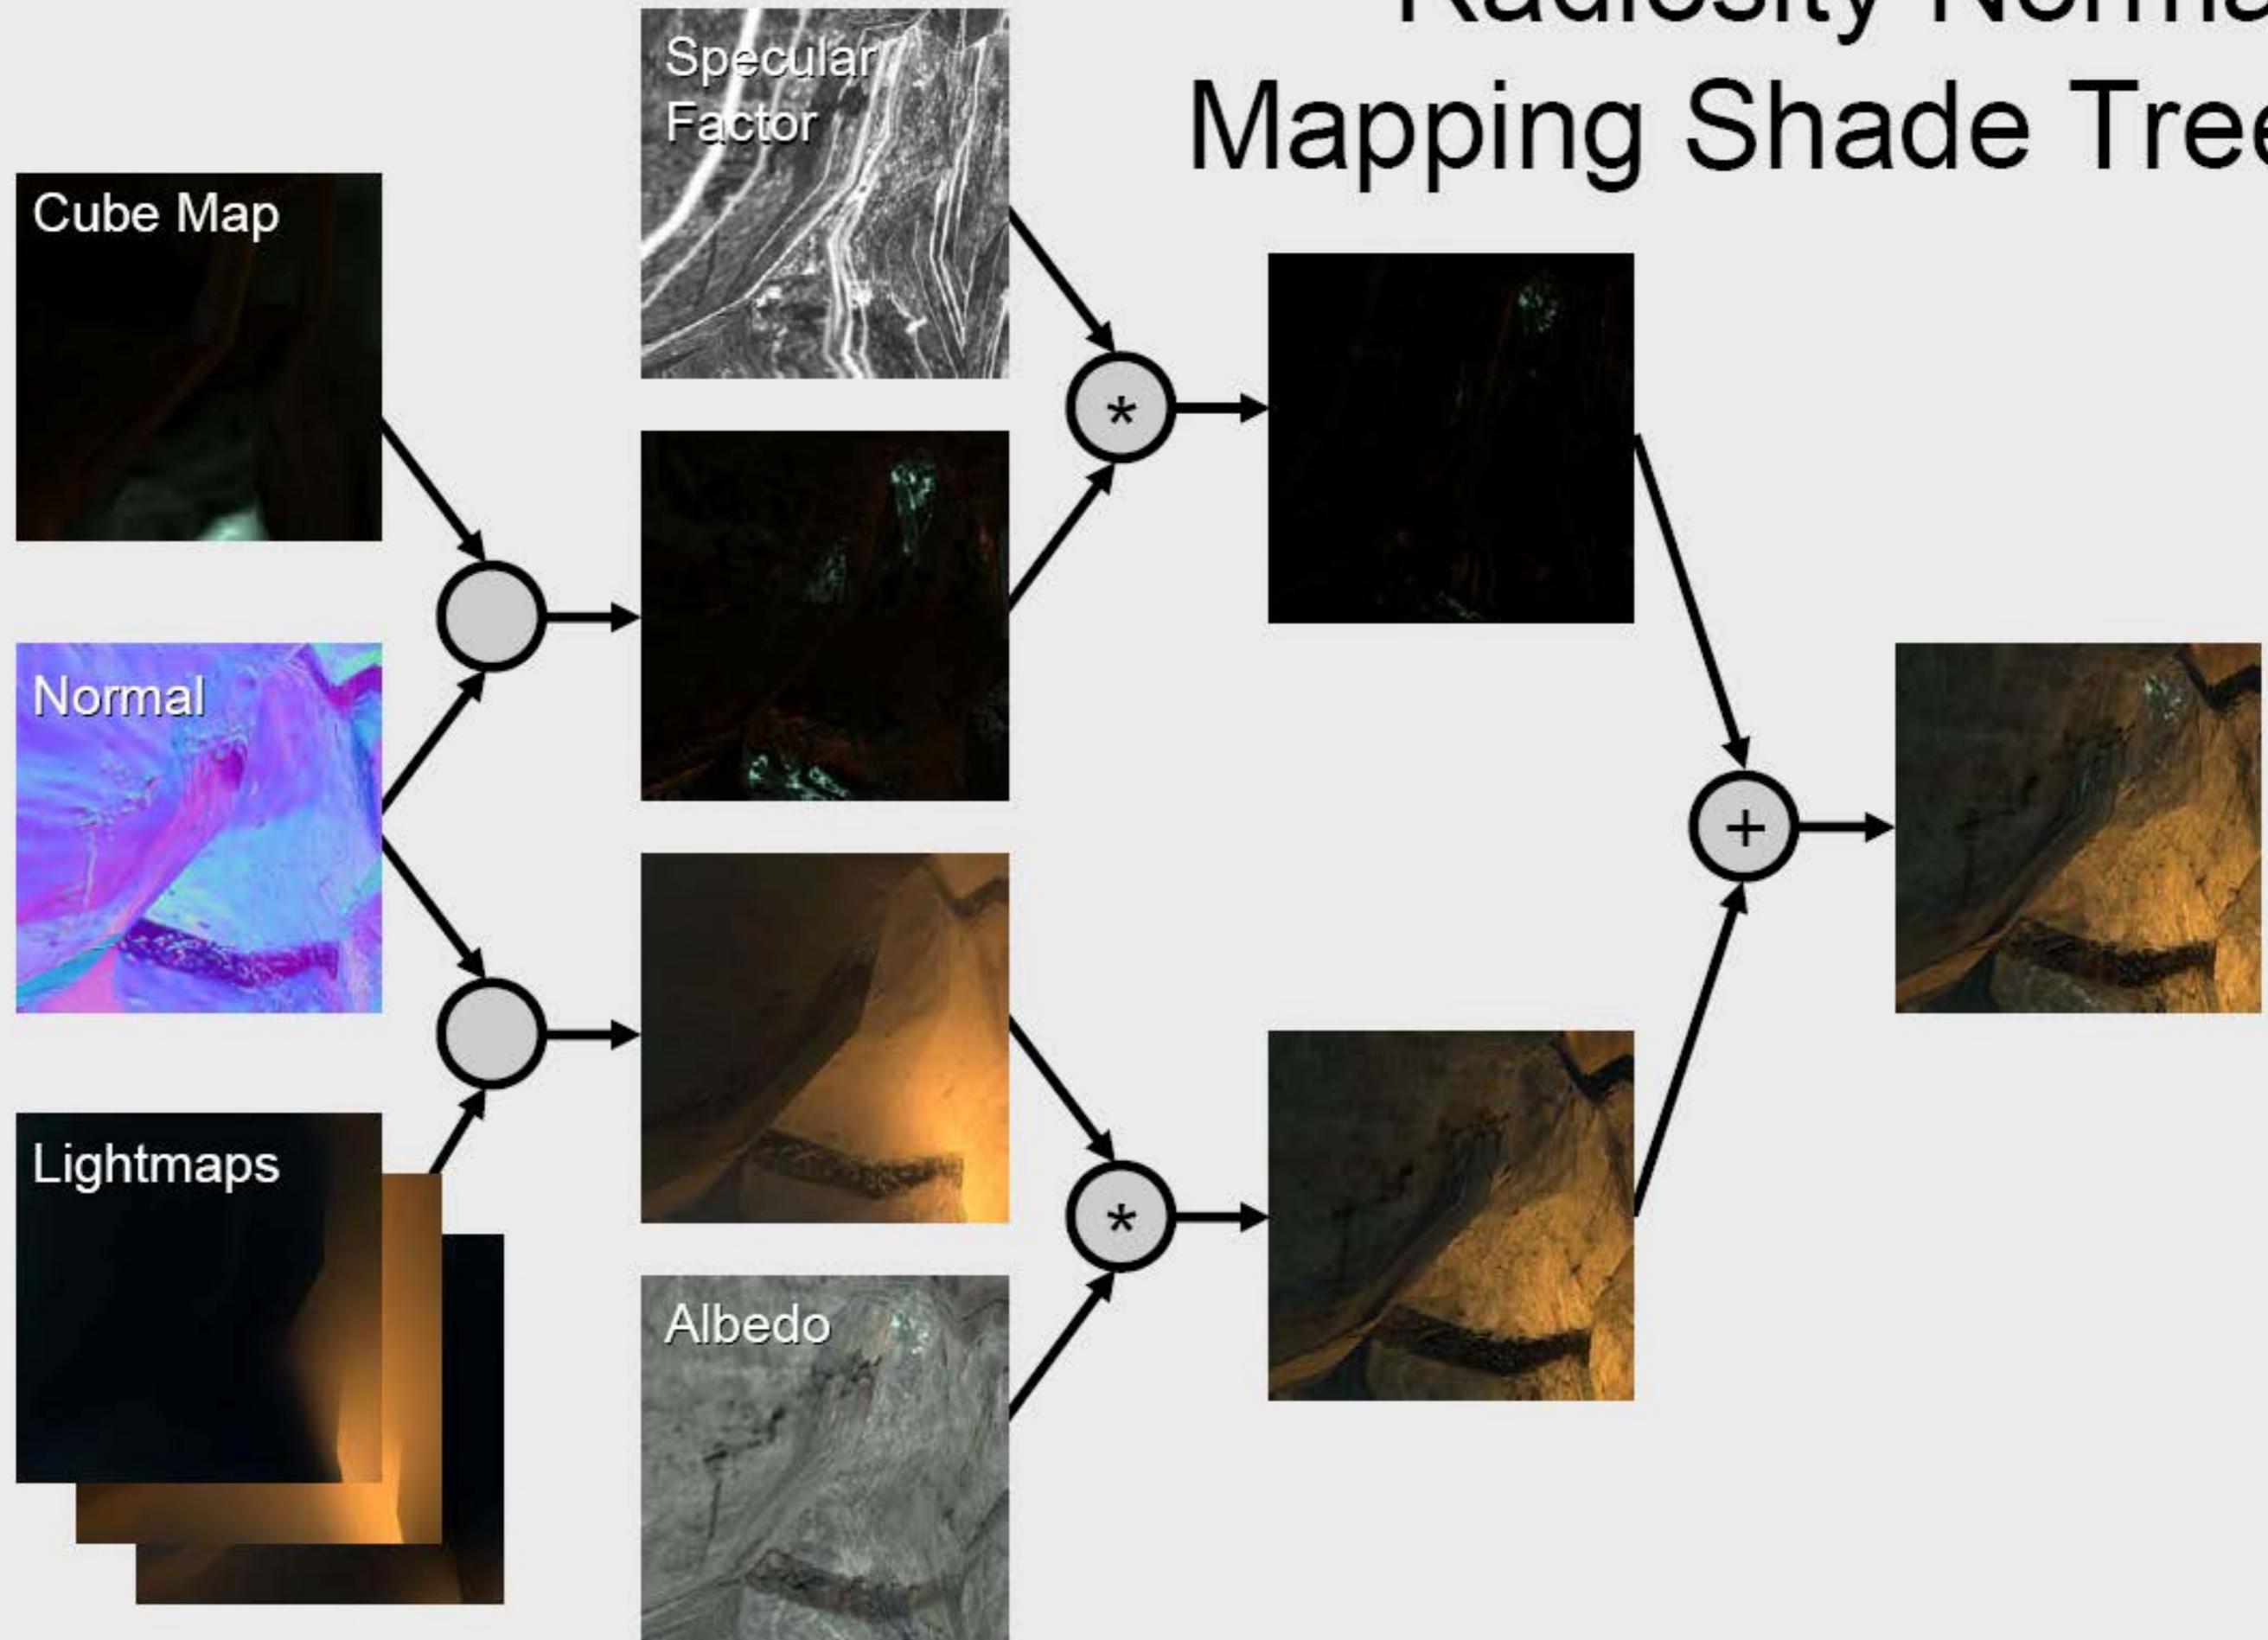


Slide by Gary McTaggart (Valve)

Final Result

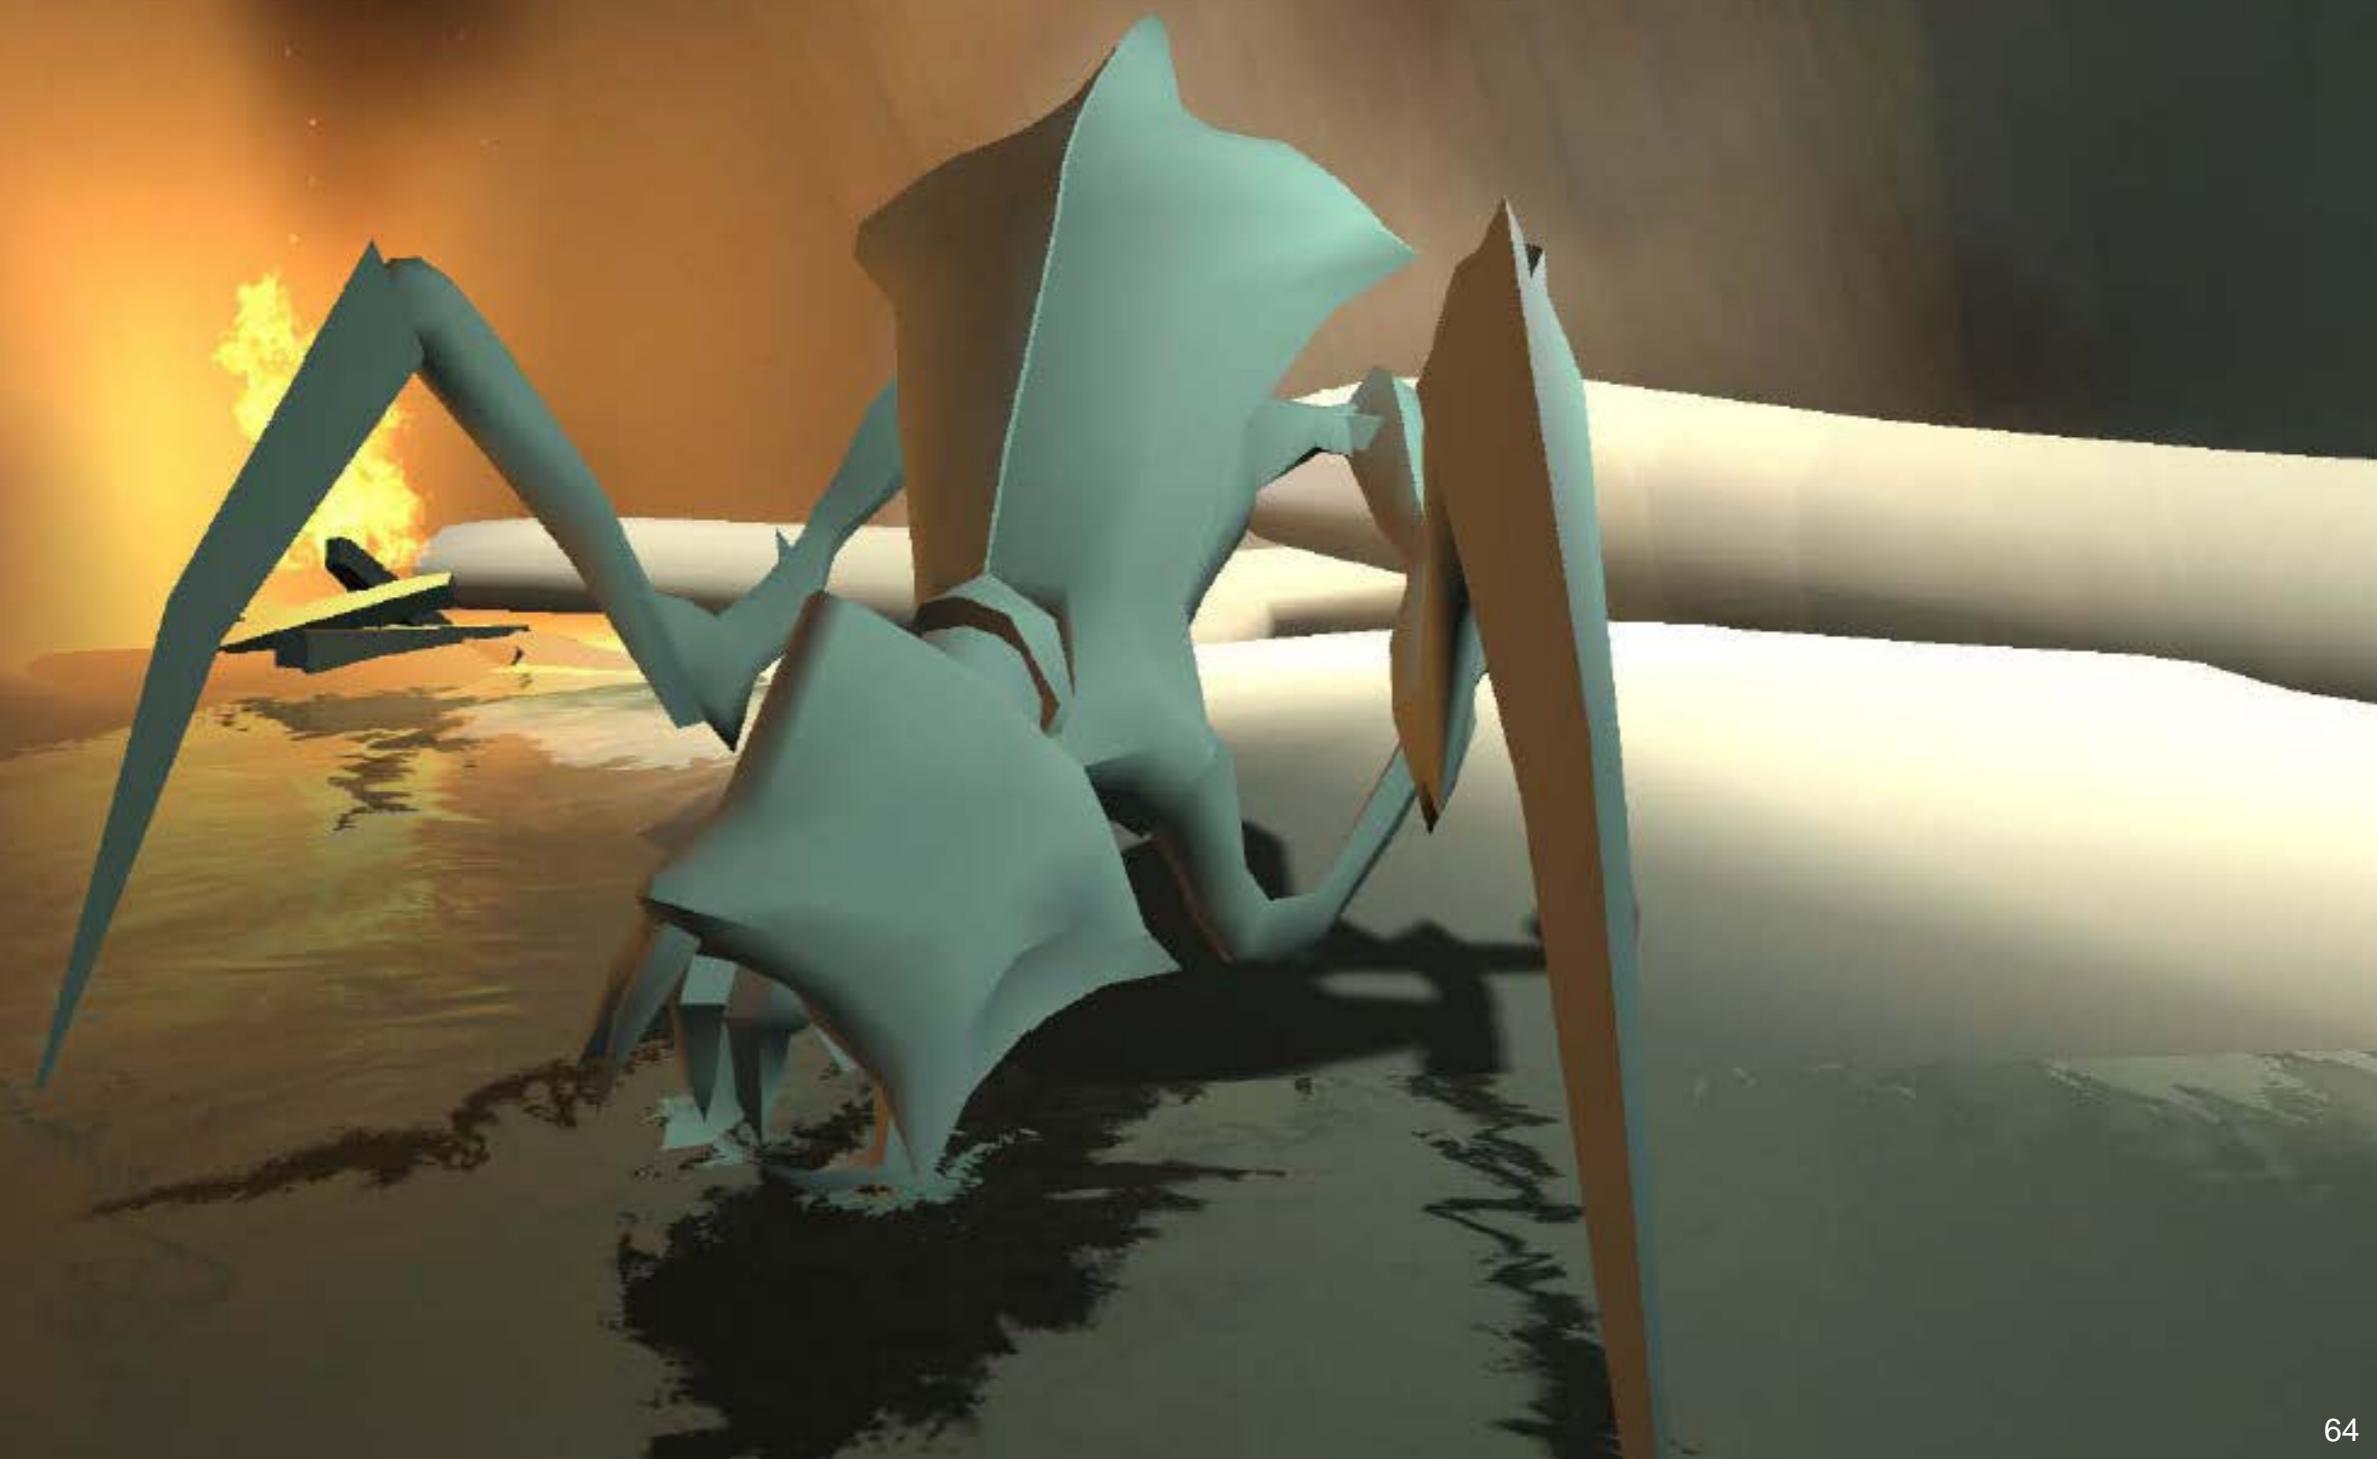


Radiosity Normal Mapping Shade Tree



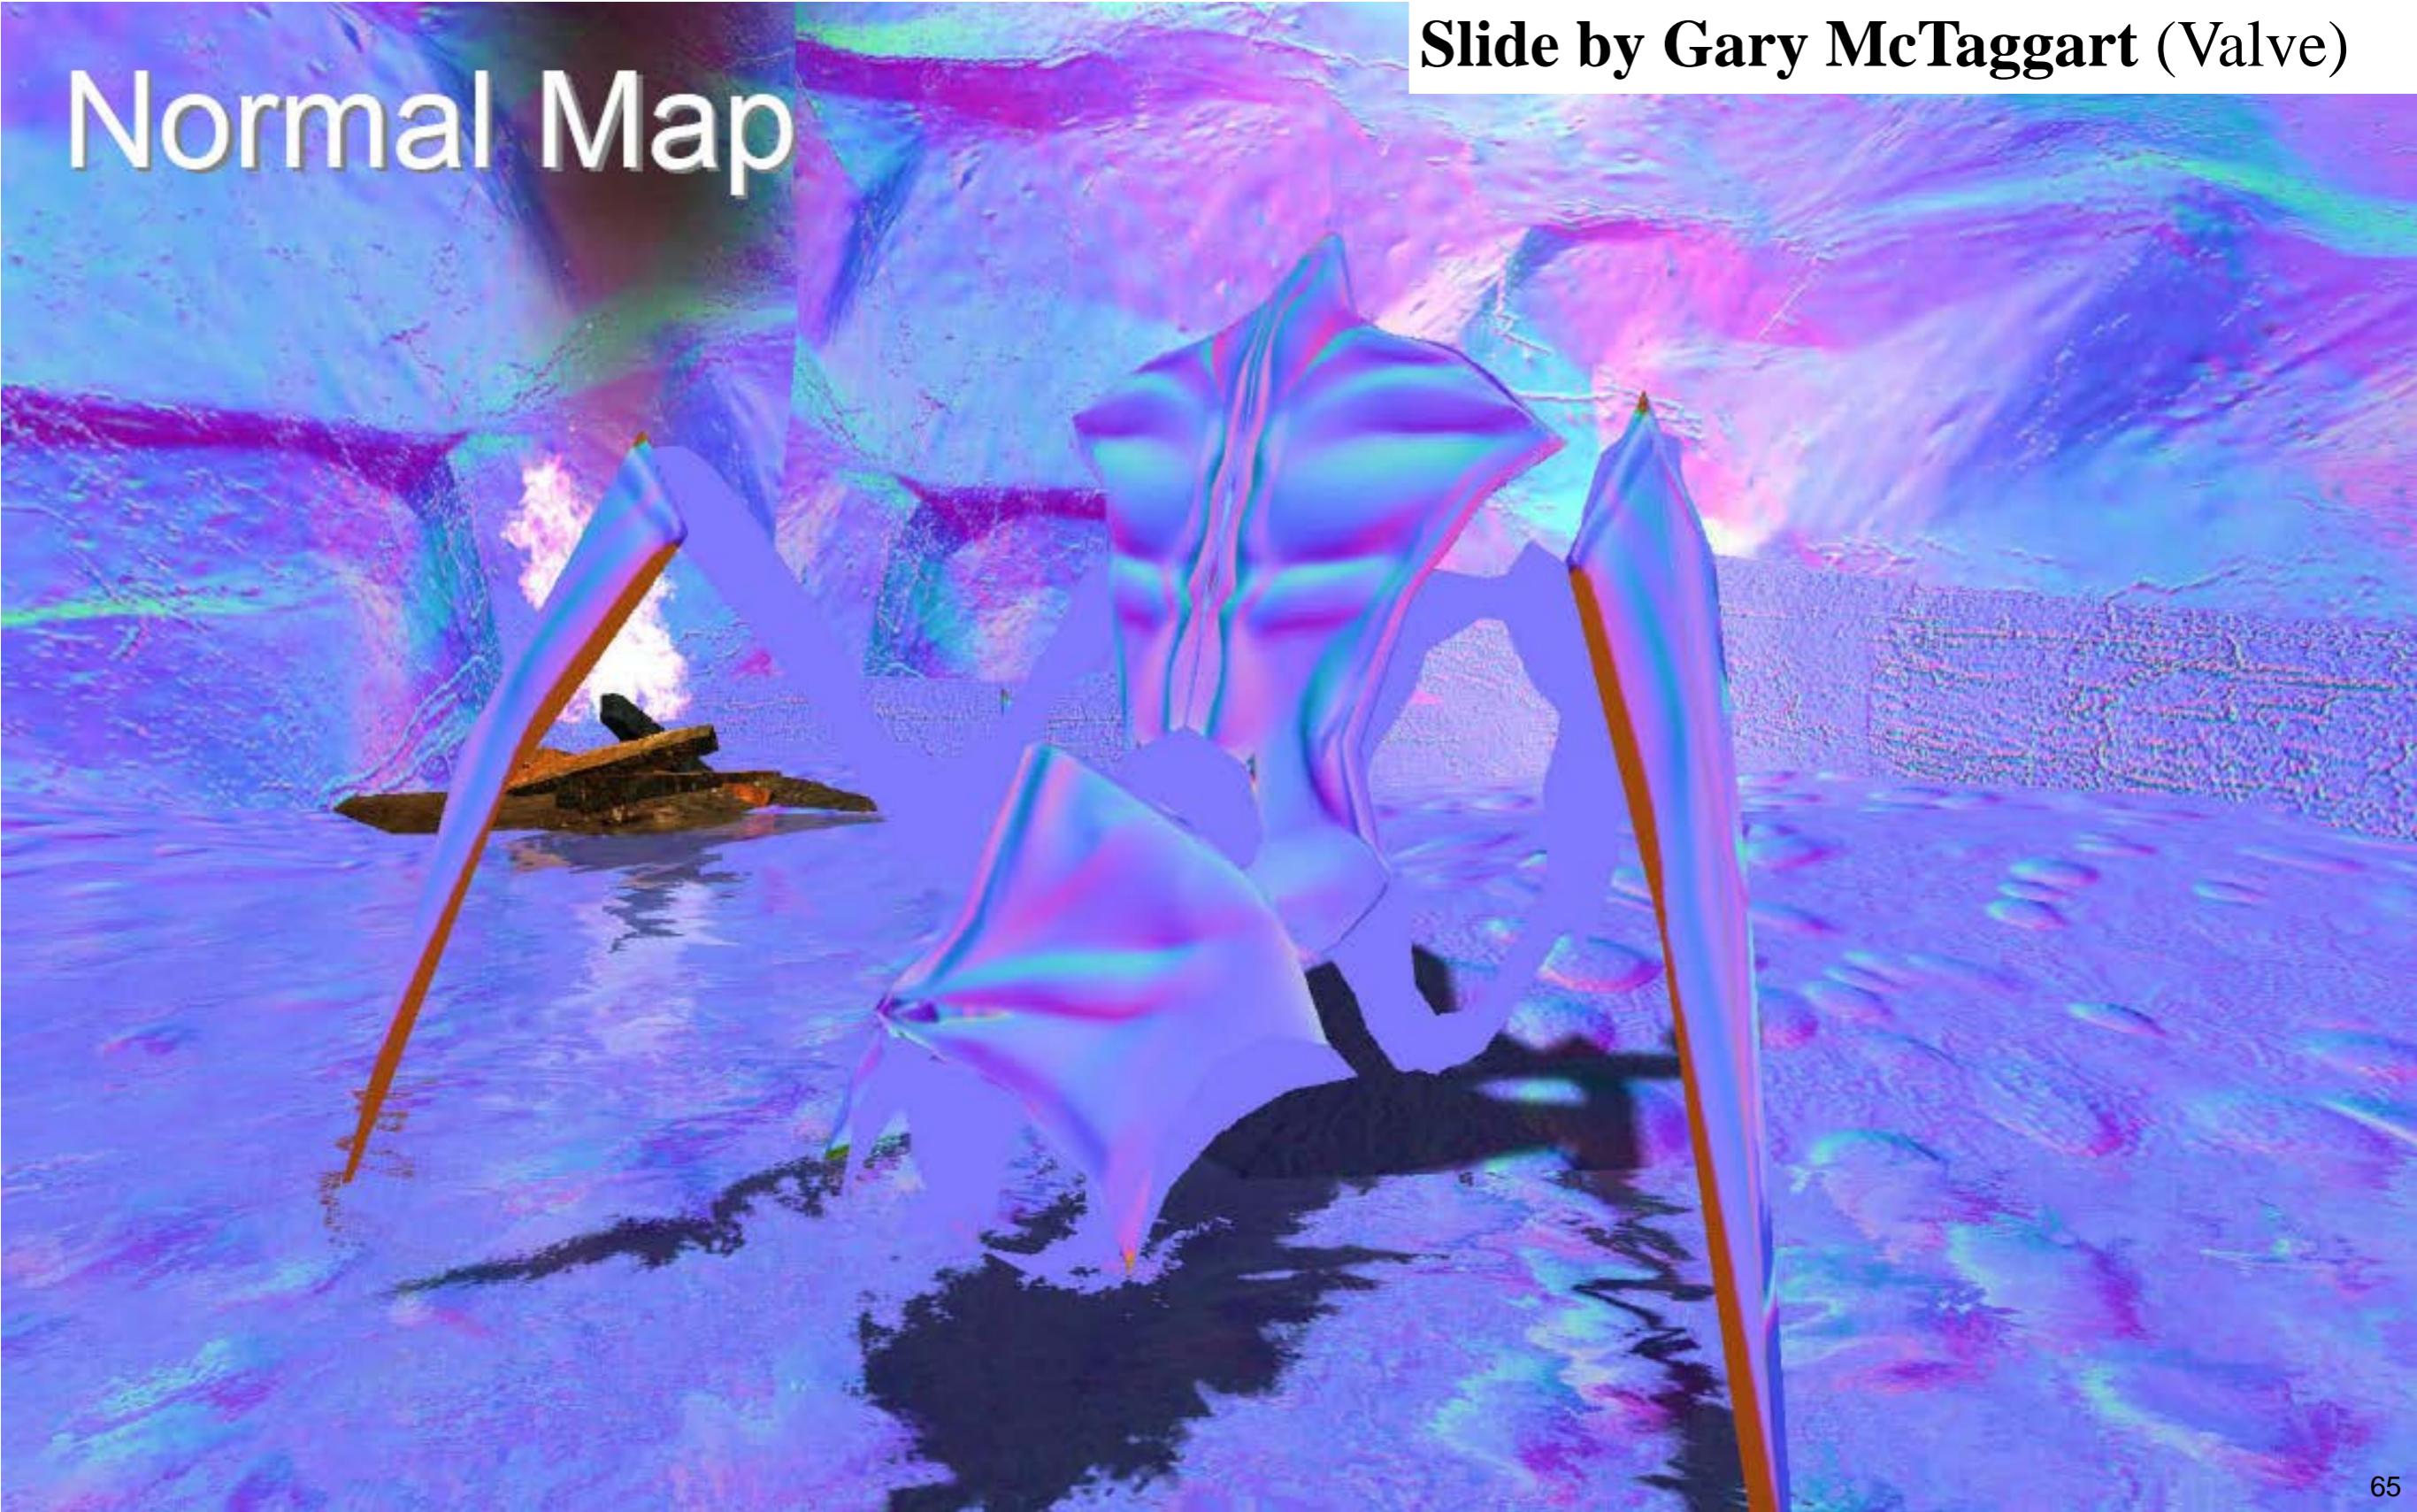
Radiosity

Slide by Gary McTaggart (Valve)



Slide by Gary McTaggart (Valve)

Normal Map



Albedo

Slide by Gary McTaggart (Valve)



Slide by Gary McTaggart (Valve)

Albedo * Normal Mapped Radiosity



Slide by Gary McTaggart (Valve)

Normal Mapped Specular



Slide by Gary McTaggart (Valve)

Specular Factor



Slide by Gary McTaggart (Valve)

Normal Mapped Specular * Specular Factor



Slide by Gary McTaggart (Valve)

Albedo * Normal Mapped Radiosity



Final Result



Slide by Gary McTaggart (Valve)

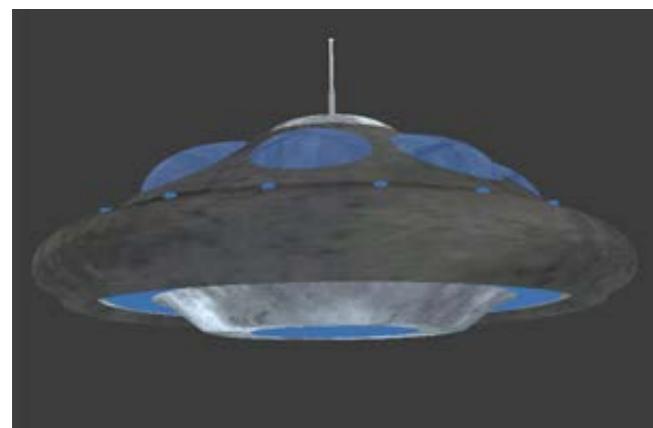
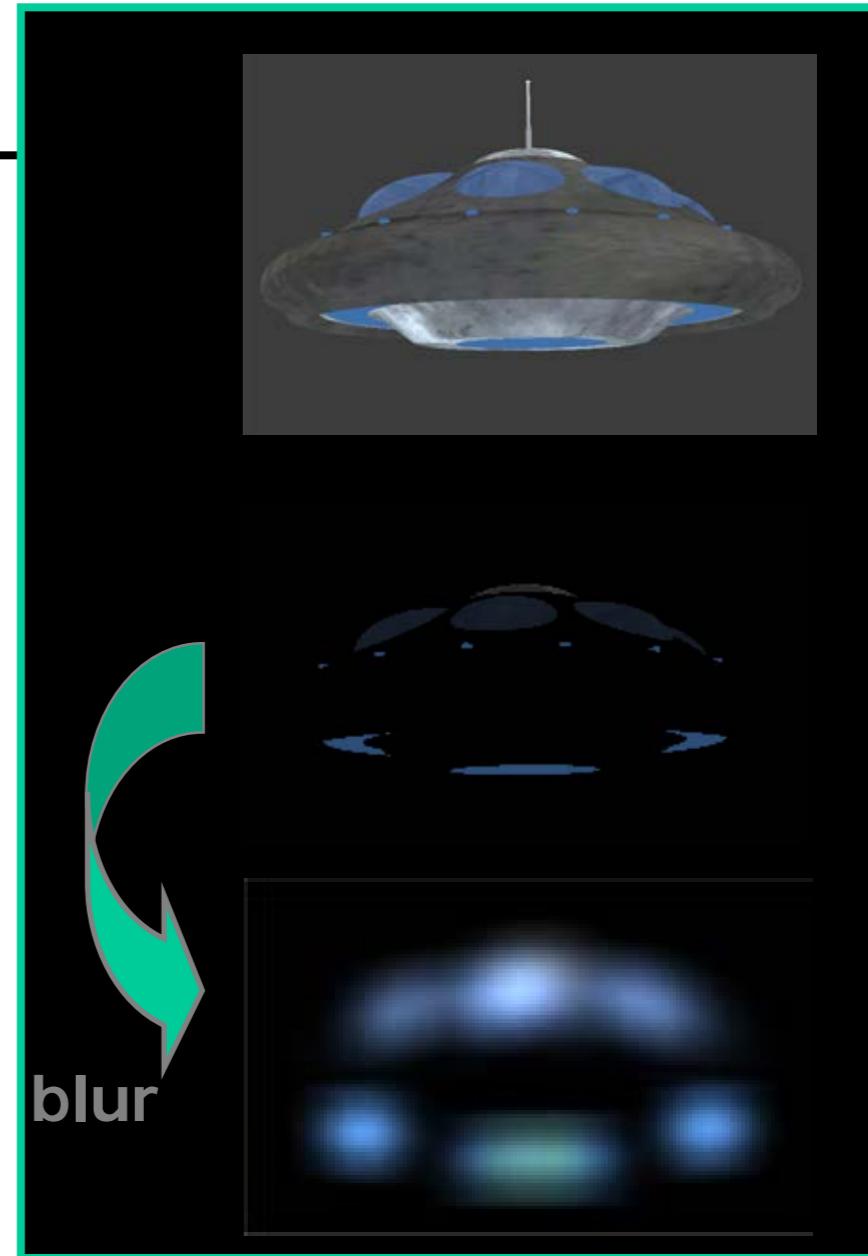
Refraction mapping (multipass)

Slide by Gary McTaggart (Valve)

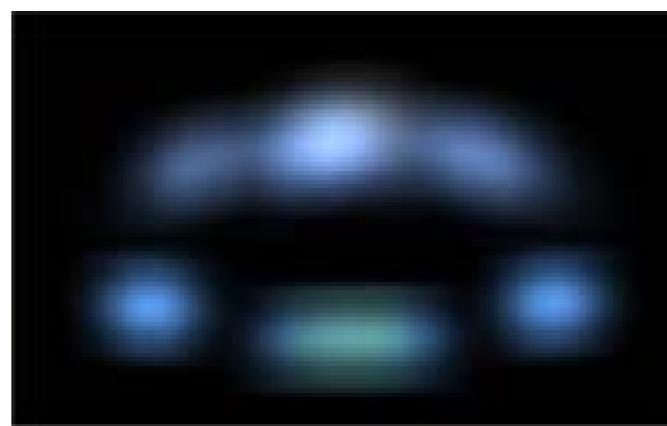


Image processing

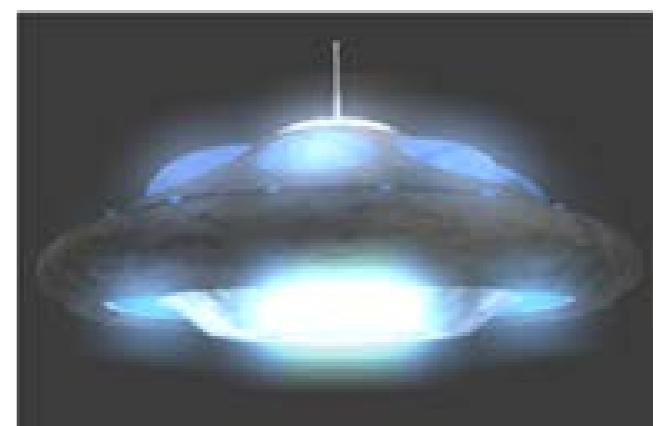
- Start with ordinary model
 - Render to backbuffer
- Render parts that are the sources of glow
 - Render to offscreen texture
- Blur the texture
- Add blur to the scene



+



=



More glow



Assets courtesy of Monolith & Disney Interactive

From “Tron”

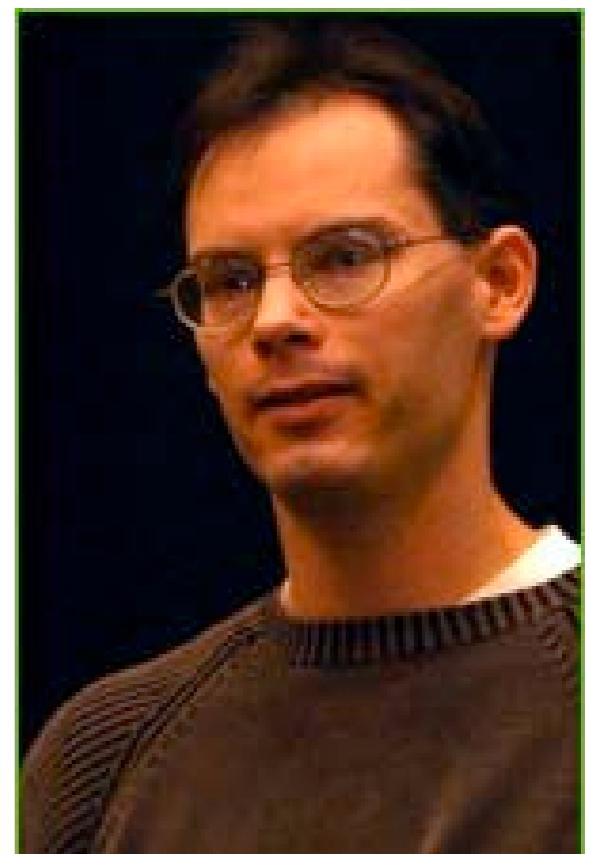
To learn more...

- <http://www.graphics.stanford.edu/courses/cs448a-01-fall/>
<https://graphics.stanford.edu/wikis/cs448-07-spring>
- <http://s09.idav.ucdavis.edu>
<http://s08.idav.ucdavis.edu>
- <http://www.beyond3d.com>
- <http://ati.amd.com/developer/techpapers.html>
- <http://developer.nvidia.com>
http://download.nvidia.com/developer/SDK/Individual_Samples/samples.html
http://download.nvidia.com/developer/SDK/Individual_Samples/effects.html
<http://developer.nvidia.com/page/tools.html>

Game Development Outline

- Game Development
 - Typical Process
- What's in a game?
 - Game Simulation
 - Numeric Computation
 - Shading

Tim Sweeney



<http://www.cs.princeton.edu/~dpw/popl/06/Tim-POPL.ppt>

Game Development: Gears of War

- Resources
 - ~10 programmers
 - ~20 artists
 - ~24 month development cycle
 - ~\$10M budget
- Software Dependencies
 - 1 middleware game engine
 - ~20 middleware libraries
 - OS graphics APIs, sound, input, etc

Software Dependencies

Gears of War
Gameplay Code
~250,000 lines C++, script code

Unreal Engine 3
Middleware Game Engine
~250,000 lines C++ code

DirectX
Graphics

OpenAL
Audio

Ogg
Vorbis
Music
Codec

Speex
Speech
Codec

wx
Widgets
Window
Library

ZLib
Data
Compr-
ession

...

Game Development – Platforms

- The typical Unreal Engine 3 game will ship on:
 - Xbox 360
 - PlayStation 3
 - Windows
- Some will also ship on:
 - Linux
 - MacOS

What is in a game?

The obvious:

- Rendering
- Pixel shading
- Physics simulation, collision detection
- Game world simulation
- Artificial intelligence, path finding

But it's not just fun and games:

- Data persistence with versioning, streaming
- Distributed Computing (multiplayer game simulation)
- Visual content authoring tools
- Scripting and compiler technology
- User interfaces

Three Kinds of Code

- Gameplay Simulation
- Numeric Computation
- Shading

Gameplay Simulation

- Models the state of the game world as interacting objects evolve over time
- High-level, object-oriented code
- Written in C++ or scripting language
- Usually garbage-collected

Gameplay Simulation – The Numbers

- 30-60 updates (frames) per second
- ~1000 distinct gameplay classes
 - Contain imperative state
 - Contain member functions
 - Highly dynamic
- ~10,000 active gameplay objects
- Each time a gameplay object is updated, it typically touches 5-10 other objects

Numeric Computation

- Algorithms:
 - Scene graph traversal
 - Physics simulation
 - Collision Detection
 - Path Finding
 - Sound Propagation
- Low-level, high-performance code
- Written in C++ with SIMD intrinsics
- Essentially functional
 - Transforms a small input data set to a small output data set, making use of large constant data structures.

Shading

- Generates pixel and vertex attributes
- Written in HLSL/CG shading language
- Runs on the GPU
- Inherently data-parallel
 - Control flow is statically known
 - “Embarassingly Parallel”

Shading – The Numbers

- Game runs at 30 FPS @ 1280x720p
- ~5,000 visible objects
- ~10M pixels rendered per frame
 - Per-pixel lighting and shadowing requires multiple rendering passes per object and per-light
- Typical pixel shader is ~100 instructions long
- ~500 GFLOPS compute power

Three Kinds of Code

	Game Simulation	Numeric Computation	Shading
Languages	C++, Scripting	C++	CG, HLSL
CPU Budget	10%	90%	n/a
Lines of Code	250,000	250,000	10,000
FPU Usage	0.5 GFLOPS	5 GFLOPS	500 GFLOPS

Gears of war 3

- (Unreal engine)
- Smart conservative collision detection
- Blades of grass using vertical textures
- Anisotropic mipmapping not bad
- Physics
- Visibility culling
- Ambient occlusion
- Softbody dynamic, fracture



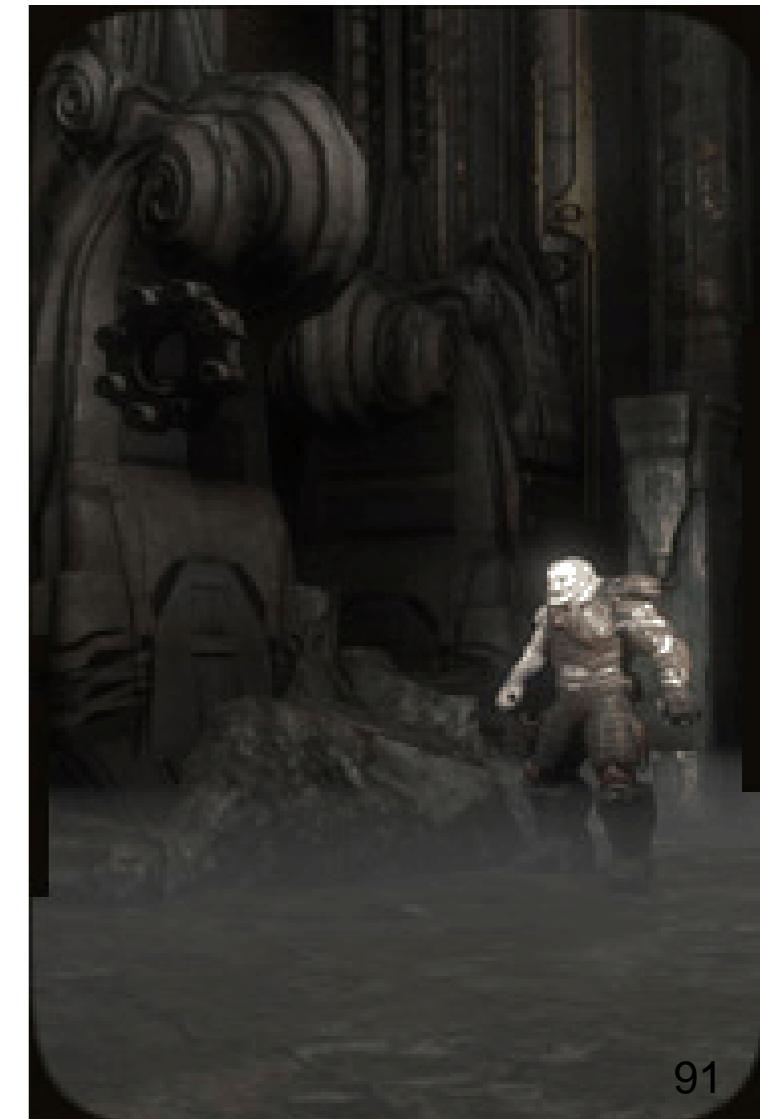
Unreal Engine

- <http://www.unrealengine.com>
- Mostly for First Person Shooters
- Animation, rendering, sound, etc.



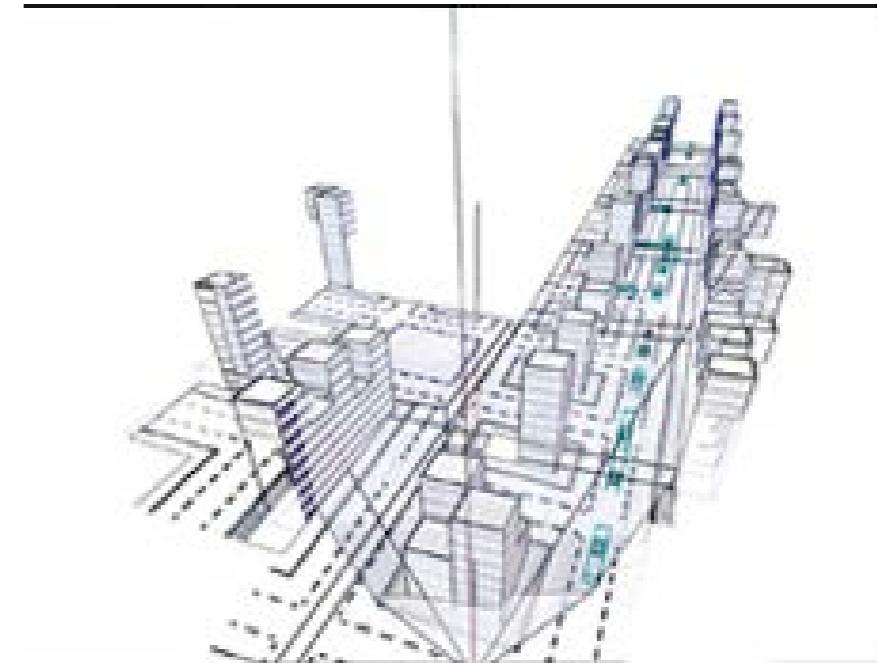
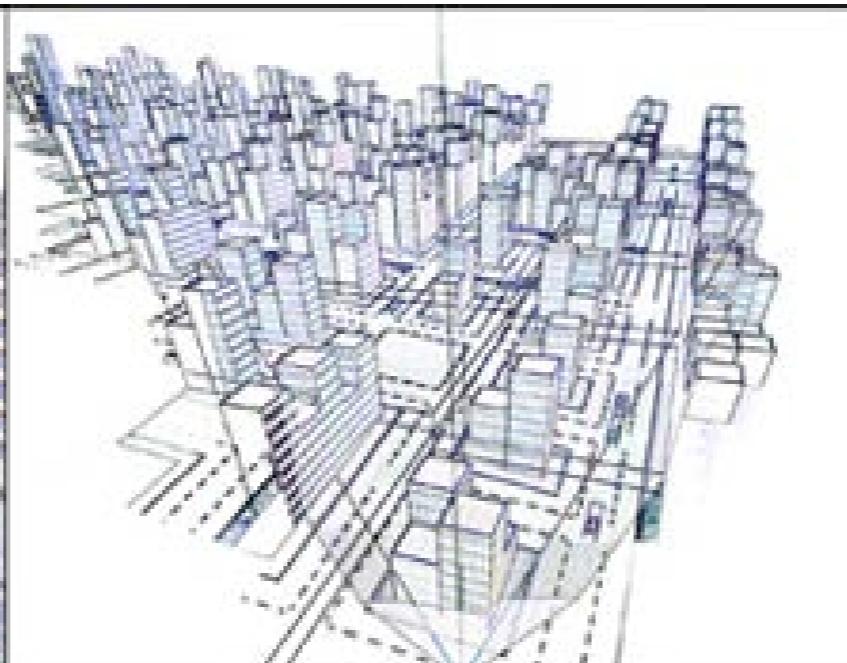
Rendering in Unreal Engine 3

- <http://www.unrealtechnology.com/features.php?ref=rendering>
- Multithreading
- Visibility culling
- High Dynamic Range (more than 0-255)
- Pixel shaders: normal map, Phong; anisotropic effects; displacement mapping; etc.
- Platform abstraction layer
- Lighting with spherical harmonics
- Shadow volume, shadow map, PCF
- Volumetric effects (fog)
- texture streaming
- reflection with dynamic environment maps
- Image post-processing: bloom, tone mapping, etc.
- Lots of shader support



Umbra's occlusion culling

- <http://www.umbrasoftware.com/index.php?products&umbra>



A view of a 3d-scene

The same view seen from higher perspective.

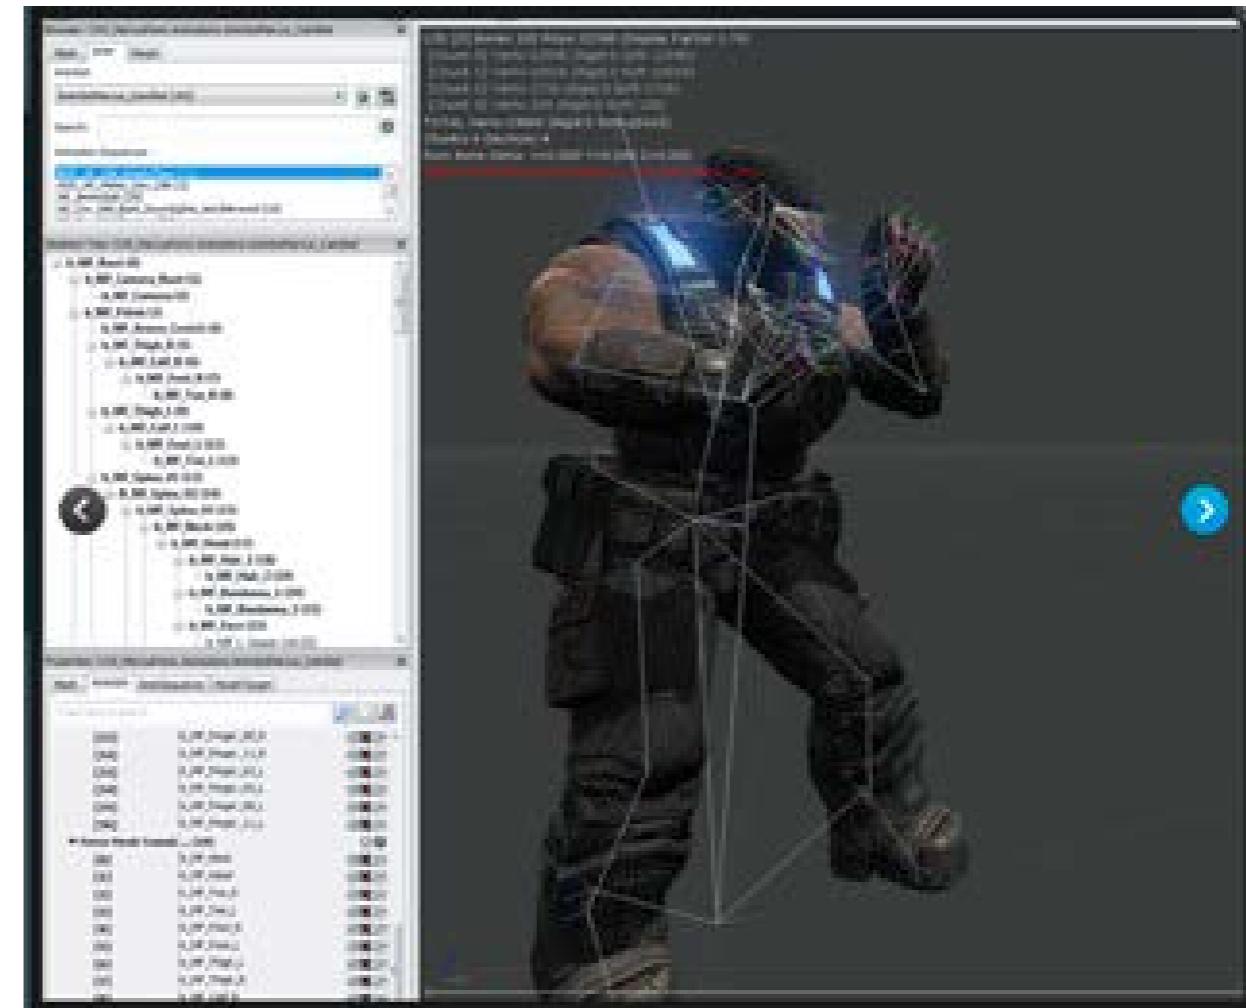
The same scene with Umbra culling off the nonvisible objects. And it works real time!

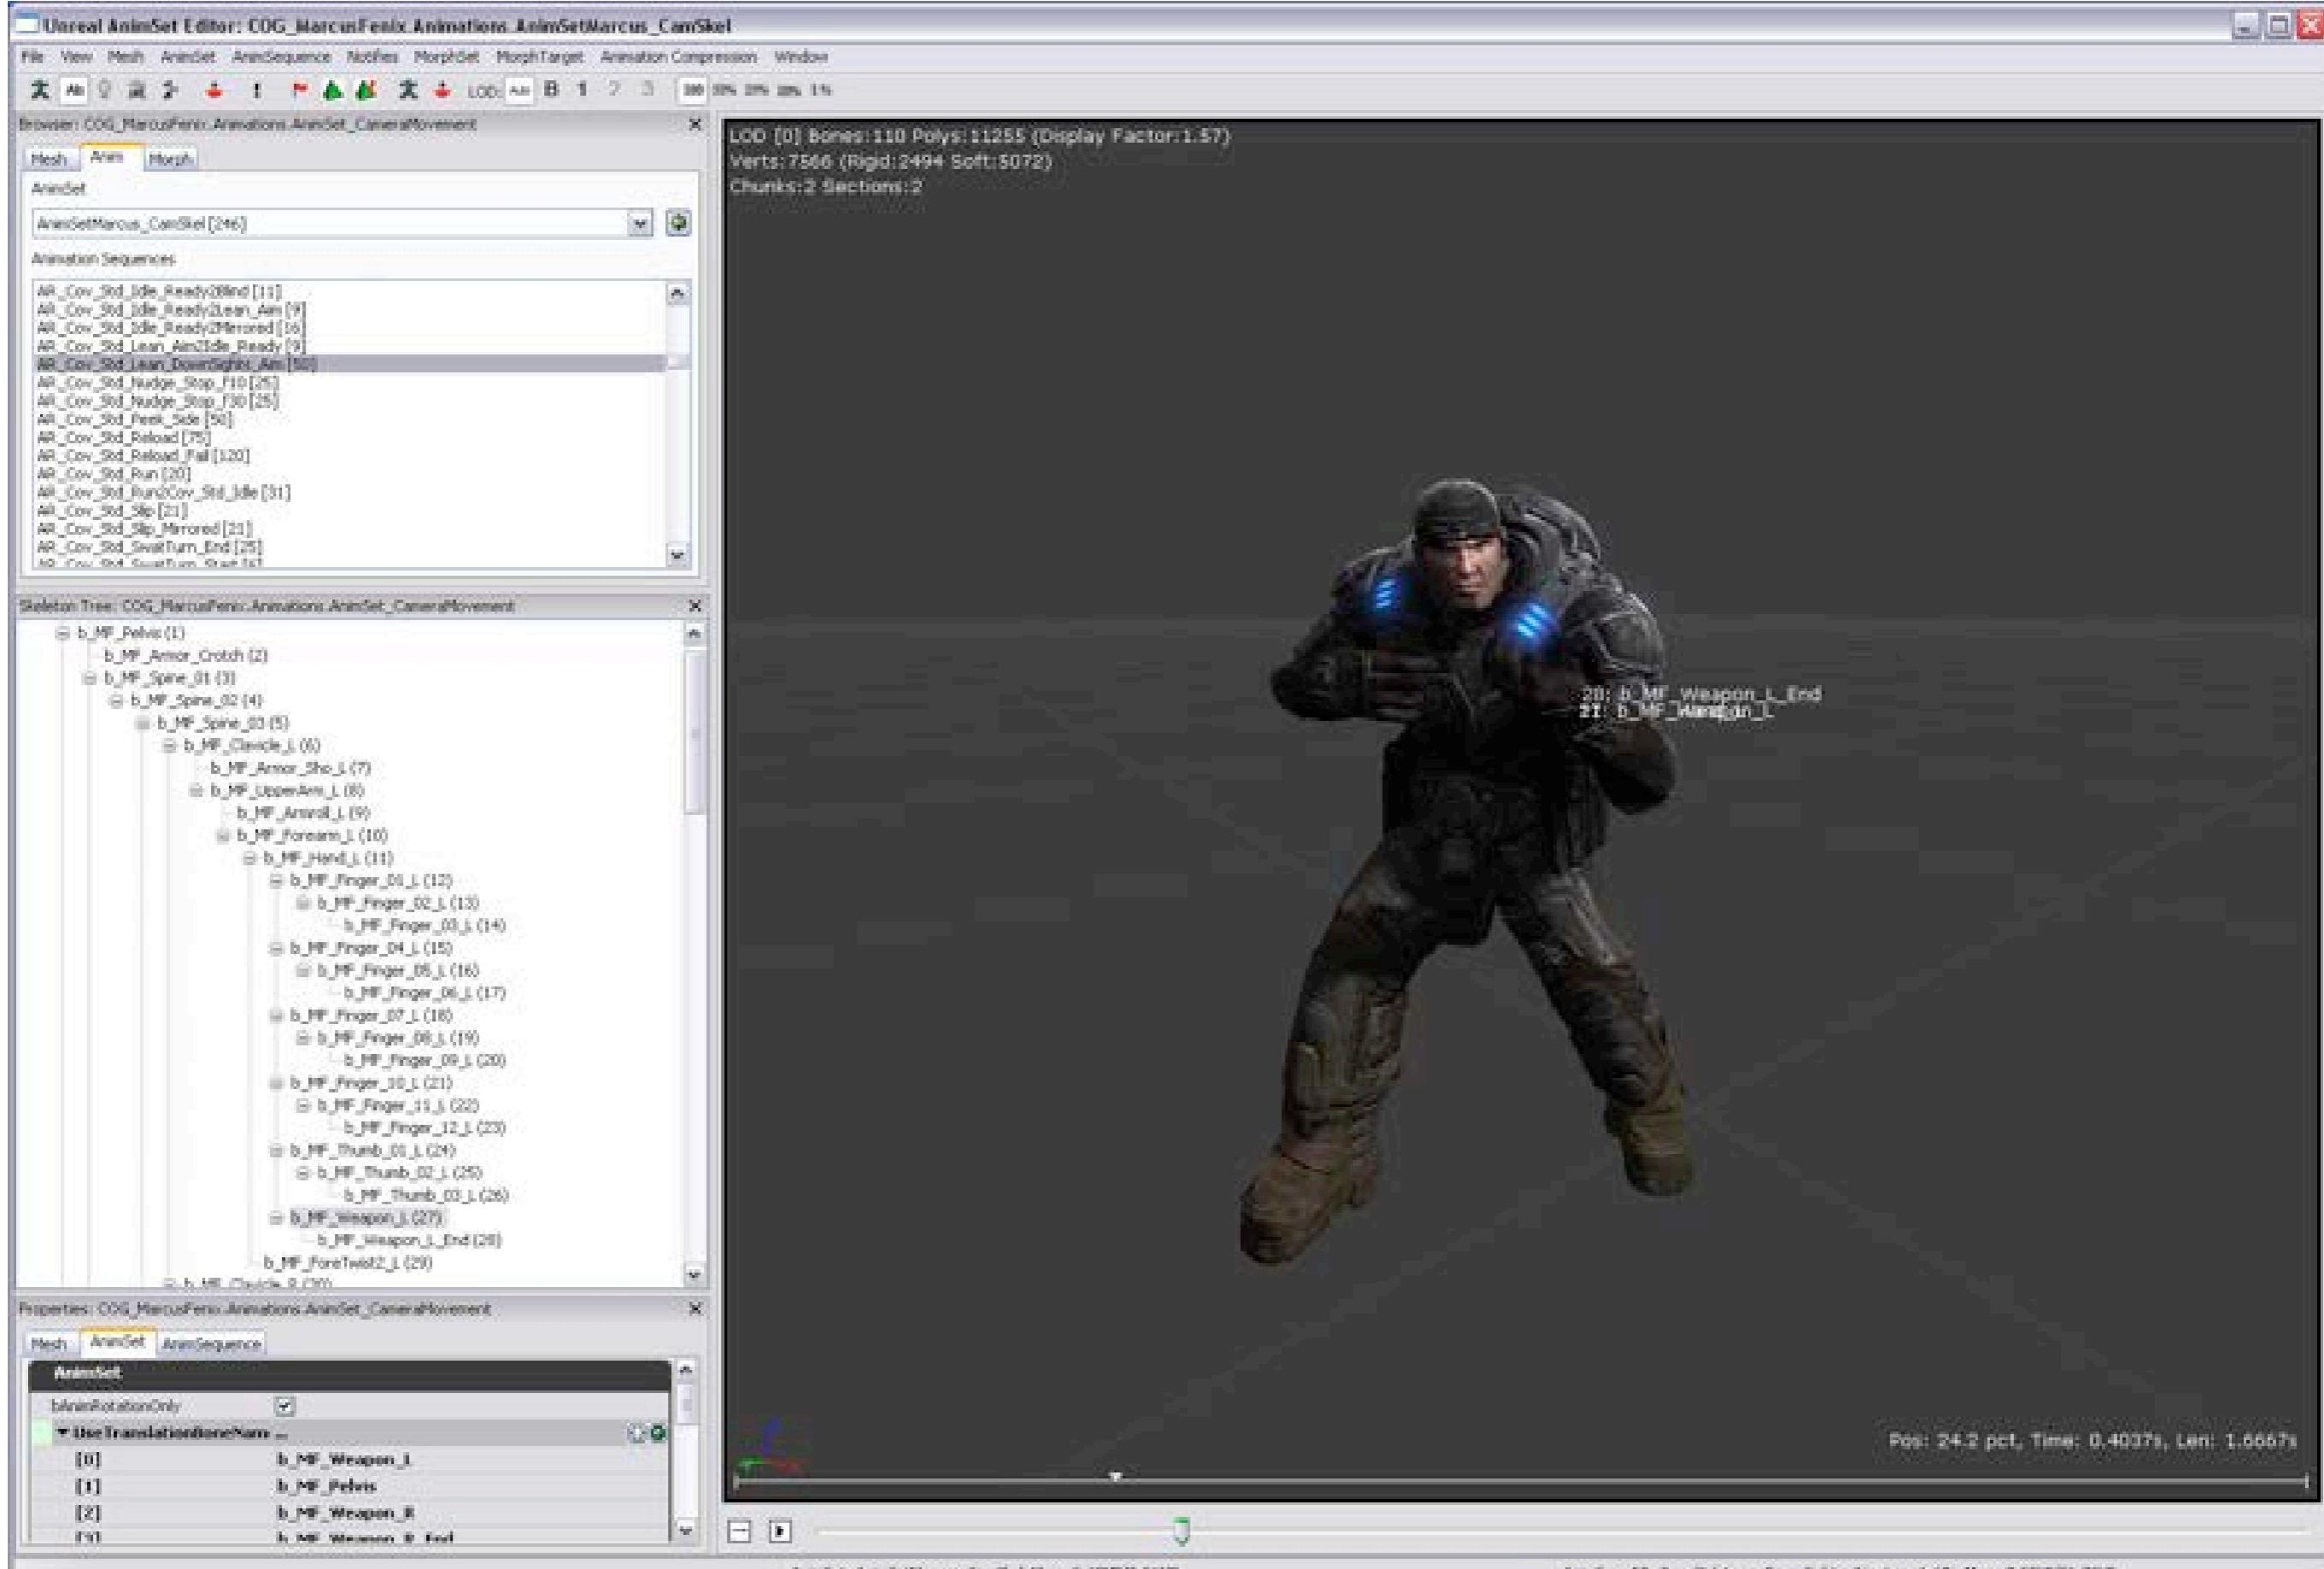


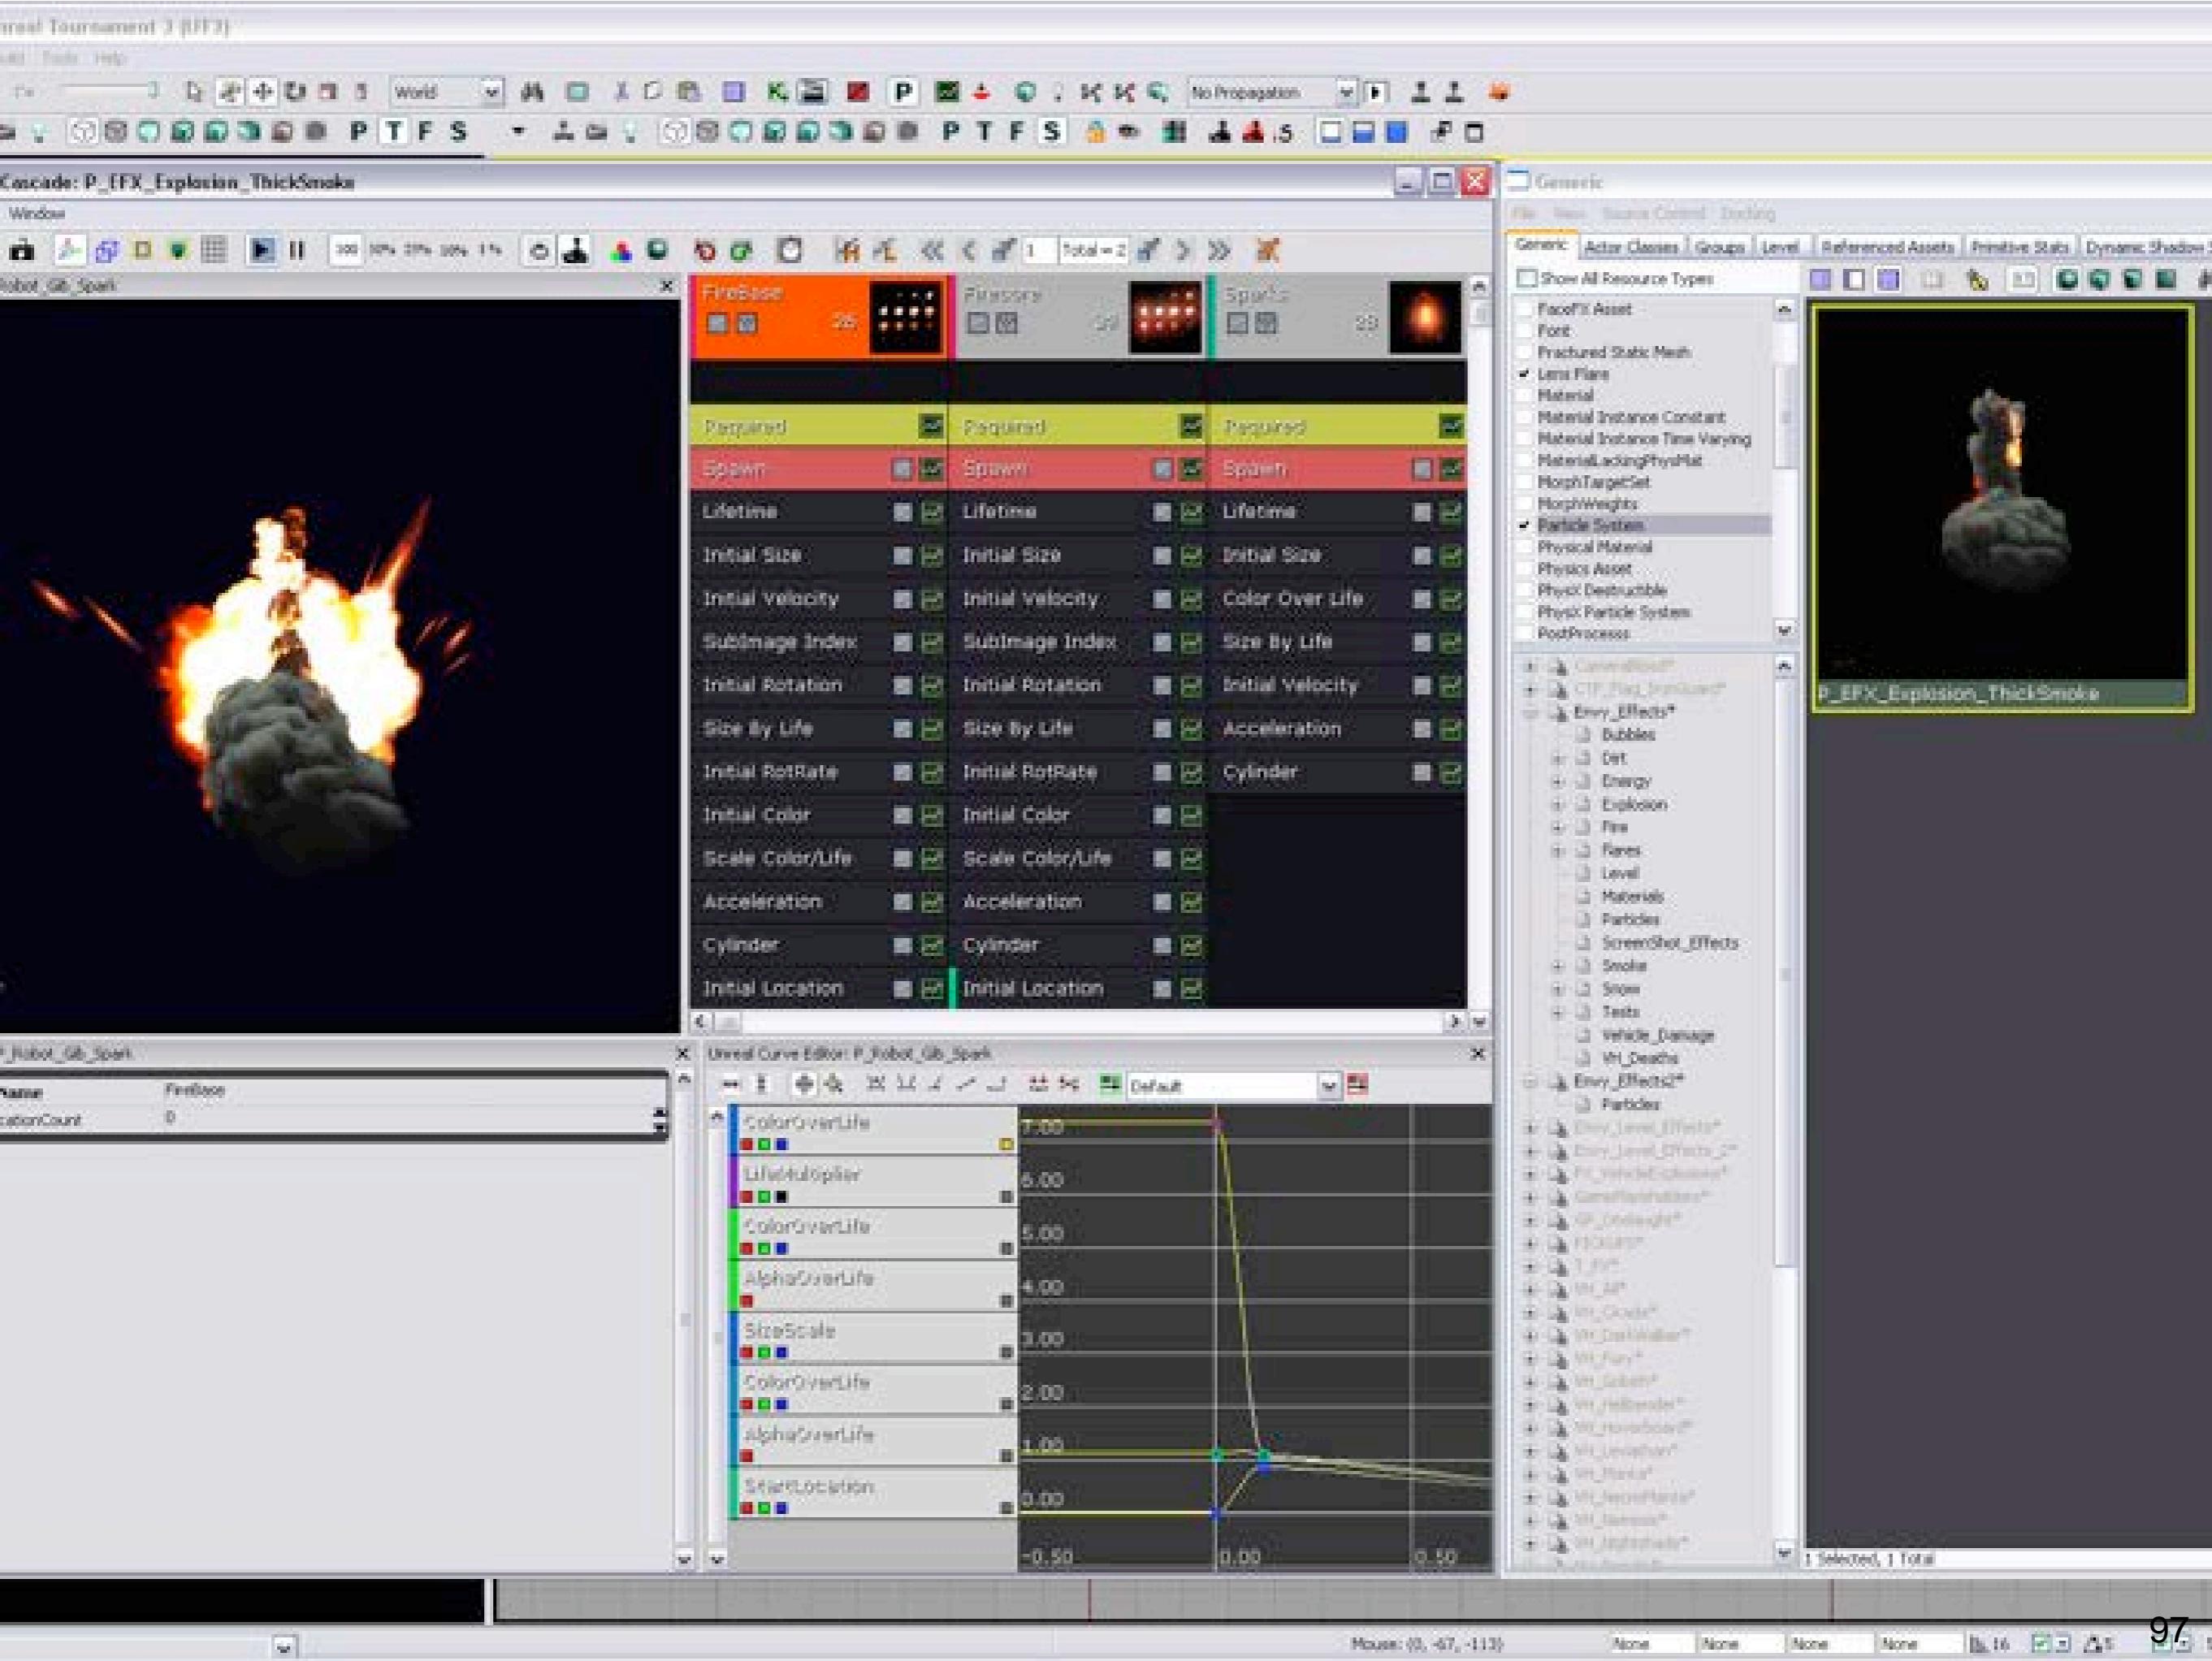


Animation

- <http://www.unrealengine.com/features/animation/>
- Skinned mesh rendering (up to 4 bones / vertex)
- Mesh and bone Level of Detail
- Blend shapes
- Keyframed animation
- Particle system
- Rigid-body physics
- Cloth simulation
- Soft-body physics
- Fracture
- Collision detection







See also Valve source engine

- Slides from Jason Mitchell
- <http://www.pixelmaven.com/jason/>